

Big Data Infrastructure Backbone System

Big Data on Kubernetes - [Day 2]



LUAN MORENO
CEO & Data Architect
Data Engineer & MVP



MATEUS OLIVEIRA
Big Data Architect
Data In-Motion Specialist



Cloud Managed Services

IaaS, PaaS and SaaS Software for Easy and Fast Development for Enterprises



Market Share by 2020

AWS = 52%

Microsoft Azure = 21%

Google Cloud Platform = 18%



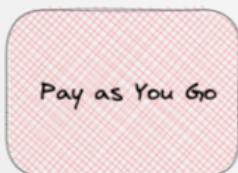
24x7 Customer Support



IaaS, PaaS and SaaS Softwares
for Easy and Fast Development



Managed by Cloud Provider



Pay as You Go

Most customers use resources on business hours and tears down after to reduce costs



Small Teams

Easy Collaboration using DevOps Principles
Autonomous and Automatic Processes
to Reduce Maintenance



Latest Versions

Latest Versions of the Products.
Battle-Tested before going to GA



Business-Oriented

Focused on Delivery and Success of Your Business

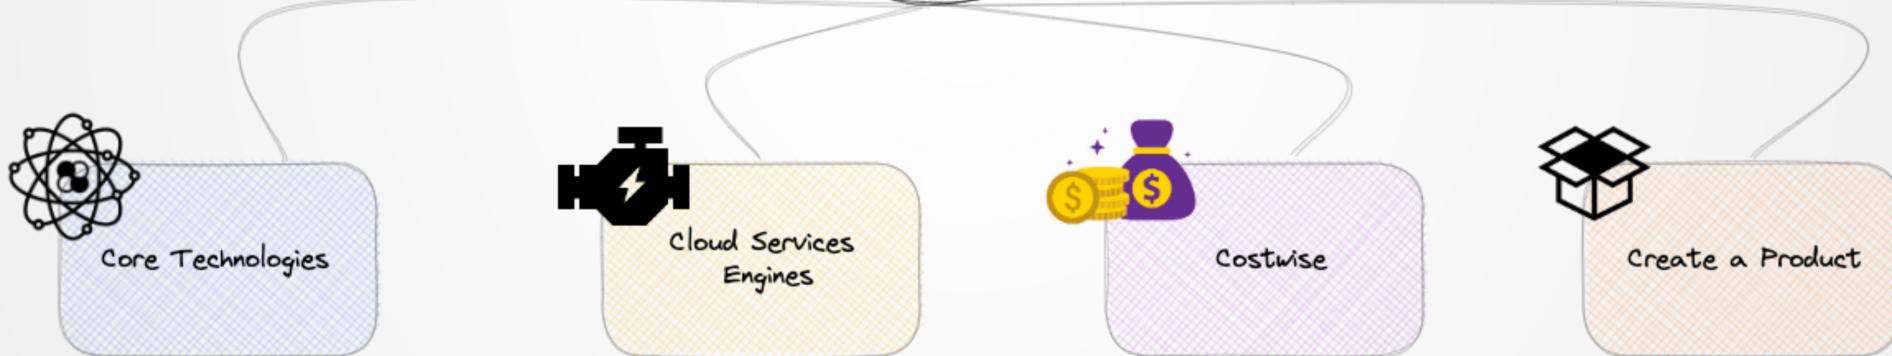
Open-Source Solutions

Best of Breed Big Data Technologies. Bleeding-Edge and Most Used in the Market



Open-Source Jobs Report

- 45% Demand for Kubernetes
- 97% Hiring OS is Priority
- 50% Increase in Hiring



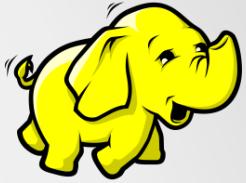
Open-Source, Worldwide Community, Feature Discussion, Constant Improvement, Vendor Neutral

Used by Cloud Providers Encapsulated as Managed Services for Easy-to-Use Approach

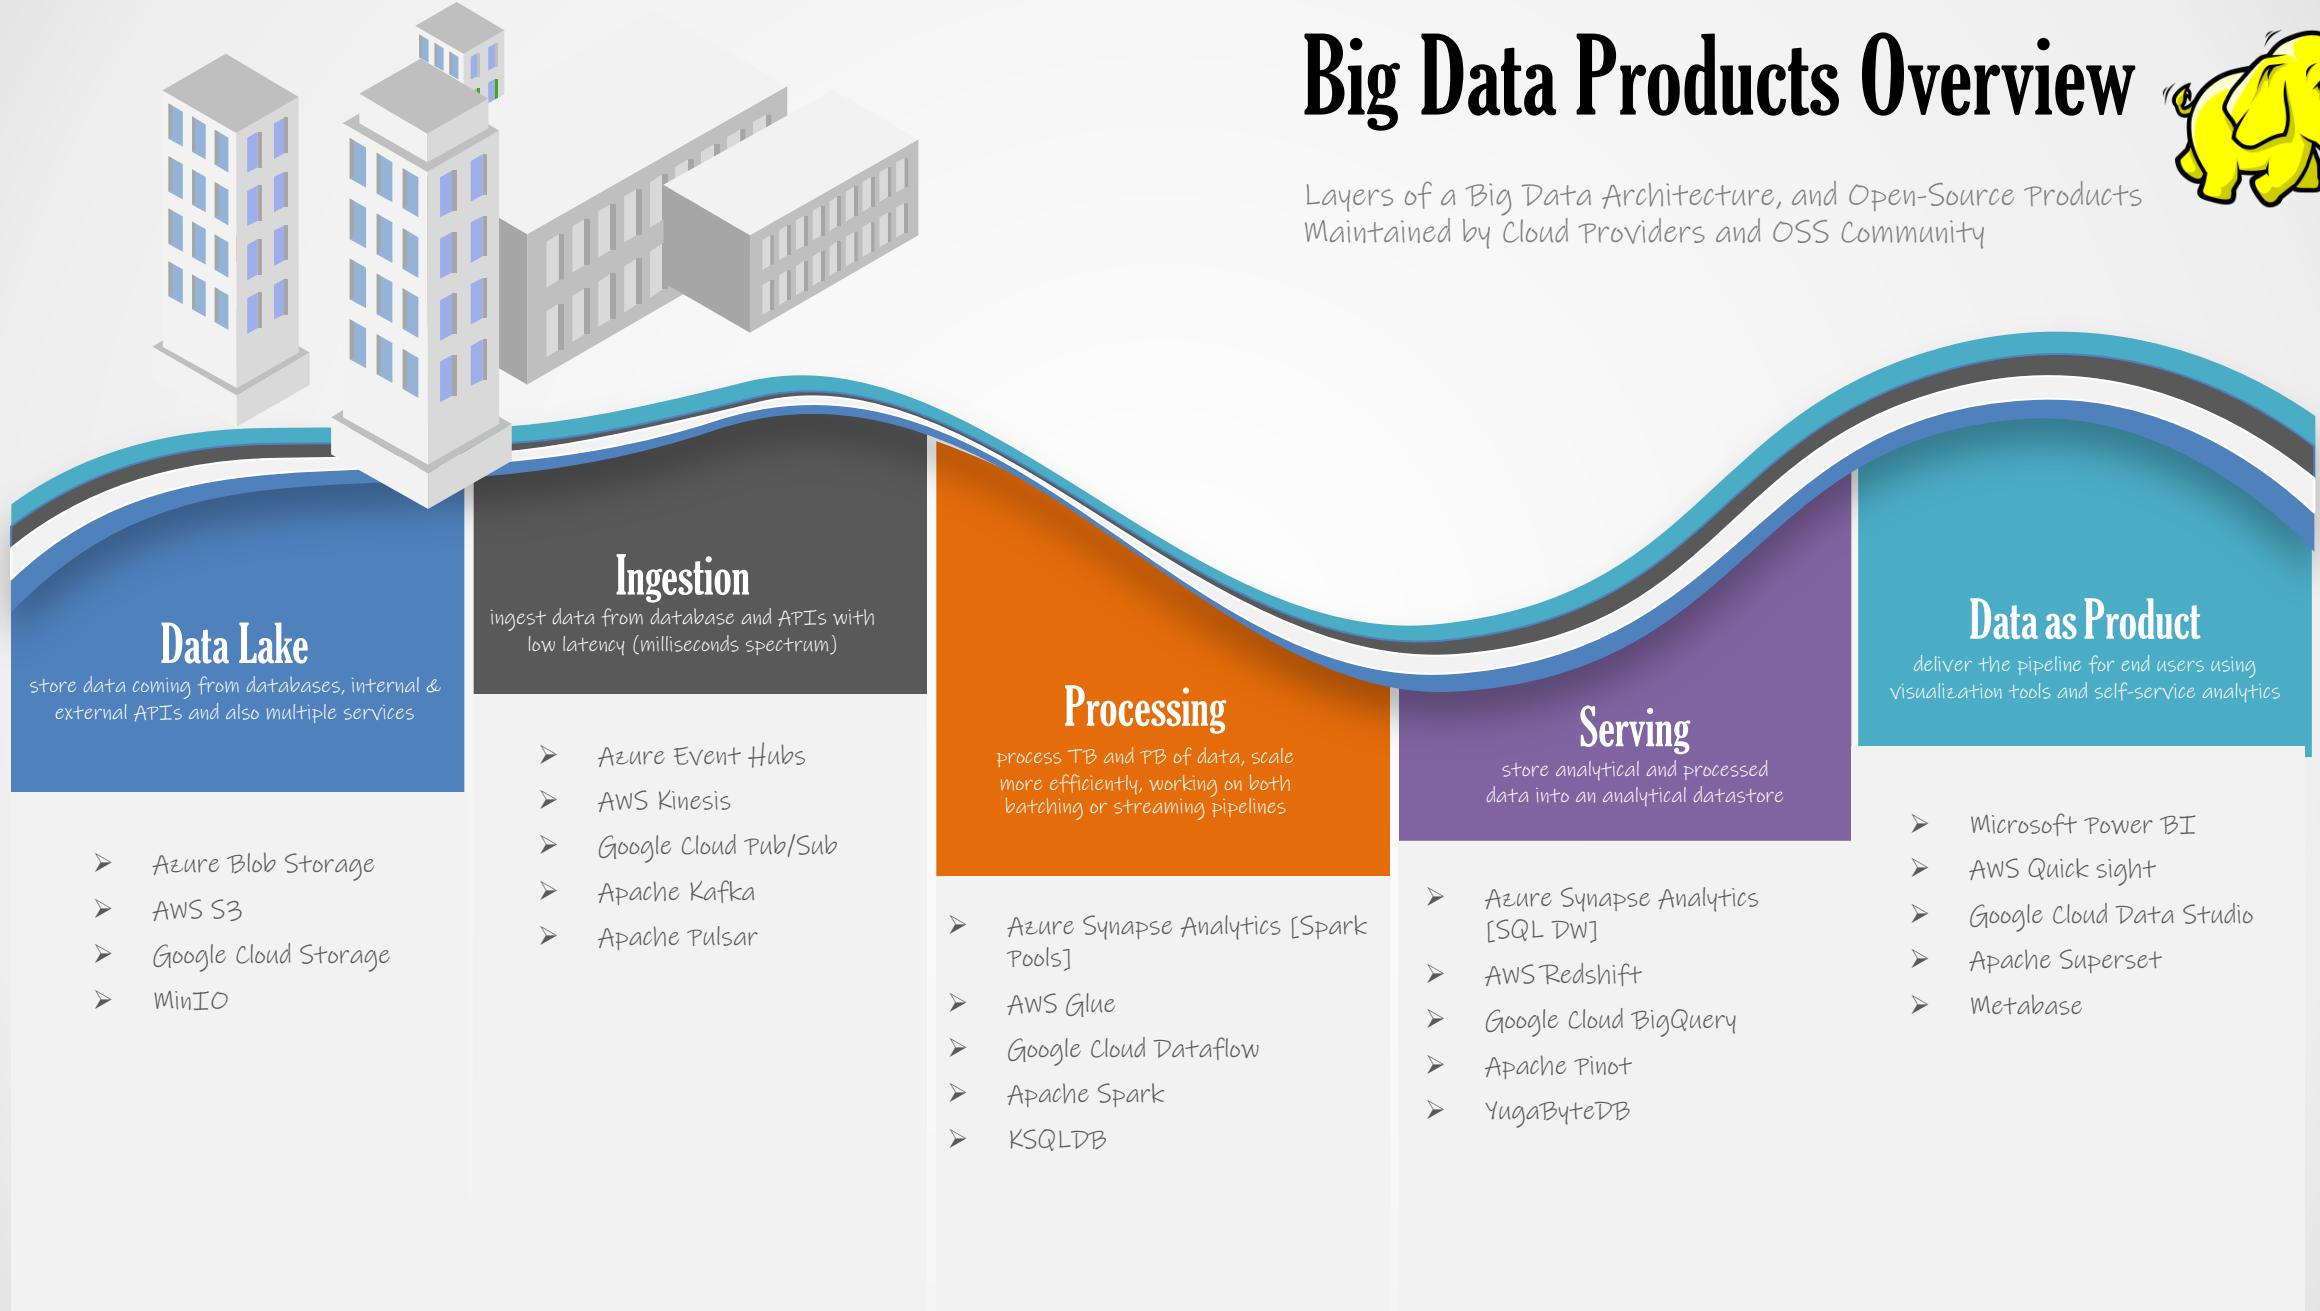
Without License Fees, No Lock-In and, Almost No Code Change for Cloud Transition

Create your Own Product, Without Lock-In and Low Entry Point for Business

Big Data Products Overview

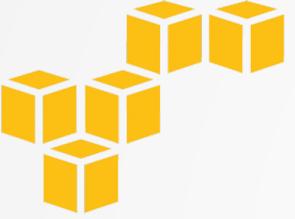


Layers of a Big Data Architecture, and Open-Source Products Maintained by Cloud Providers and OSS Community



Cost Comparison

Compare Cloud Provider's Costs against
Kubernetes Based Big Data Deployment



Total Cost for Data Pipelines on Azure & AWS

- Storage Layer = R\$ 2.414
- Data Processing Layer = R\$ 6.556
- Data Serving Layer = R\$ 14.580
- **Total Monthly Cost - R\$ 23.550**

Total Cost for Big Data Stack

- **Layers** = Ingestion, Processing, Serving, Visualization, Integration, Exploration, Orchestration, Logging, Monitoring, CI
- **Products** = Apache Kafka, Apache Spark, KSQLDB, MinIO, Apache Pinot, YugaByteDB, Trino, Zeppelin, Metabase, Superset, Apache Airflow, ELK, Prometheus, Grafana & ArgoCD
- **Category** = General Purpose
- **Series** = DSV4
- **Instance** = 4 vCPUs & 16 GB of RAM - [24 vCPUs & 96 GB of RAM]
- **VMs** = 6
- **Hours** = 730 Hs
- **Total Monthly Cost - R\$ 7.150 ~ 3x of Savings for ALL Layers**

Decision Points

Better Strategy to choose the right Infrastructure for Big Data



Moment of business

Evaluable the end goal of your pipelines and what are the main partnerships. Understand how to position the solution offered



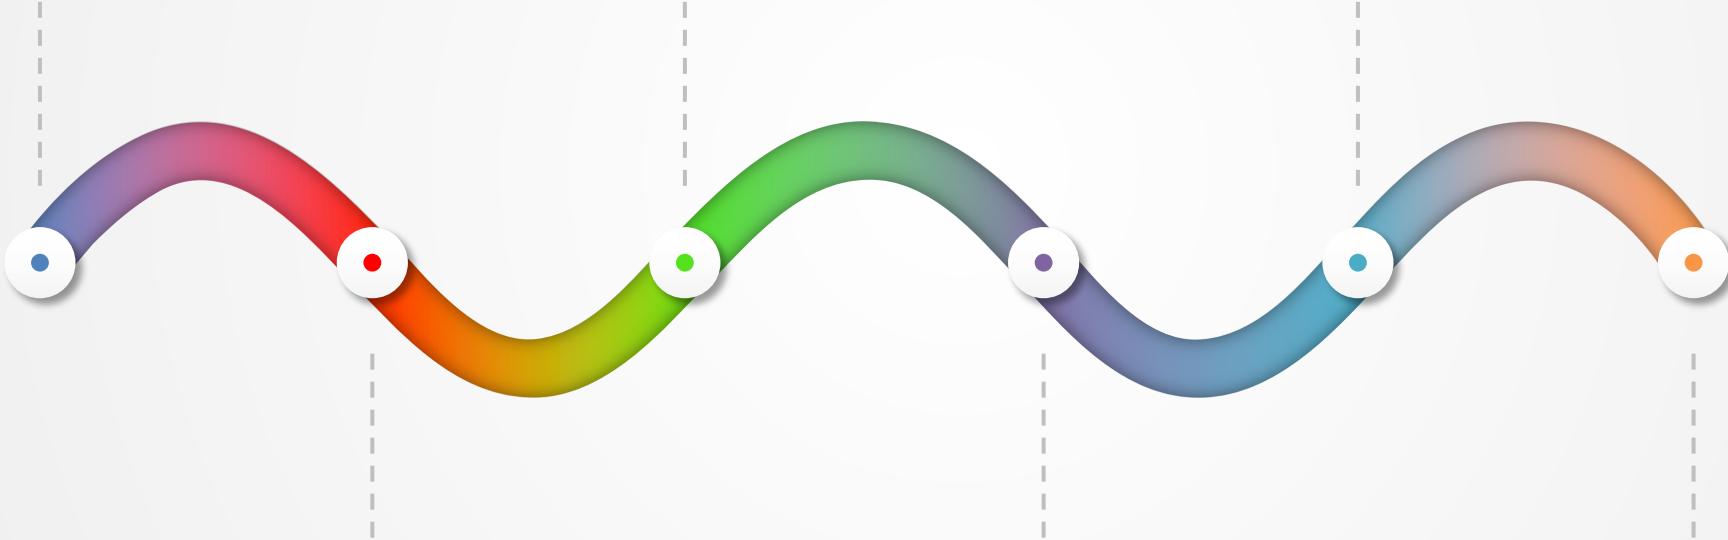
Budget

Not only understand, how much is going to cost but how much will cost to hire or train people to operate the solution



Sight of the future

Need to know where you want to be, and which technologies can deliver the best insights, and fast to migrate if the business need



Teamwise

Your team must be aligned with all the drawbacks in this new approach for big data architecture



Automate

Understand your data and the level you can automate to deliver pipelines that reduces overall complexity

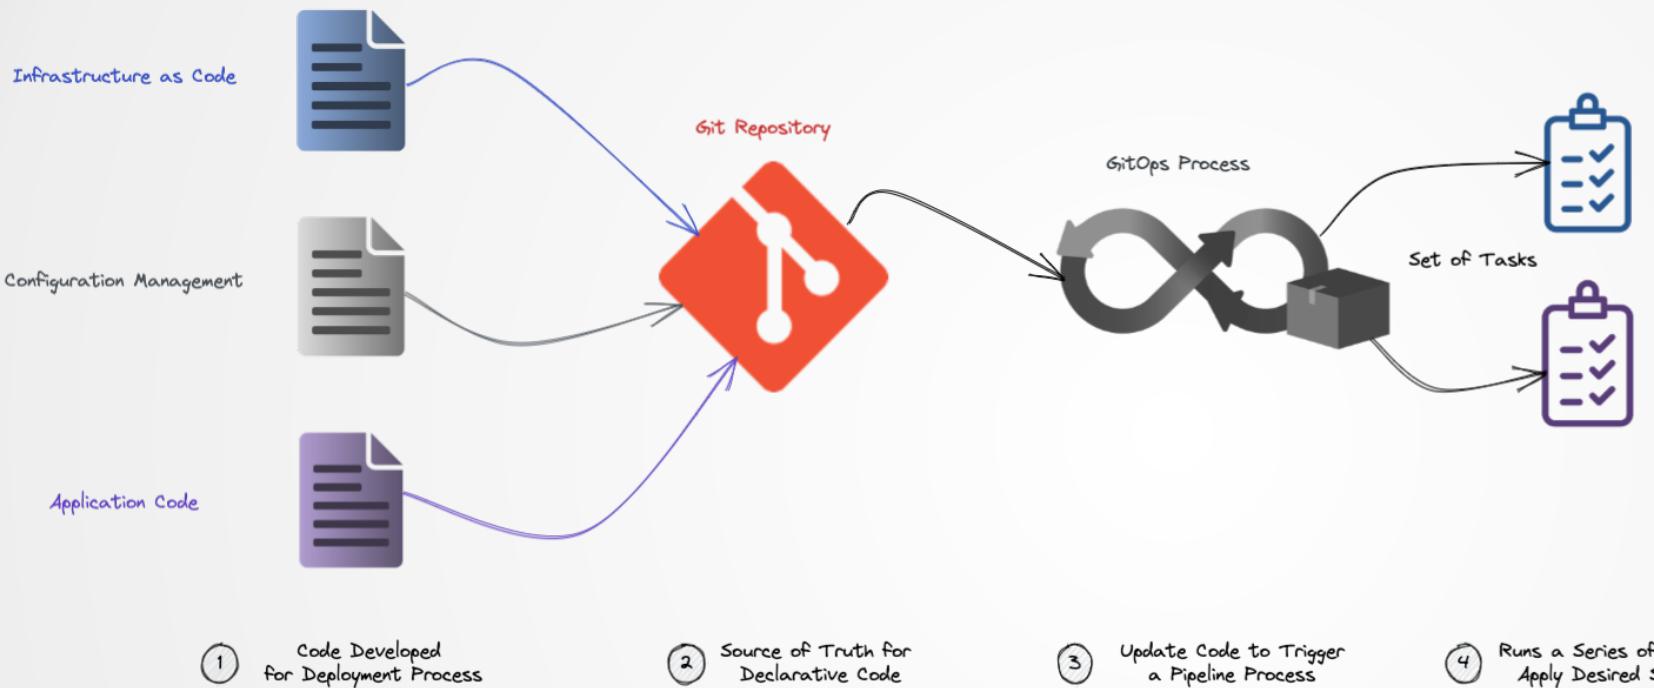


Goal

More that know the future is to now the present goal, what can be better for the business today and think about cost of any nature.

The Principles of GitOps Model

Set of Practices that Empowers Developers to Perform Tasks for CI/CD Lifecycle Process



DevOps

- Set of Practices to Combine Dev & IT Ops
- Shorten System Development Life Cycle
- Provide CI & CD Approaches
- Usually, Use of Agile Methodology

Continuous Integration (CI)

- Software Development Strategy
- Increase Speed of Development
- Ensure Quality of Code
- Continually Commit Code in Small Increments
- Automatically Built and Tested Before Merged
- Use of a Git Repository

Continuous Delivery (CD)

- Software Delivered to Production at Any Time
- Produce Software in Short Cycle
- Reliably Release at Any Time
- Automatic Process
- Build, Test and Release

GitOps Workflow Process



- Set of practices to Manage Infrastructure and Application
- Use Git as a Single Source of Truth for Declarative Process
- Git Pull Requests to Automatically Manage Deployments
- Apply Desired State of System
- Enable CD for Kubernetes

Used for Kubernetes and Cloud Native Applications



Automate Continuous Delivery (CD) using ArgoCD Tool

Declarative GitOps Tool for Kubernetes. Application Deployment and Lifecycle Management

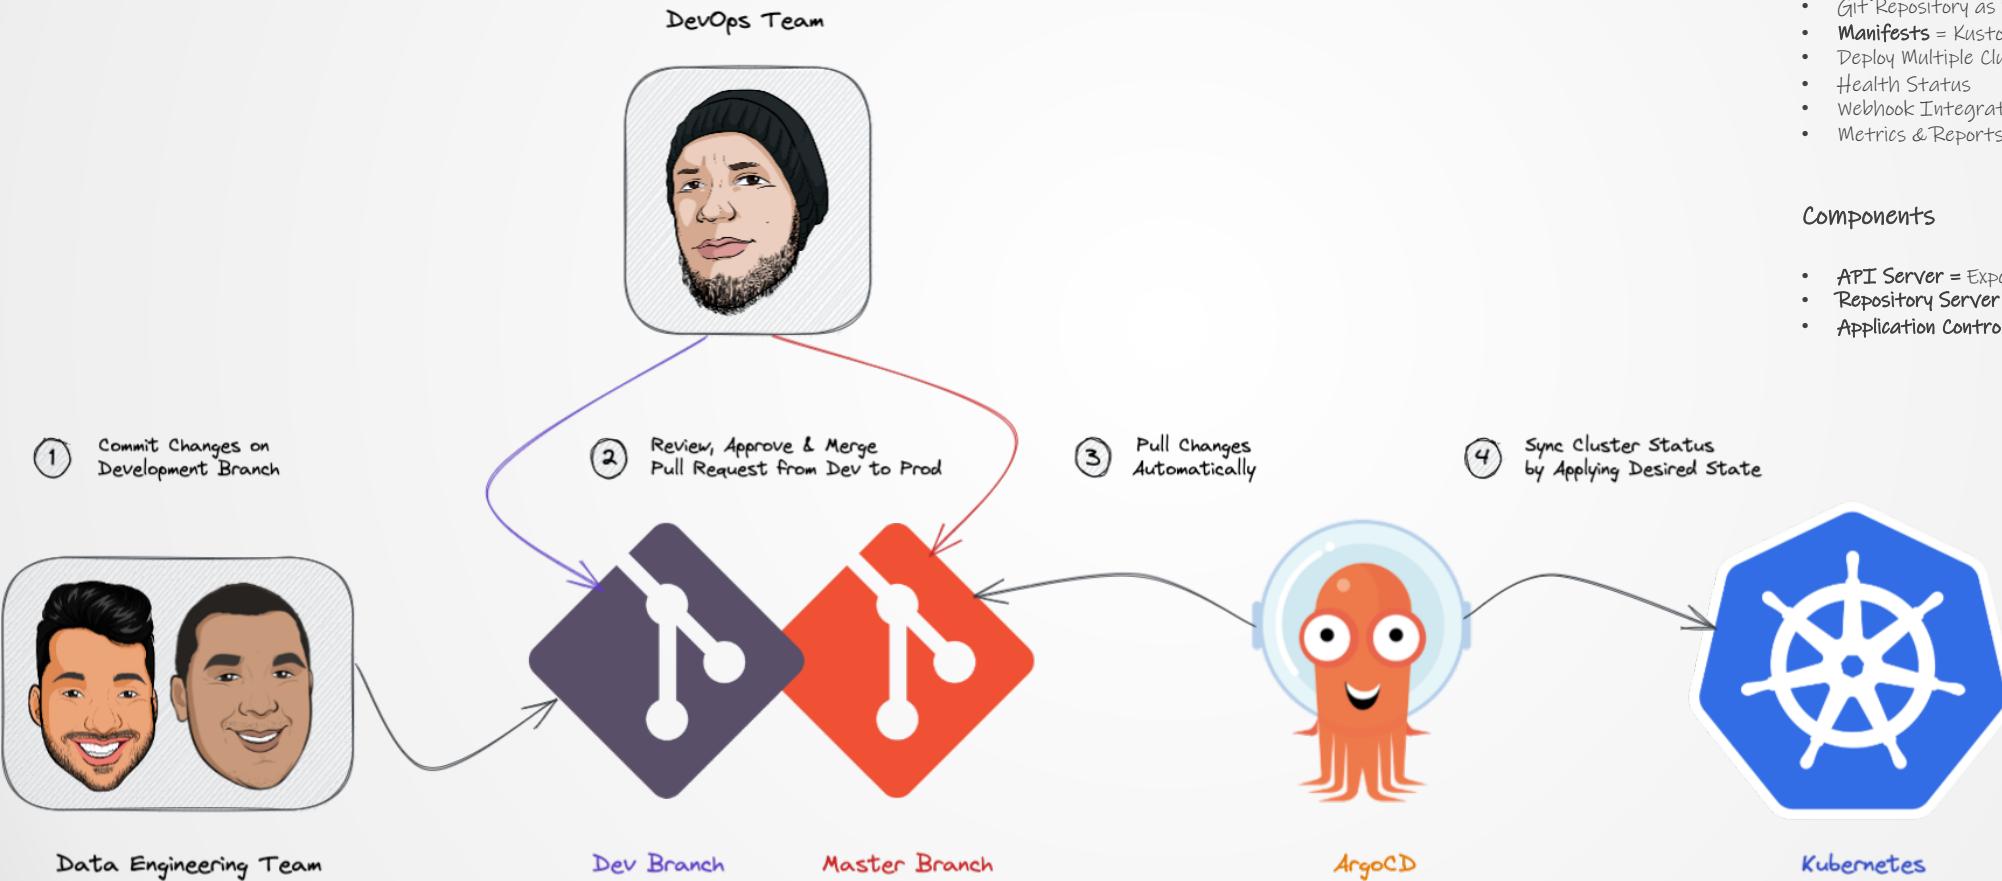


ArgoCD

- Git Repository as Source of Truth
- **Manifests** = Kustomize & Helm
- Deploy Multiple Clusters
- Health Status
- Webhook Integrations
- Metrics & Reports

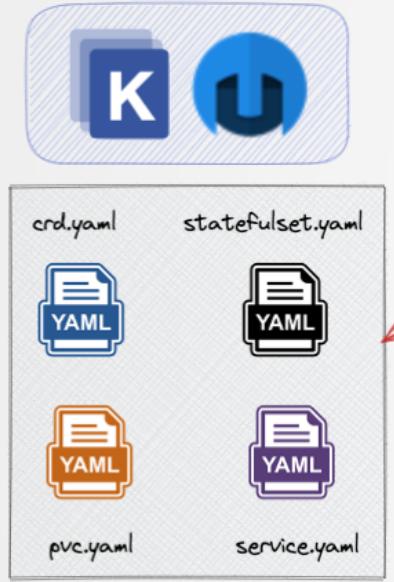
Components

- API Server = Expose API for Application Management
- Repository Server = Maintains Local Cache of Git
- Application Controller = Continuously Monitors Apps

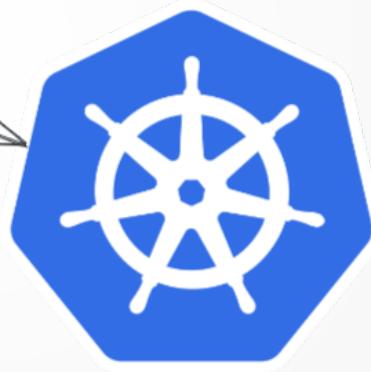


Big Data Products Deployment Workflow

Using ArgoCD to Deploy Big Data Stack using GitOps Process



ArgoCD

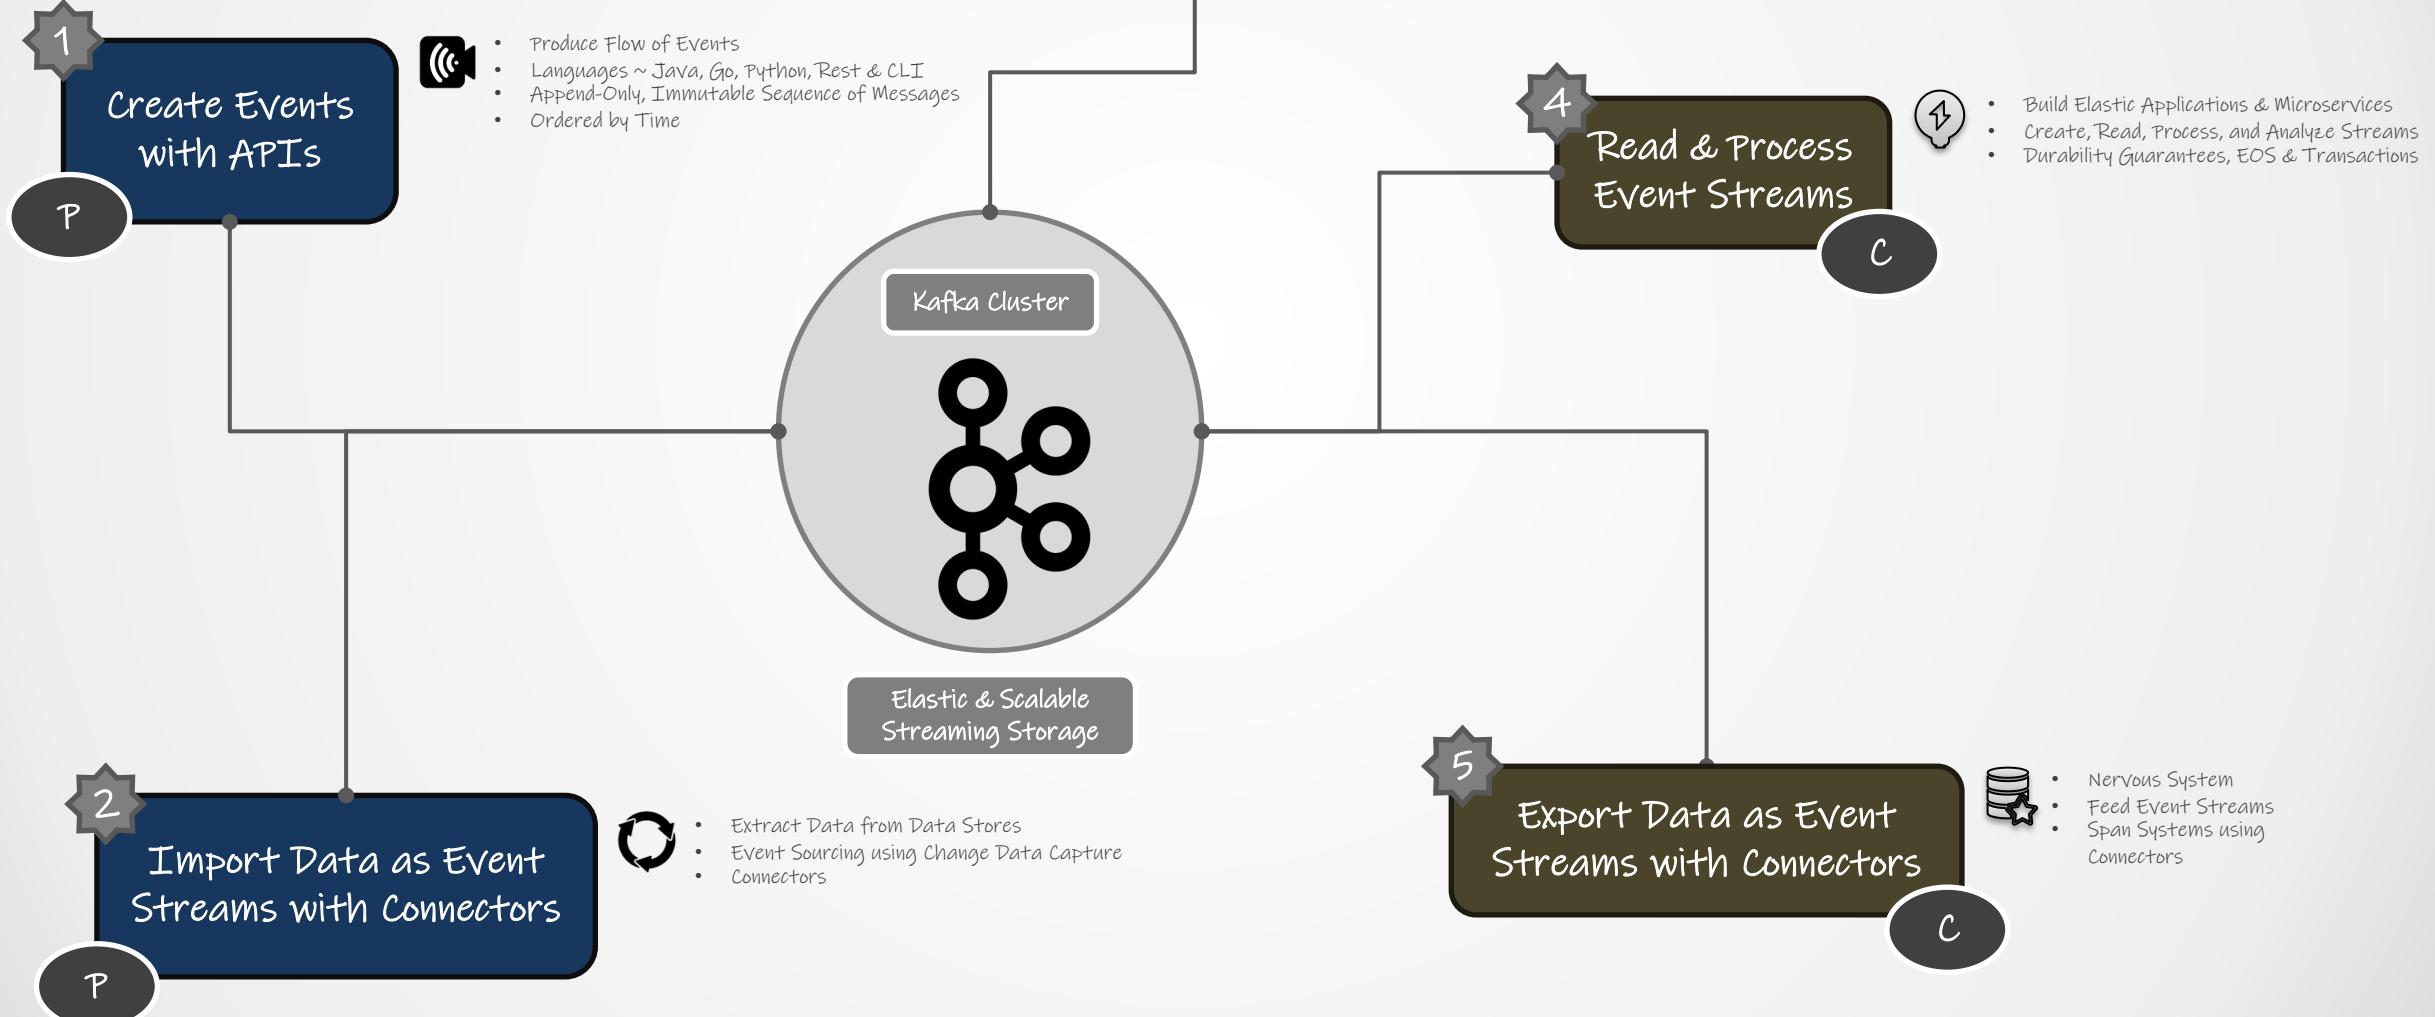


Kubernetes



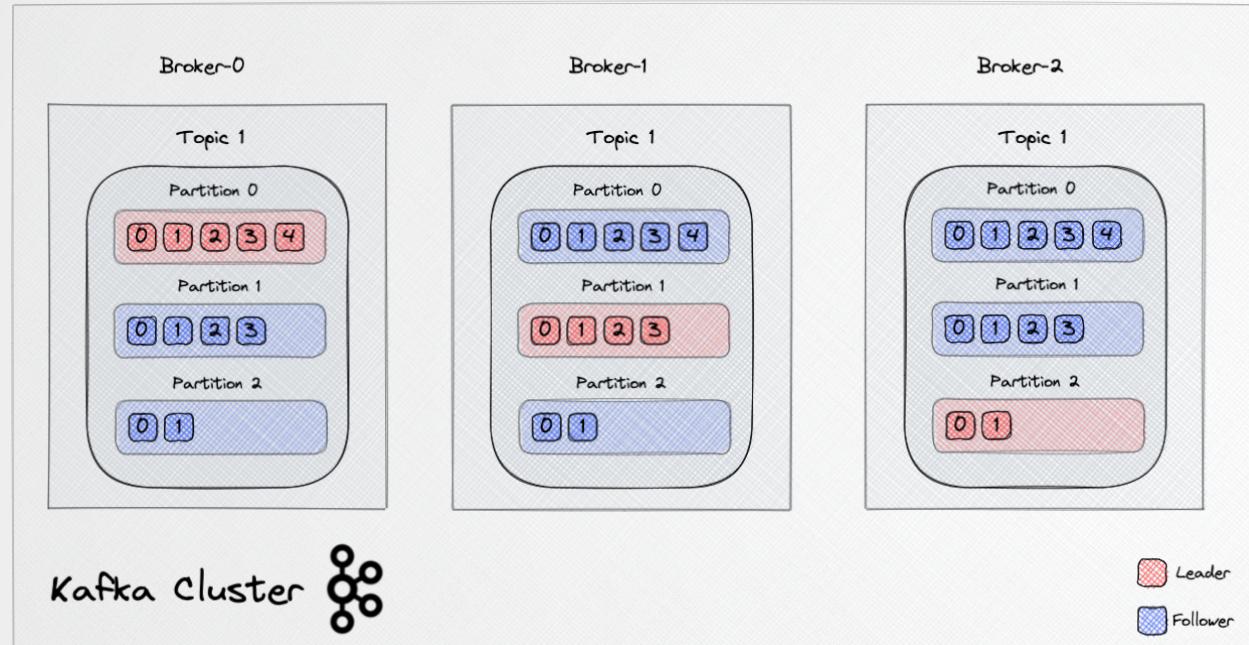
Backbone System

De-Facto Distributed Event Streaming Platform for Data Pipelines, Analytics, Data Integration and Mission-Critical Apps



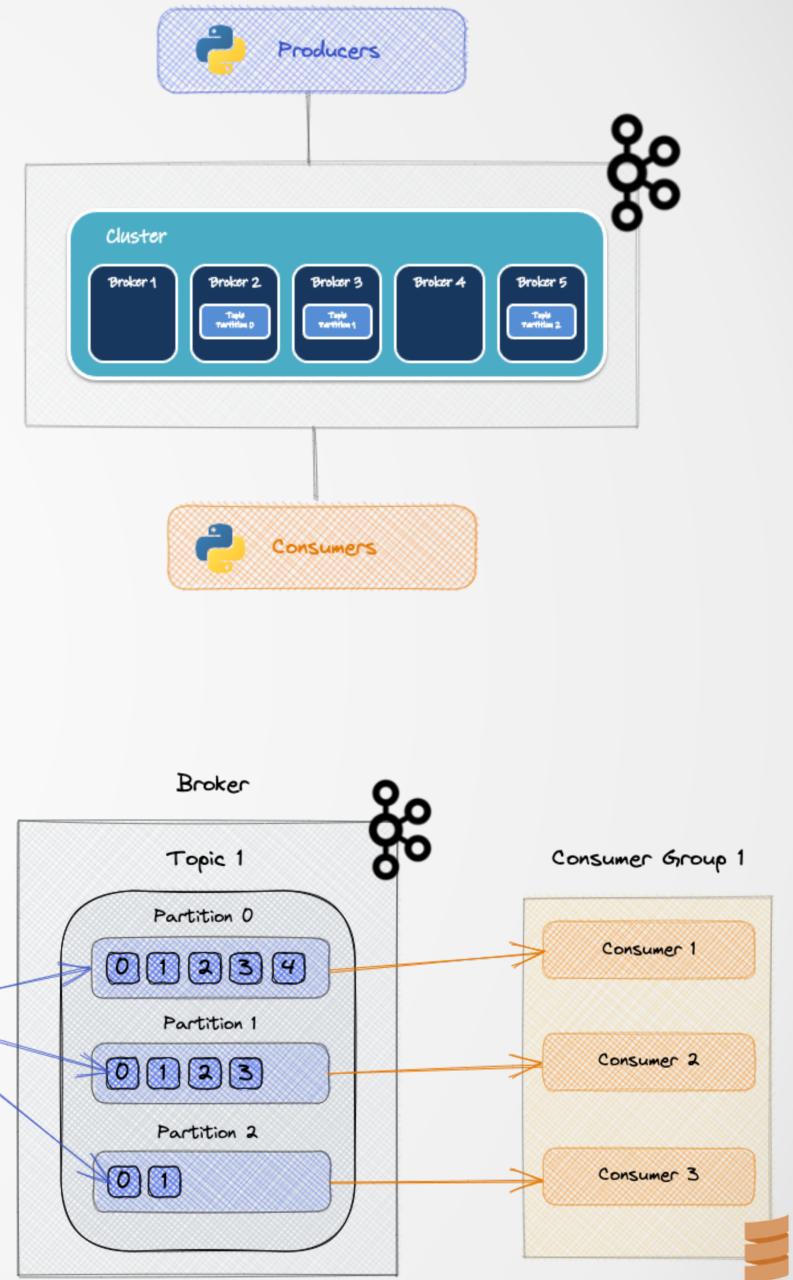
Components & Architecture

Main Components and Architecture of Apache Kafka, Available Features and Characteristics



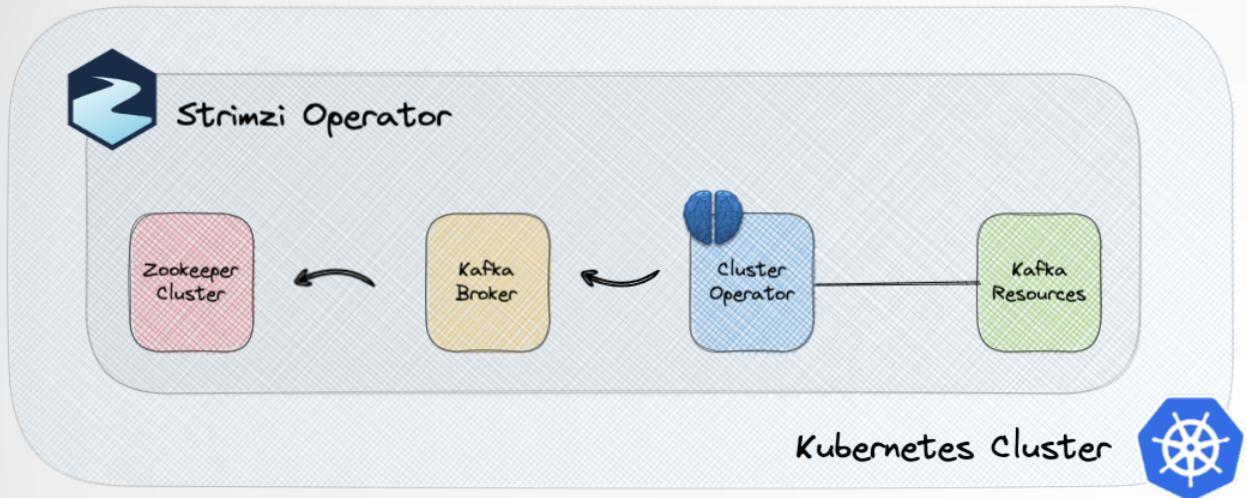
Characteristics

- Kafka Broker = Server or Node, Orchestrates Storage
- Cluster = Group of Broker Instances
- Topic = Main Entry, Split into One or More Partitions
- Partition = Horizontal Partitioning for Writes
- Partition Leader = Handles ALL Producer Requests
- Partition Follower = Replicates Partitions from Leader



Strimzi - Kafka on Kubernetes

The Best Way to Run Kafka on Kubernetes in Various Deployment Flavors, From Development to Production Effortlessly



Storage Types

Apache Kafka & Zookeeper Store Data on Disks

- * Ephemeral
- * Persistent
- * JBOD



Listener Types

How Client Connects to a Kafka Cluster

- * Route
- * Load Balancer
- * Node Port
- * Ingress
- * In-Transit mTLS
- * SALS SCRAM-SHA-512
- * OAuth 2.0 Token Based
- * ACLs



Security Options

Perform a Secure Deployment of Apache Kafka on Kubernetes



- Cluster = Deploys and Manages Resources
- Entity = Topic & User Operator
- Topic = Manages Kafka Topics
- User = Manages Kafka Users

- Kafka Broker = Cluster of Broker Nodes
- Zookeeper = Replicated Instances
- Kafka Connect = External Data Connections
- Kafka MirrorMaker = Mirror for Secondary Cluster
- Kafka Exporter = Extract Metrics for Monitoring
- Kafka-Bridge = HTTP-Based Request Component
- Cruise Control = Optimizes Rebalance and Anomalies

Features

Kubernetes Native Multi-Cloud Open-Source Object Storage Available on Public, Hybrid and Private Cloud



Information

- **Interfaces** = Client, Admin, Server, Gateway and Console
- **Notifications** = ElasticSearch, Kafka, Postgres, Redis
- **Gateway** = NAS, Azure, S3, HDFS, GCS



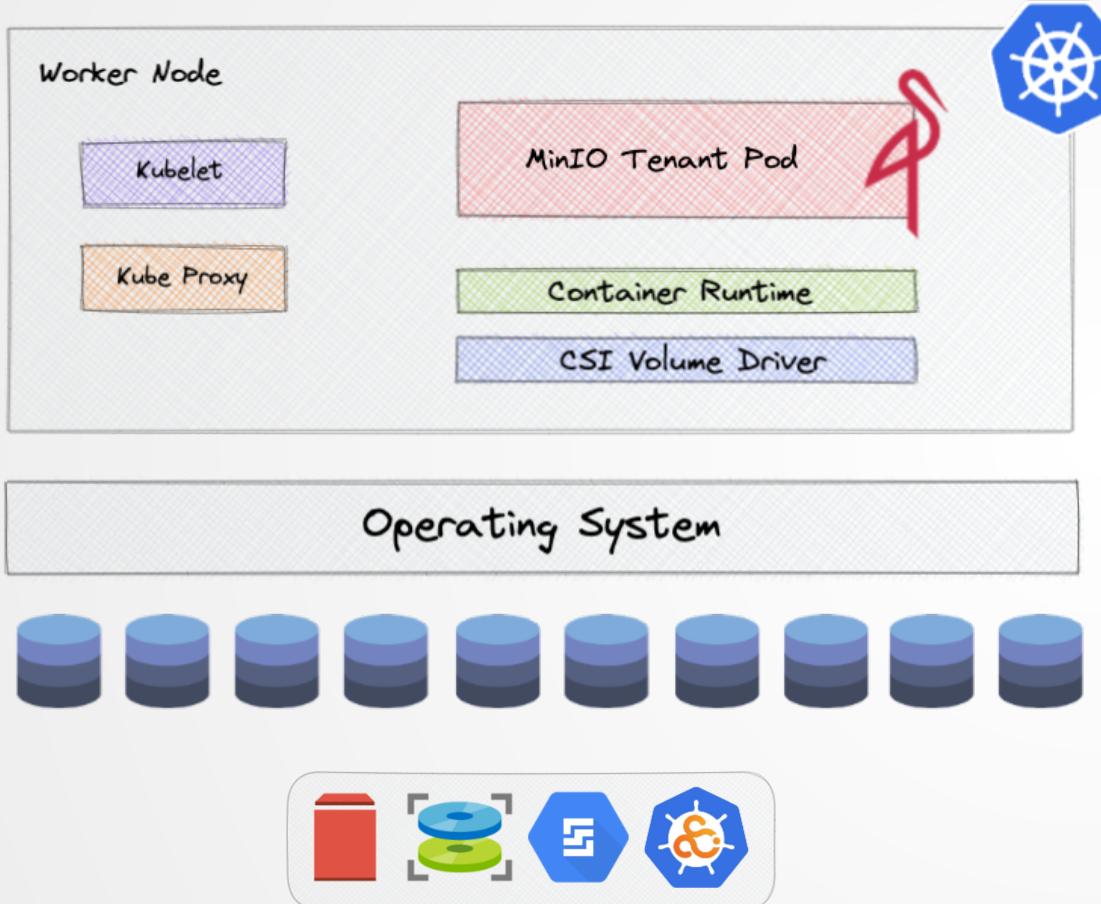
High Performance

World's Fastest Object Storage Server

Read & Write of **183 GiB/s** and **171 GiB/s** on Standard Hardware

MinIO and Kubernetes

Deployment Spec for MinIO on Kubernetes,
Architecture Diagram and Killer Features



Volume Type

Directory Determined
by a Volume Type

- * AWS EBS
- * Azure Disk
- * Google Persistent Disk
- * Container Storage Interface

Features

- Active-Active Replication
- Identity & Access Management
- Encryption
- Bucket & Object Immutability
- Bucket & Object Versioning
- Data Lifecycle Management and Tiering
- Automated Data Management Interfaces
- Monitoring and Scalability

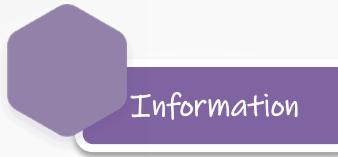
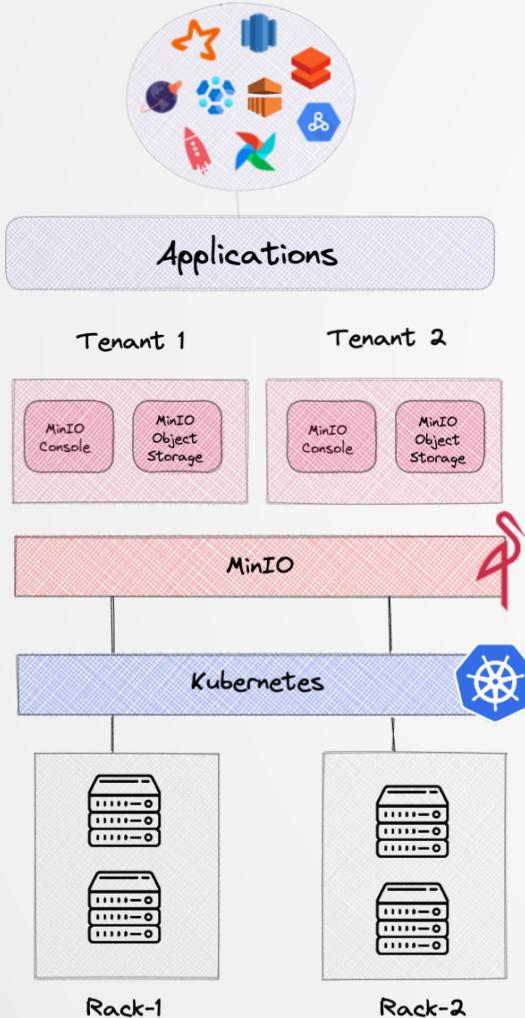
Container Storage Interface (CSI)

Standard for Exposing Arbitrary Block and File Storage Systems
for Container Orchestration Systems (COs)

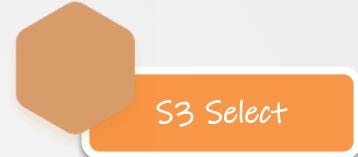


MinIO Operator Architecture

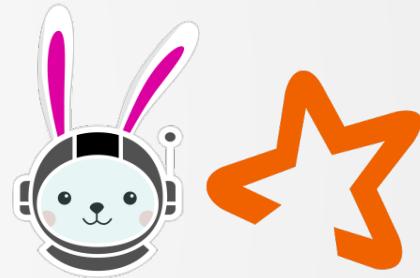
Kubernetes Friendly Object-Storage System for Big Data Workloads with High-Speed Response Time



- Kubernetes Native DNA
- Use of Operator Pattern (CRDs) using Helm Chart
- Resource Orchestration
- Non-Disruptive Upgrades
- Cluster Expansion
- High-Availability

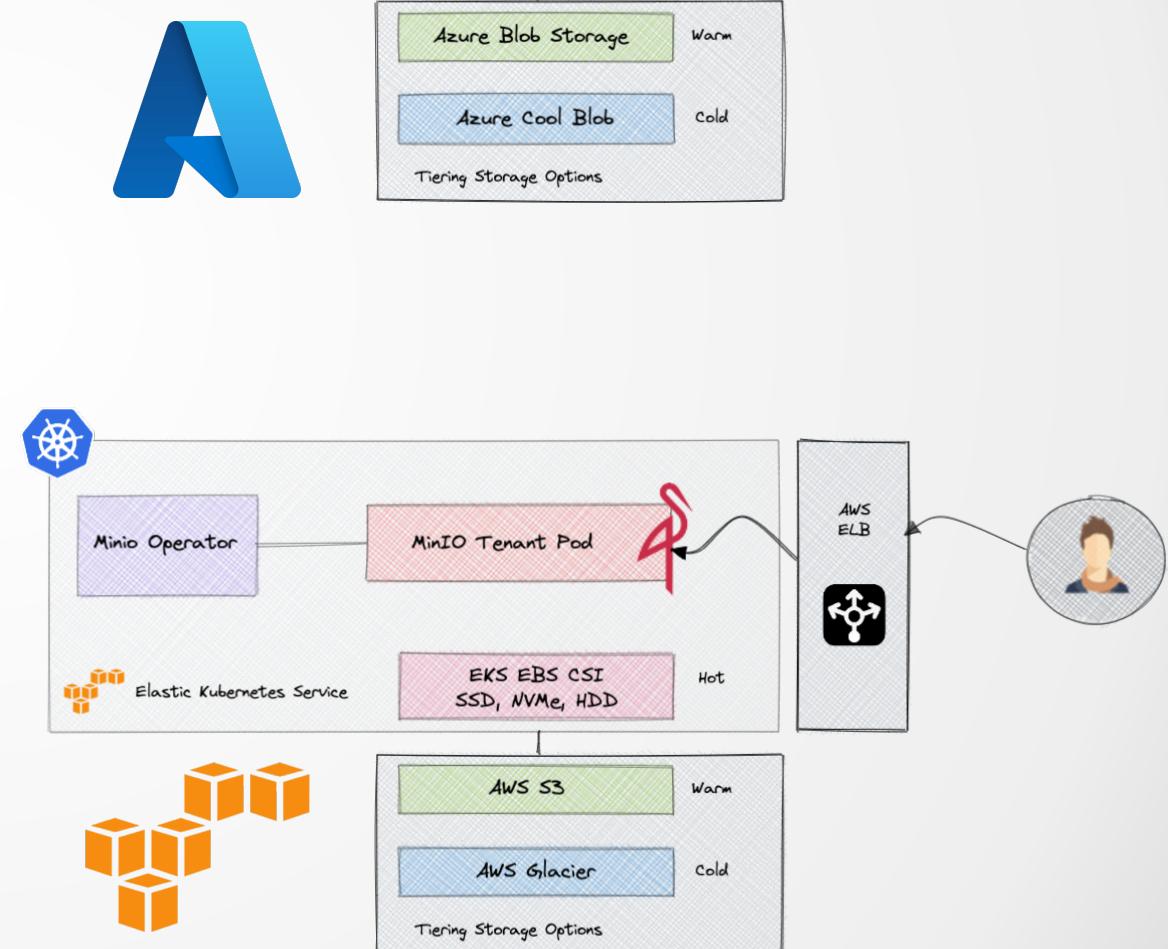
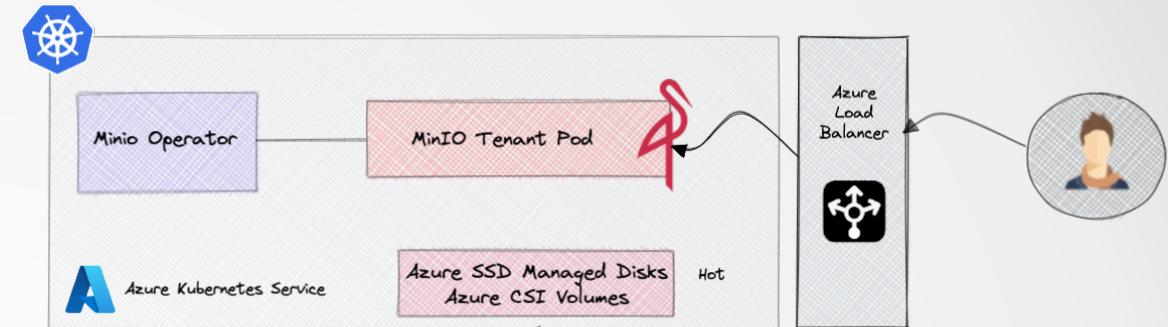
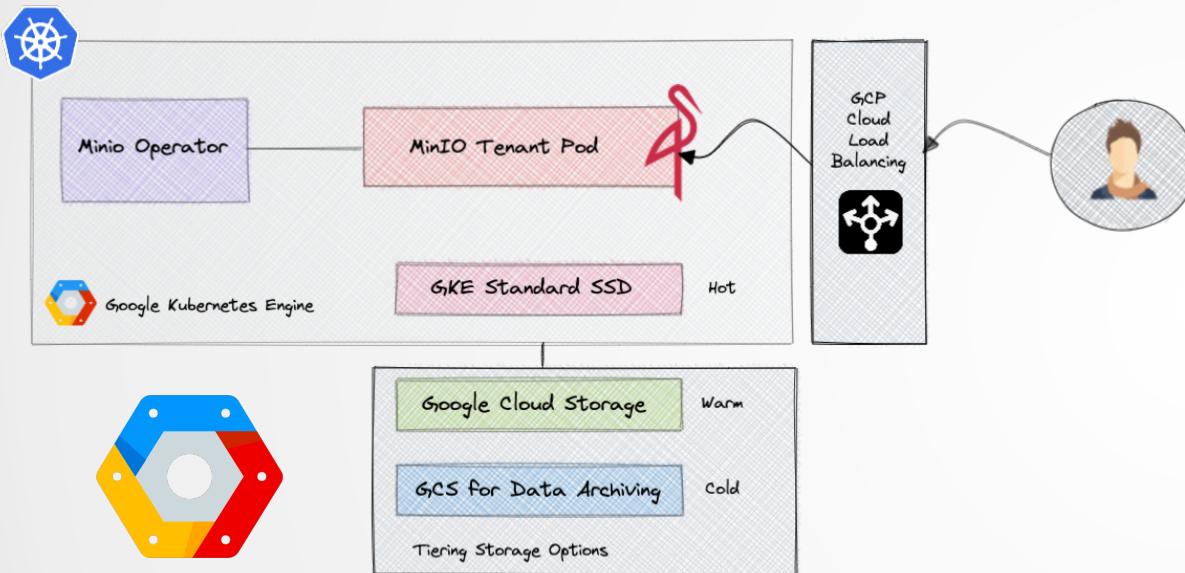


- SQL Expressions to Retrieve a Subset of S3 Object
- SELECT & WHERE Clause to Fetch Data
- CSV, JSON, or Apache Parquet
- GZIP or BZIP2 Compression
- Python, Spark & Presto Engines



MinIO & Managed Kubernetes

Managed Kubernetes Offerings - GKE, AKS & EKS and
MinIO's Tight Cloud Integration



Deployment Modes



Methods in how to Submit an Apache Spark Application
to the Server. Different Modes available

Local

Runs on a Single JVM ~ Laptop or Single Node,
Executor and Cluster Manager Runs on Same Host



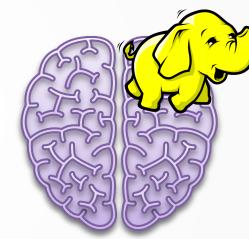
Standalone

Run on Any Node in Cluster. Each Node Launch Own
Executor JVM. Allocated Arbitrarily to Any Host



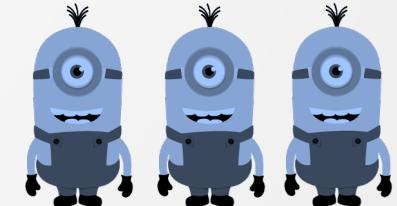
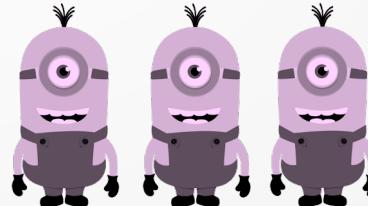
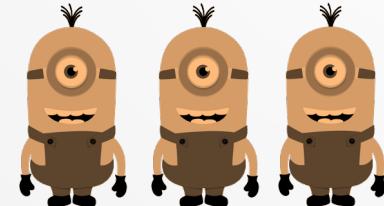
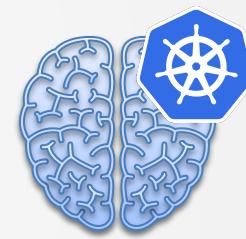
Yarn

Runs with YARN (Cluster) Application Master. Node
Manager & Resource Allocate Containers for Execution



Kubernetes

Runs in a Kubernetes Pod. Each Worker Runs Within on Pod
Context, use Kubernetes Master for Cluster Management



Deployment Options

Different Options to Deploy Apache Spark Cluster either on Cloud or K8S



Azure HDInsight

Parallel Processing Framework with In-Memory Support to Boost Performance of Big Data Analytics

- HDInsight 4.0
- Ubuntu 16.0.4 LTS
- Apache Spark 3.0.0



Kubernetes Operator for Apache Spark: Running Spark Applications at Scale with Low and Affordable Cost

- Apache Spark 3.1.1



Spark Pools

Microsoft's Implementation of Apache Spark in Azure Synapse Compatible with Azure Storage

- Apache Spark 3.0.0
- Delta Lake Integration
- .NET for Apache Spark



Amazon EMR

Industry-Leading Cloud Big Data Platform for Processing Vast Amount of Data Using Open-Source Tools

- EMR 6.3.0
- Linux AMI
- Apache Spark 3.1.1



Google DataProc



Fully Managed and Highly Scalable Service for Running Apache Spark & +30 OSS

- BORG = Kubernetes
- Debian, Ubuntu & CentOS
- Apache Spark 3.1.1



Data + AI Company. Open-Source Apache Spark, Founded in 2013 by Original Creators of Apache Spark

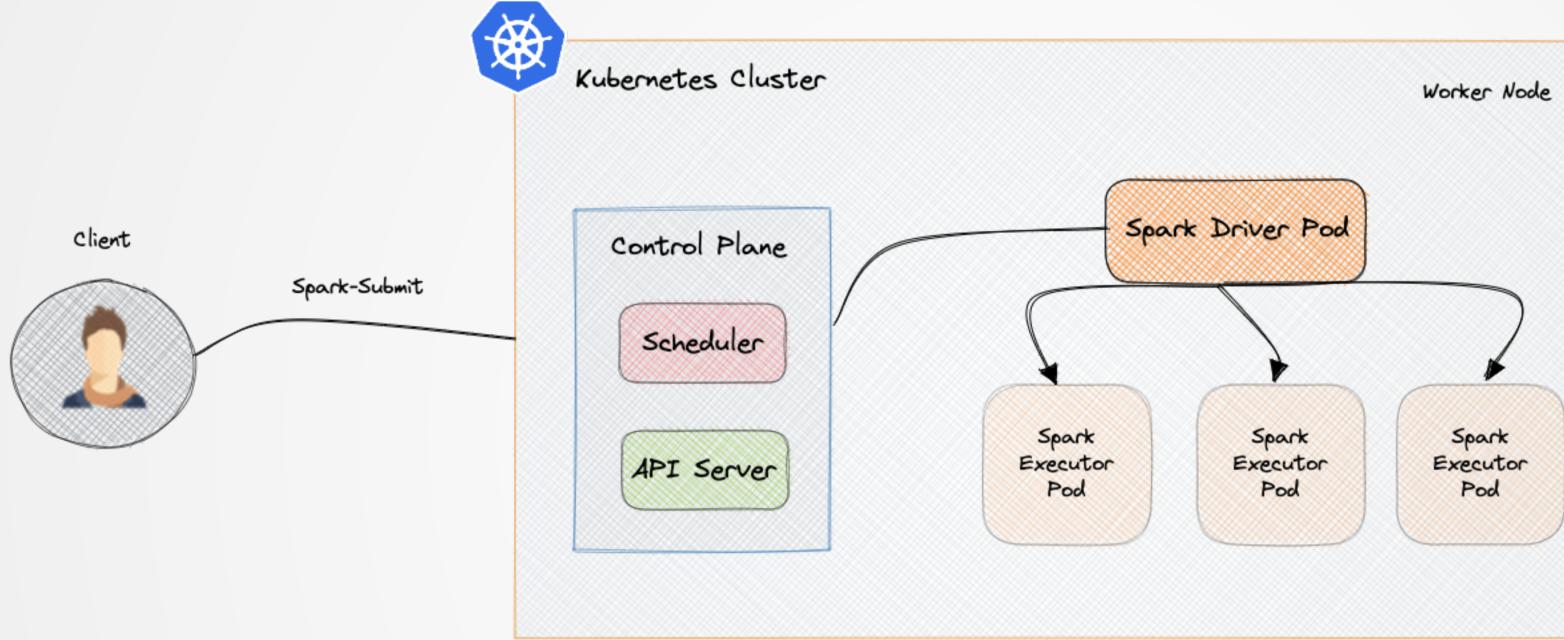
- Apache Spark
- Delta Lake
- MLFlow
- SQL Analytics
- Runtime 8.3
- Apache Spark 3.1.1

Spark on Kubernetes

Use of Native Kubernetes Scheduler for Manage Spark Applications in a Lightweight Fashion

Apache Spark GA in 3.1

Typically - 50% to 70% of Cost Savings



- Use of SSDs or Large Disks for Best-Shuffle Performance
- Optimize Pod Sizes to Avoid Wasting Capacity (90%)
- `spark.executor.cores=4`
`spark.kubernetes.executor.request.cores=3600M`
- Enable App-Level Dynamic Allocation
`spark.dynamicAllocation.enabled=true`
`spark.dynamicAllocation.shuffleTracking.enabled=true`
- Enable Cluster-Level Autoscaling
- Use Spot Nodes to Reduce Cloud Costs

Spark-on-Kubernetes ~ Operator

Operator to Deploy and Manage the Lifecycle Process of an Apache Spark Application on Kubernetes



Spark-Submit

Submit an Apache Spark Program and Launches on Cluster



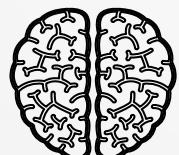
Spark-on-K8s-Operator

Kubernetes Operator ~ Method of Packaging, Deploying and Managing an Application. Extends Functionality to Create, Configure, and Manage Instances of Complex Applications

Spark-on-K8s-Operator ~ Aims to Specifying and Running Apache Spark Applications as Easy and Idiomatic as Running Kubernetes Workloads Instances

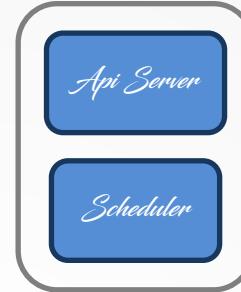


`spark-pi.yaml`



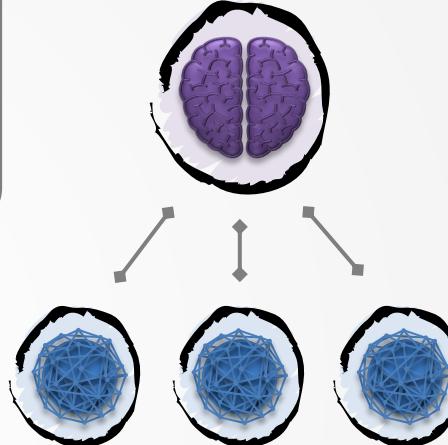
Spark Operator

- Controllers
- Submission Runner
- Pod Monitor
- Mutating Admission



Spark Driver

Creates Spark Context, Session, & Executes User Code

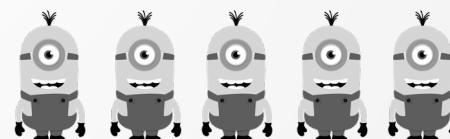


Spark Executor

Executes Individual Tasks & Return Result Set to Driver



Kubernetes

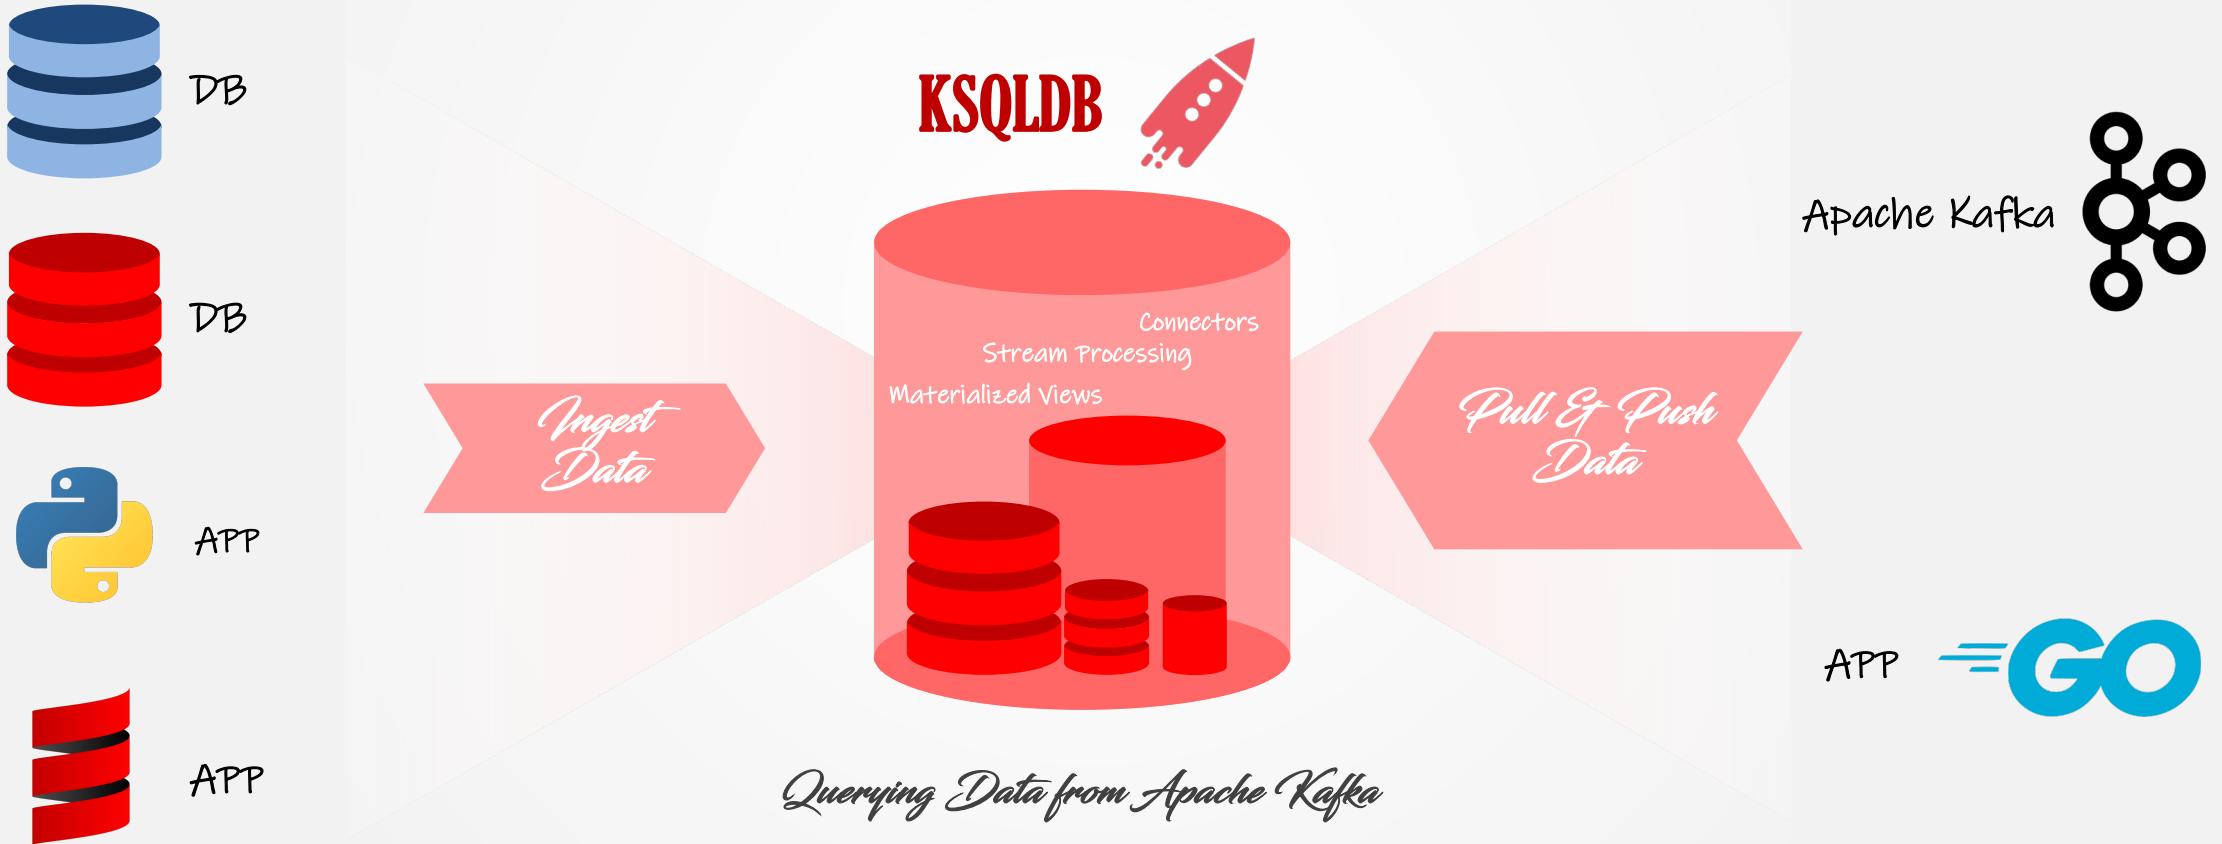


Deployment Mode

Runs in a Kubernetes Pod. Each Worker Runs Within on Pod Context, use Kubernetes Master for Cluster Management

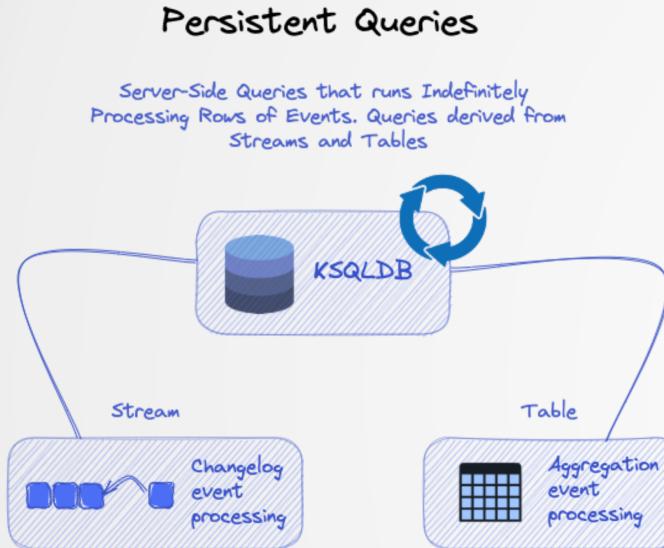
Single Mental Model for Event Streams

The Database Purpose-Built for Stream Processing Applications
that uses Apache Kafka as Backbone System



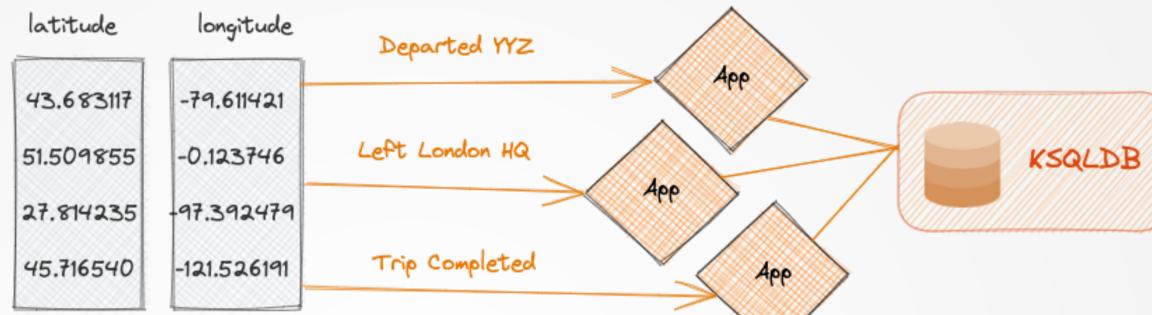
Query Modes

Kinds of Query Models: Push, Pull and Persistent. Different Ways to Work with Rows of Events in Real-Time



Push Queries

Subscribe to a query's result as it changes in real-time. Emit refinements to react quickly



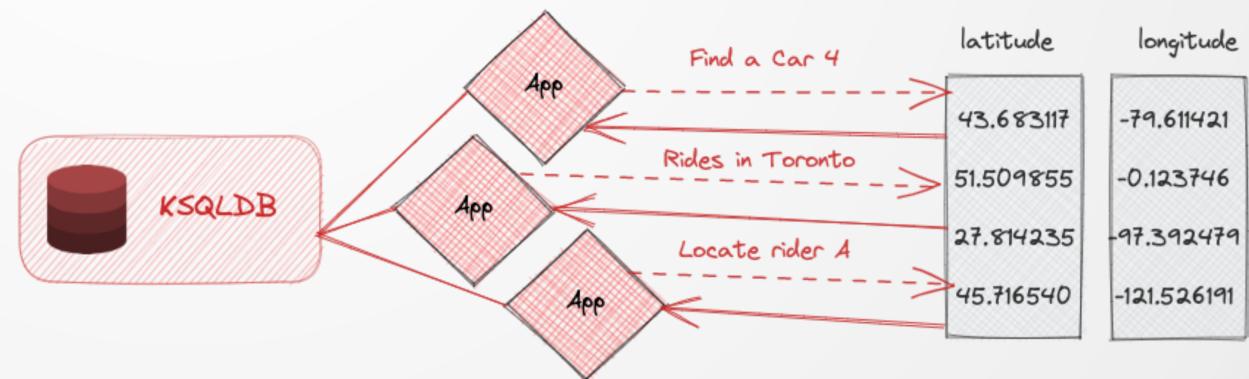
```
SELECT * FROM pageviews  
WHERE ROWTIME >= 1510923225000  
AND ROWTIME <= 1510923228000  
EMIT CHANGES;
```

Pull Queries

Get current values from the Materialized View and Terminates. Result not Persisted in a Kafka Topic

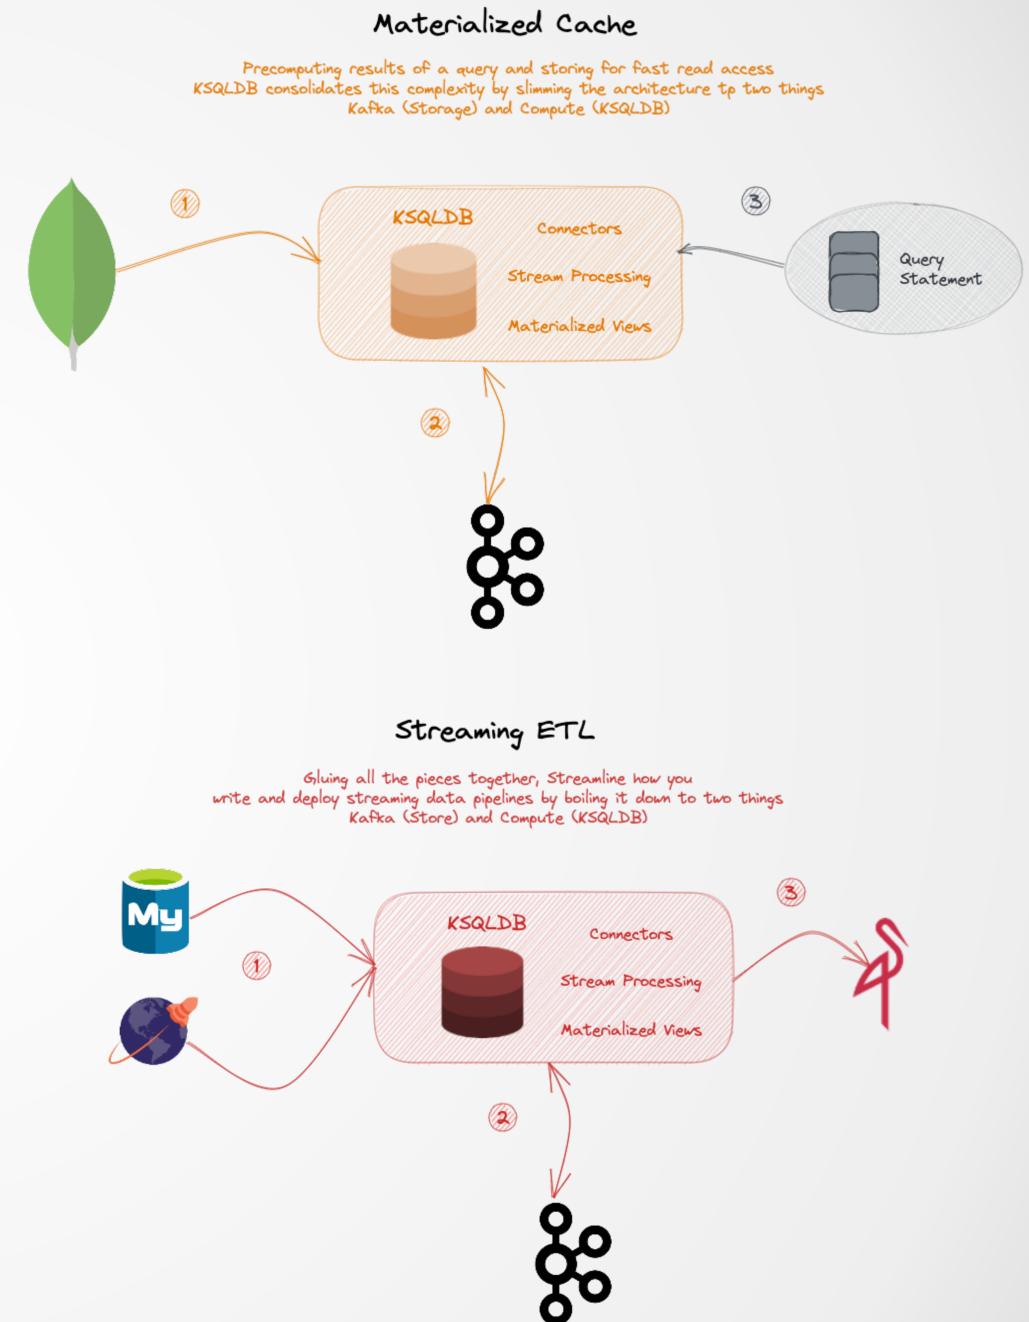
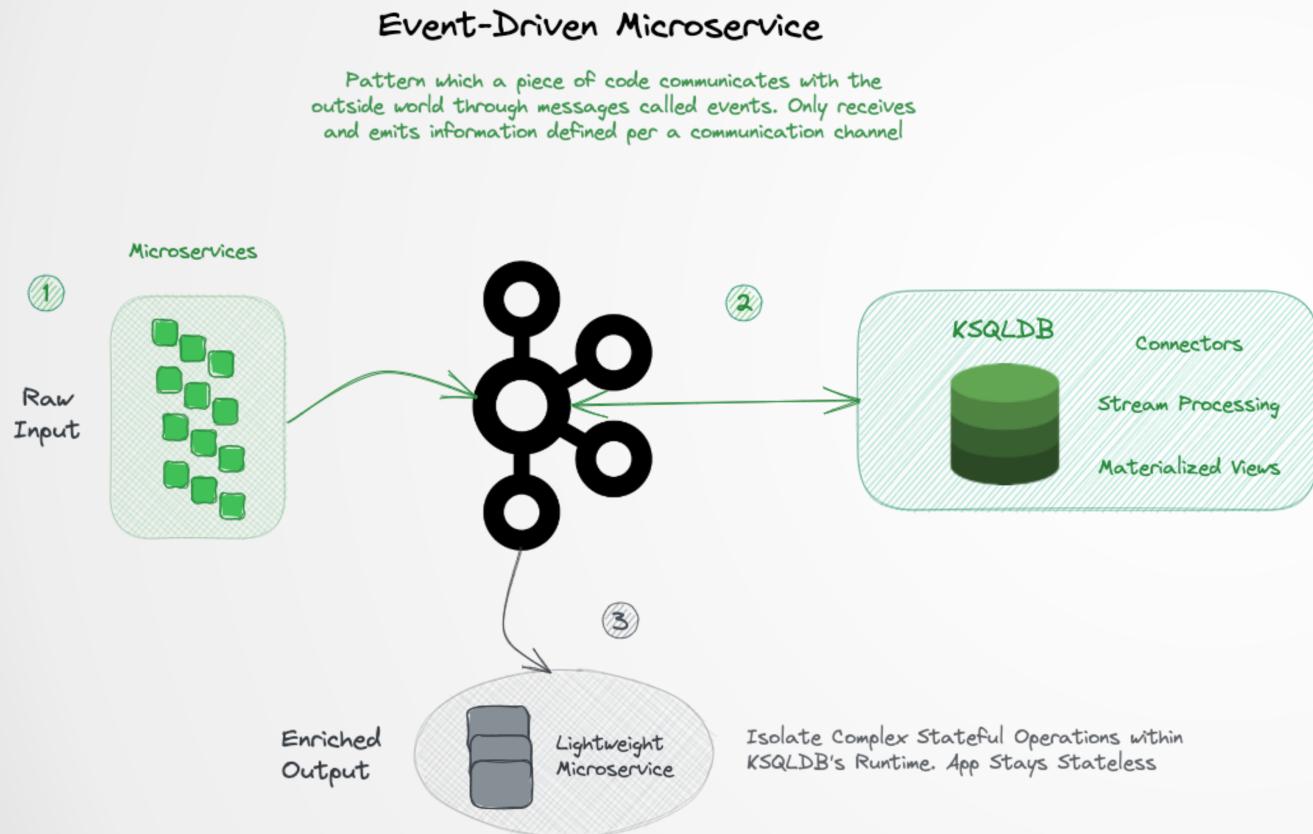
```
CREATE TABLE GRADES  
(  
    ID INT PRIMARY KEY, GRADE STRING, RANK INT  
)  
WITH  
(  
    kafka_topic = 'test_topic',  
    value_format = 'JSON',  
    partitions = 4  
)
```

```
SELECT * FROM TOP_TEN_RANKS  
WHERE ID = 5;
```



Use-Cases

Different Approaches in how to leverage the Best of KSQLDB
for Microservices, Big Data and Analytics



Features

Kubernetes Native Data Platform for RDBMS & NoSQL Databases in a Single Deployment Unit



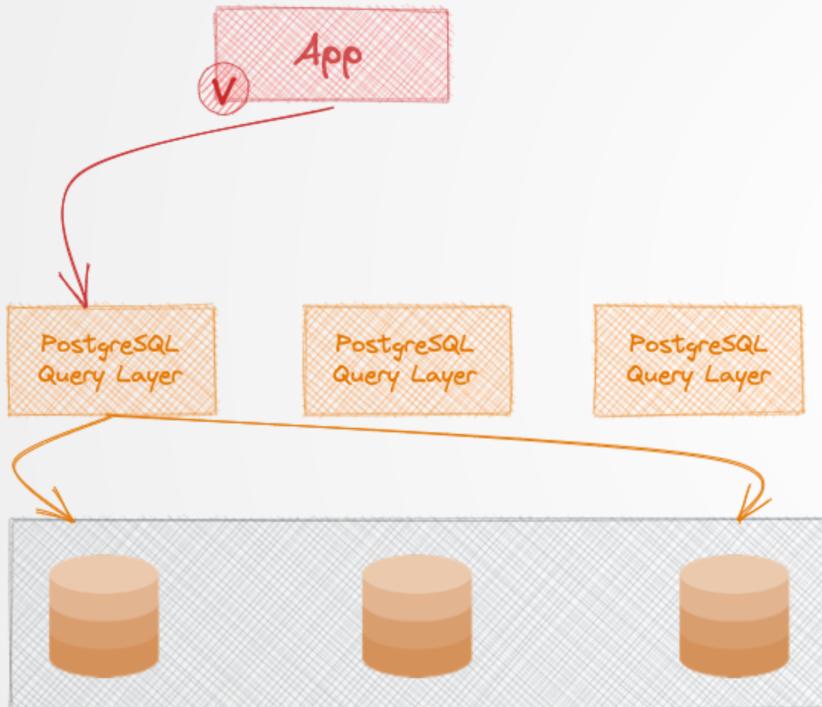
Investment

- Series C Funding of \$188 Million, Company Value of \$1.3B
- 4K+ Slack Users
- 5.5K+ GitHub Stars
- 1M+ Clusters Deployed
- 220+ Contributors



Relational Database for Cloud Era

Reimagine the Relational Database for the Cloud. Uniquely Combines Enterprise-Grade Relational Database with Horizontal Scalability



Characteristics

PostgresSQL Compatibility

- Wire Compatible
- Query Layer Reuse
- Lift & Shift Migrations

Advanced Relational Database Capabilities

- Triggers, Functions, SPS, Secondary Indexes
- Distributed ACID Transactions

Resilience and High Availability

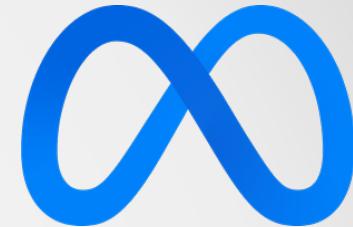
- Node, Zone, Region and Data Center Failures
- Fast Failover with Re-Replicating Data Automatically
- Zero Downtime during Upgrades or Patching

Horizontal Scalability

- Scale-Out without Disruption or Downtime
- Adding New Nodes into Cluster
- Automatically Rebalances Load

Geographic Distribution

- Flexible Deployment Options for Distribution
- Sync and Async Data Replication
- Geo-Partitioning



Layered Architecture

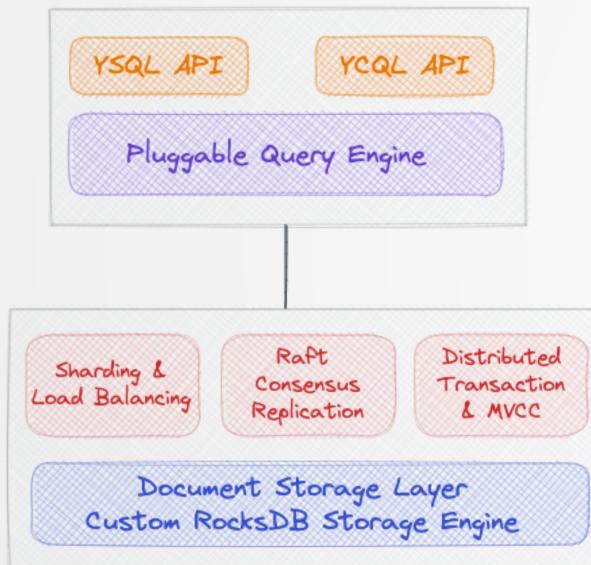
2 Logical Layers - YugaByte Query Layer and DocDB - Distributed Document Store



Characteristics

YugaByte Query Layer (YQL)

Applications interact directly using drivers
Query command issued from client



DocDB Document Store

Distributed Document Store that guarantees:

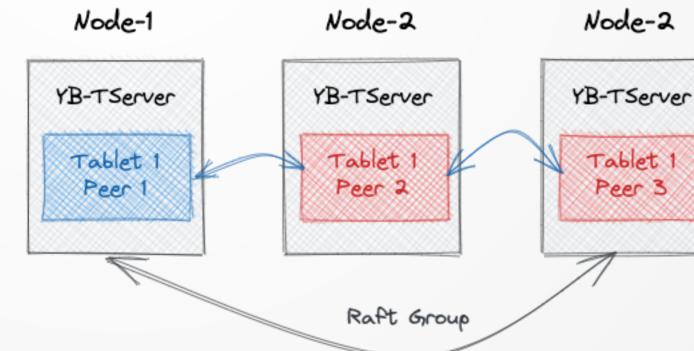
- * Strong write consistency
- * Extremely Resilient to failures
- * Automatic Sharding and Load Balancing
- * Zone, Region, and Cloud Aware Data Placement
- * Tunable Read Consistency

YugaByte SQL (YSQL) & Cloud SQL (YCQL)

- YSQL = Distributed SQL API Built to Reuse PostgreSQL Language
- YCQL = Semi-Relational Language that Roots Apache Cassandra

DocDB

- Sharding = Table Sharded into Number of Tablets
- Replication = Replication Factor using Raft Consensus Algorithm, Tablet Level
- Persistence = Log-Structured Row & Document-Oriented Storage
- Transactions = Single-Row and Multi-Row Transactions using ACID Theorem



Tablet Leader
User-Facing for Write Requests



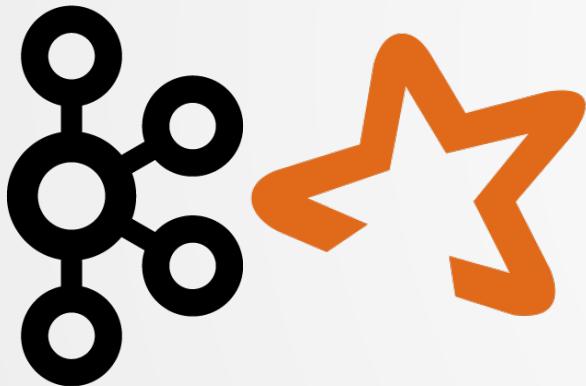
Tablet Followers
Remaining Tablet-Peers of the Raft Group

* Table Split into Tablets
* Tablets = Set of Tablet-Peers

YB-Master & YB-TServer

Distributed Architecture on Top of a Relational Database System
Intelligent System for Distributed Transactions

Integrations

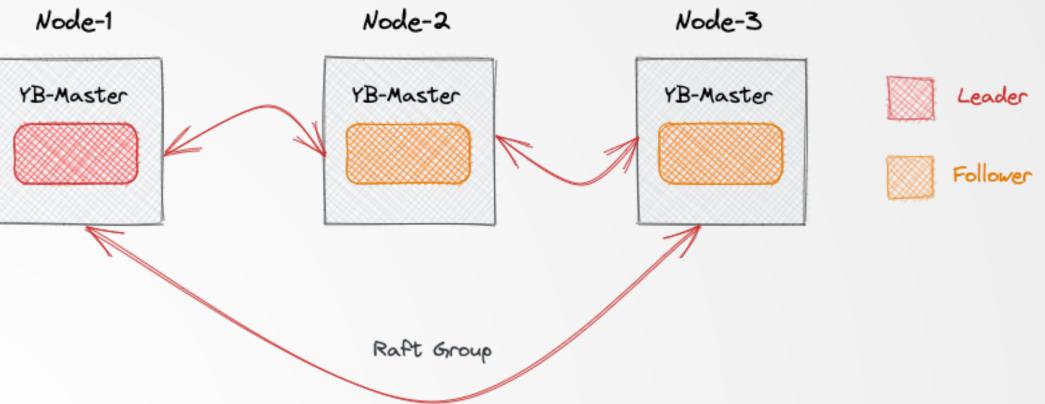


YugaByte JDBC Driver

- Distributed YSQL Built On Postgres JDBC Driver
- Cluster-Aware & Topology-Aware,

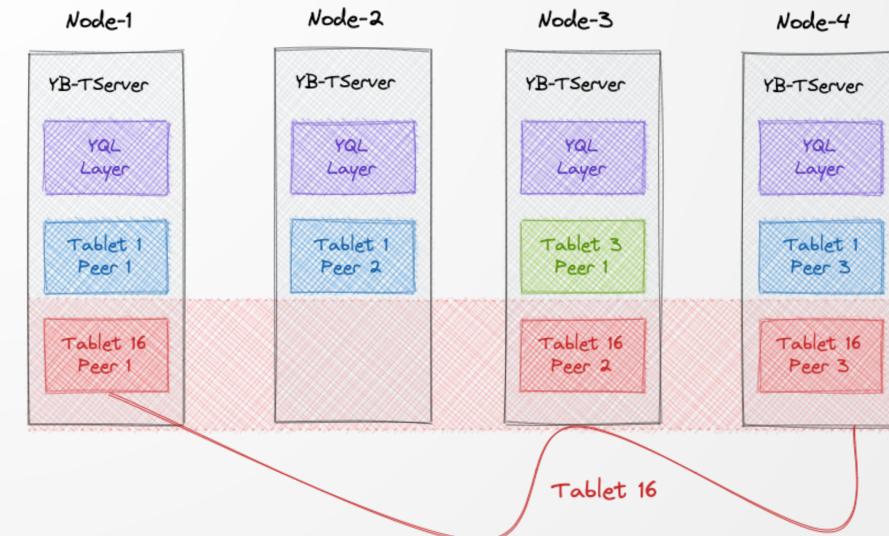
YB-Master Service

Keeper of Metadata and Records, Responsible for Coordinating Background Ops (LB, Replication) and Administrative Operations



YB-TServer Service

Tablet Server Responsible for IO Requests. Table is Split or Sharded, into Tablets. Table is composed of One or More Table-Peers, depending of the replication factor established



SQL vs. NoSQL

Best of SQL and NoSQL Databases together into one Unified Platform to simplify development of scalable cloud services

Characteristics	SQL	NoSQL	YugaByteDB
Data Model	Well-Defined Schema (Tables, Rows, Columns)	Schemaless	Both
API	SQL	Various	Fully-Relational SQL + Semi-Relational SQL
Consistency	Strong Consistency	Eventual Consistency	Strong Consistency
Transactions	ACID Transactions	No Transactions	ACID Transactions
High Write Throughput	No	Sometimes	Yes
Tunable Read Latency	No	Yes	Yes

Characteristics	SQL	NoSQL	YugaByteDB
Automatic Sharding	No	Sometimes	Yes
Linear Scalability	No	Yes	Yes
Fault Tolerance	No - Manual Setup	Yes - Smart Client Detects Failed Nodes	Yes - Smart Client Detects Failed Nodes
Data Resilience	No	Yes - Rebuilds Cause High Latencies	Yes - Automatic, Efficient Data Rebuilds
Geo-Distributed	No - Manual Setup	Sometimes	Yes
Low Latency Reads	No	Yes	Yes
Read Replica Support	No - Manual Setup	No - No Async Replication	Yes - Sync and Async Replication Options

Overview & Features

Realtime Distributed OLAP DataStore, Designed to Answer OLAP Queries with Low Latency



Characteristics

- Blazing Fast OLAP Engine
- LinkedIn, Uber, Target, Slack, Stripe
- Designed for Real-Time Analytics



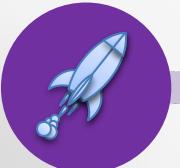
History

- Pinot Noir ~ Name of Grape for Wine
- Developed Internally at LinkedIn in 2014
- Open-Sourced in 2015
- Entered in Apache Incubator in 2018
- Top-Level Project in 2021



Key Capabilities

- Columnar-Oriented Storage
- Pluggable Indexing Technologies
- Horizontally Scale & Fault-Tolerant
- Performs Anomaly Detection using ThirdEye
- Joins using Presto & Trino



Apache Pinot

Realtime Distributed OLAP DataStore, Designed for Answering OLAP Queries with Low-Latency

Performance

- 100K + Queries per Second at ms Latency
- Ingestion of Million of Events per Second
- 50+ User-Facing Analytics



Users

- PQL ~ SQL using Apache Calcite
- API ~ RestAPIs for Broker & Controller
- External Clients ~ JDBC, Java, Python & Go

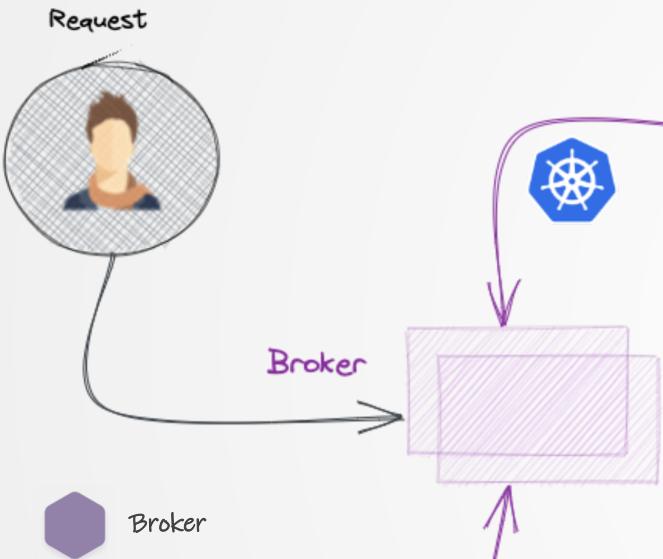


Core Components

- Controller (Helix & Zookeeper) ~ Manage State & Health
- Broker ~ Route Queries
- Server ~ Host Segments (Data) ~ Realtime & Offline Tables
- Minion ~ Purge Data

Components

Major Components and Logical Abstractions used in Apache Pinot Engine



Broker

- Handles Pinot Queries
- Accept Queries from Clients
- Consolidate Result Set



Real-Time Server

- Directly Ingests from Real-Time Systems
- Kafka, Pulsar, Kinesis
- Segments In-Memory Based on Thresholds
- Persisted on Segment Store

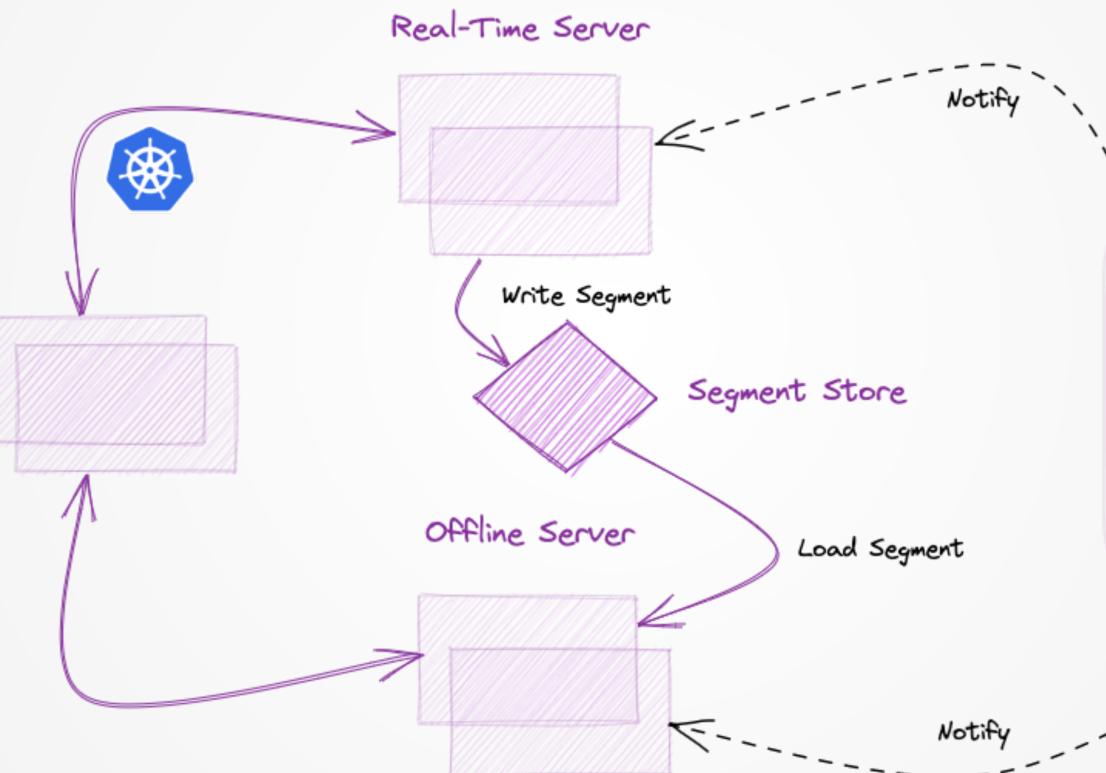
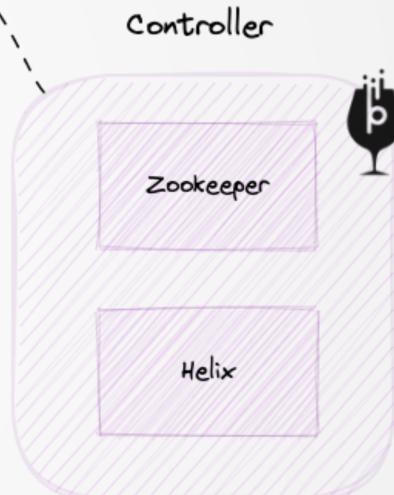


Controller

- Global Metadata
- Manage Components
- Mapping Segments and Servers
- Admin Endpoint

Apache Helix & Zookeeper

- Replicates, Partition Resources
- Uses Apache Zookeeper to Store State and Metadata

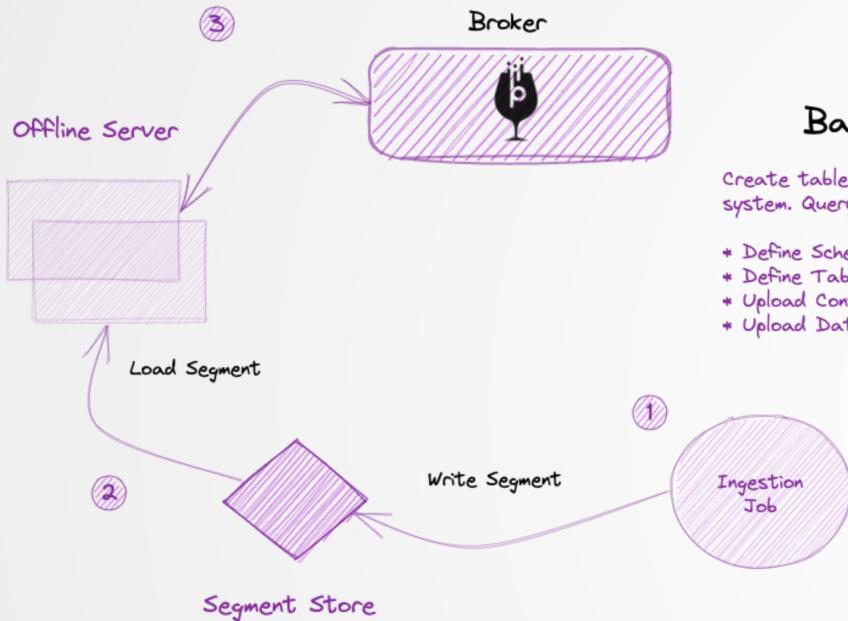


Offline Server

- Downloading Segments from Store
- Host and Server Queries Off

Ingestion Modes

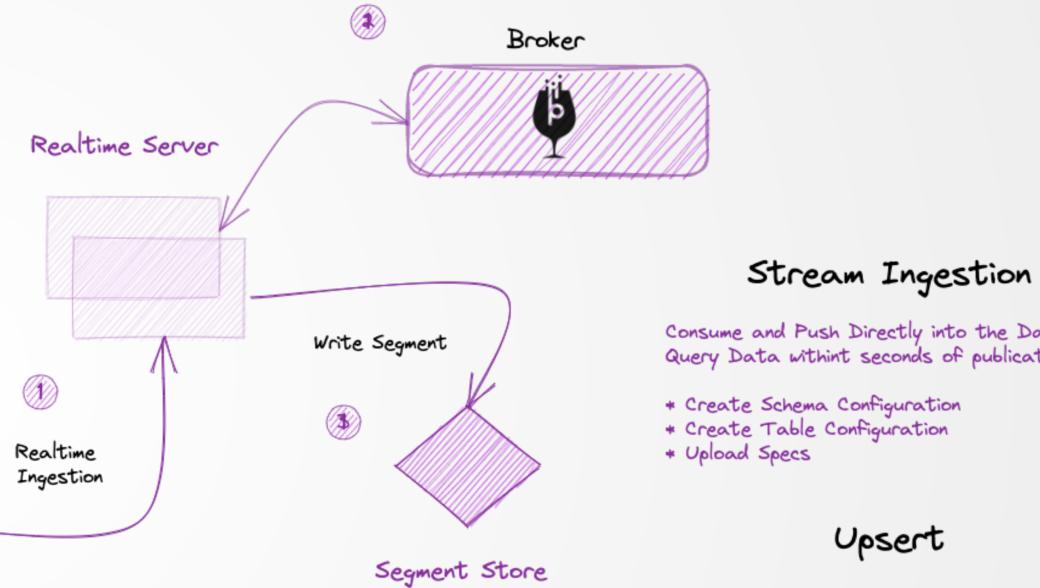
Batch and Stream Ingestion Processes of Apache Pinot Real-Time OLAP Engine



Batch Ingestion

Create tables using data present in a file system. Query large data with minimal latency

- * Define Schema
- * Define Table Config
- * Upload Configs
- * Upload Data



Stream Ingestion

Consume and Push Directly into the Database
Query Data within seconds of publication

- * Create Schema Configuration
- * Create Table Configuration
- * Upload Specs

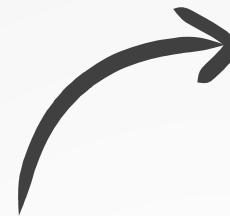
Upsert

Native Support of Upsert during Real-Time Ingestion

- * Define Primary Key
- * Enable Upsert Config

Data Virtualization

Access Data in Any Source, Shape or Structure without Copying from Source



Processing Query Engine

Connects, Discover, Plan and Analyze Different Data Stores for Query at Scale



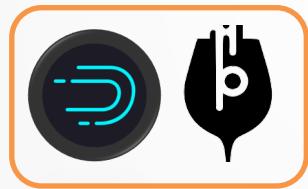
Modern Data Warehouses

Redshift, BigQuery



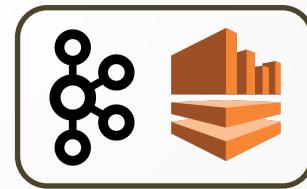
Real-Time OLAP

Apache Druid & Apache Pinot



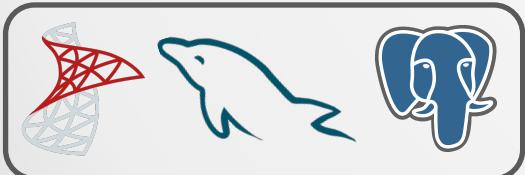
Real-Time Ingestion

Apache Kafka, Kinesis



Relational Databases

SQL Server, MySQL, Postgres



NoSQL

Cassandra, MongoDB, Redis, ElasticSearch



Data Lake

HDFS, S3, Blob Storage, GCS & MinIO [S3*]



Overview & Features

Fastest Bunny on Forest, Eating
Any Source of Data



Trino

- Speed & Scale
- Simplicity
- Versatile
- In-Place Analysis
- Query Federation
- Runs Everywhere
- Trusted
- Trino Software Foundation [OSS]



Designed For

- Access Data using SQL
- Data Warehousing
- Analytics
- OLAP



Features

- Memory Management & Spilling using Compression and Encryption
- Resource Groups
- Distributed Sort
- Dynamic Filtering ~ Query Performance, Network Traffic Reduction & Remote Load Reduction
- Cost Based Optimization - CBO
- Query Pushdown

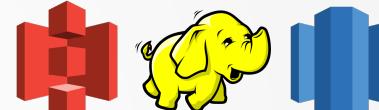


Trino is a Tool Designed to Efficiently Query Vast Amounts of Data using Distributed Queries in a Terabytes & Petabytes Range



Query Execution Model

Turn SQL Statements into Queries, Executes Across a Distributed Cluster of Coordinator and Workers



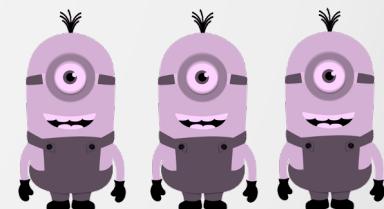
Coordinator a.k.a The Brain

Parsing Statements, Planning Queries, & Managing Trino Worker Nodes



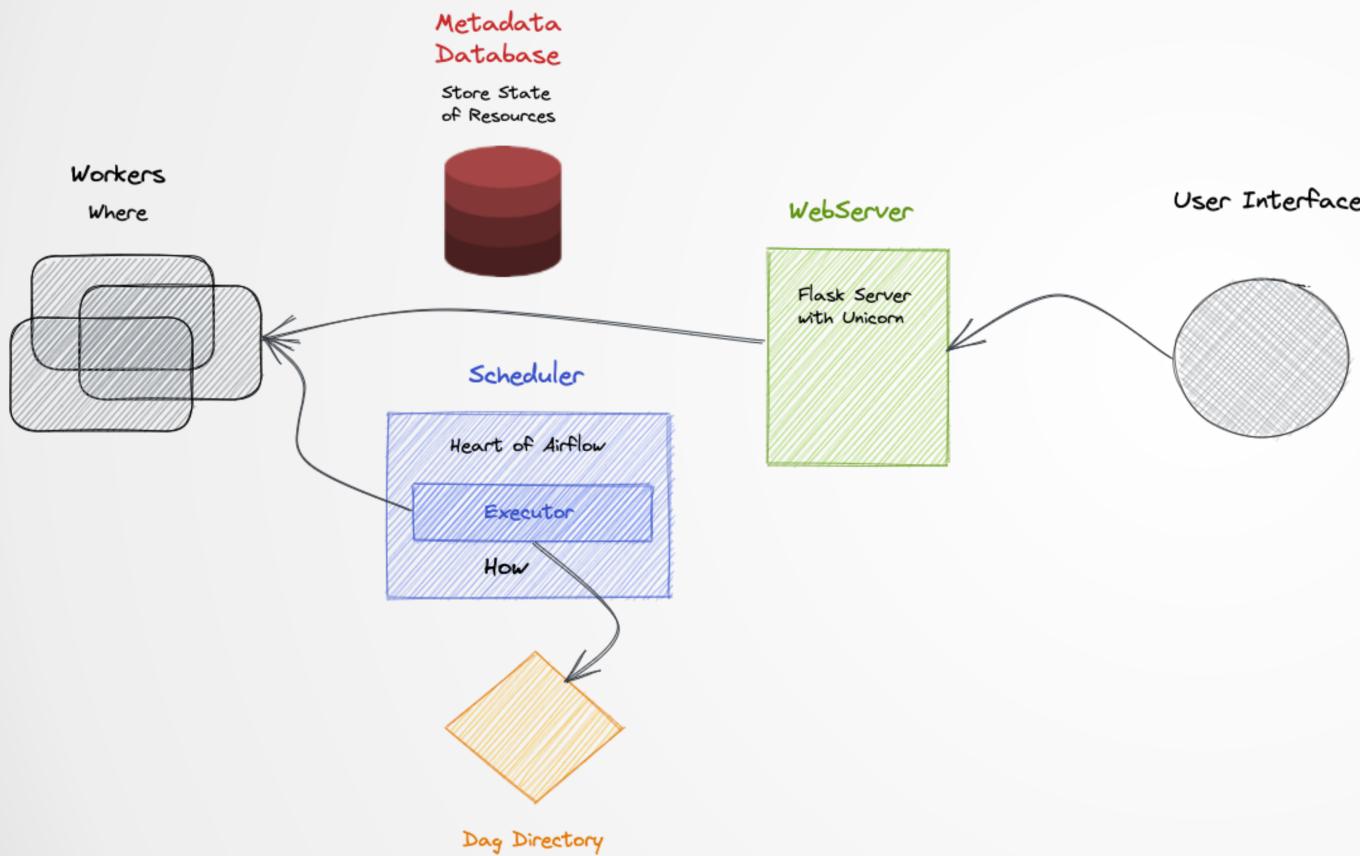
Workers a.k.a The Slave

Responsible for Executing Tasks & Processing Data. Fetch Data from Connectors and Exchange Intermediate Data in Parallel

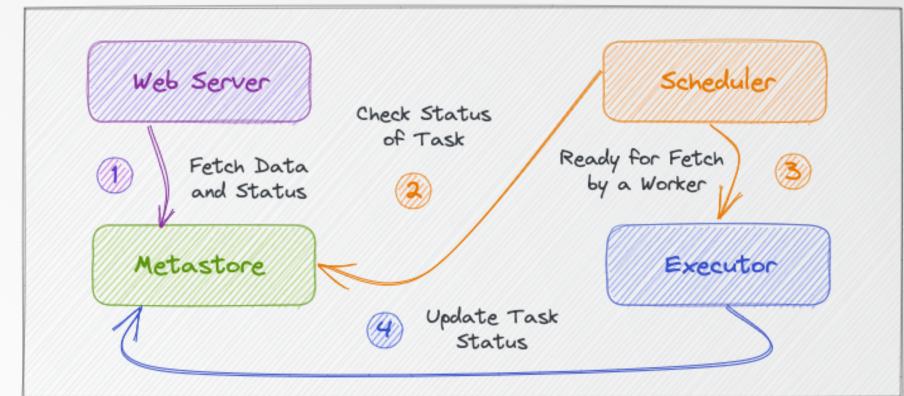


Core Components & Architecture

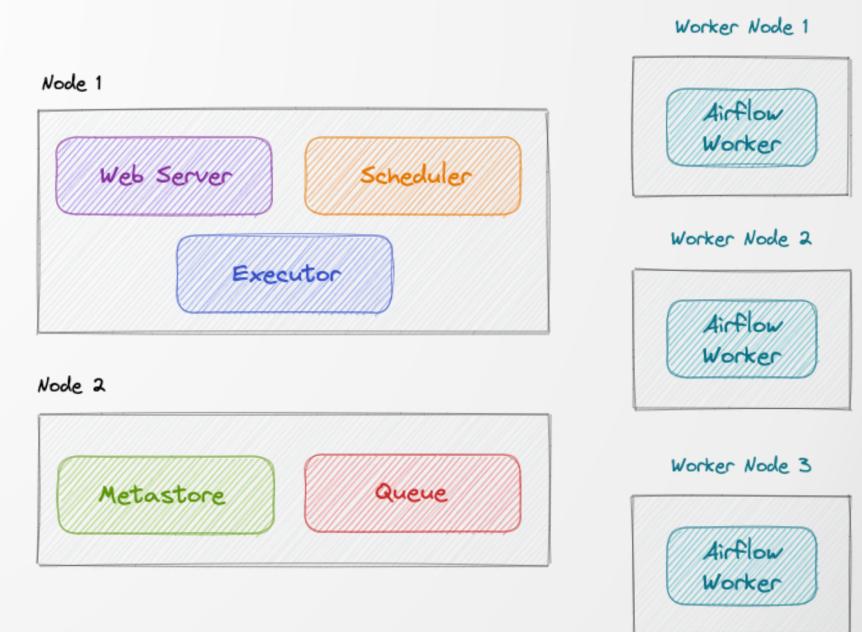
Architecture that Consists of Several Scalable Components to Execute DAGS



Single Node Architecture



Multi Nodes using Celery



Executor Types & Astronomer Registry

Different Flavors in how to Provision and Work with Apache Airflow in Dev and Production Environments



Local Executor

- Execute Tasks in Separate Processes
- Single Machine
- Non-Distributed Execution
- For Small Deployments and Dev Environment



Celery Executor

- Most Popular
- Uses Celery Queue System
- Deploy Multiple Workers
- Reads from Queue System (Redis or RabbitMQ)
- Compute Heavy Operations
- Used for Production Environments



Kubernetes Executor

- Widely Used for Production-Ready Executor
- Spawns a New Pod with Instance
- Can Cause Additional Overhead for Short Running Tasks



Celery & Kubernetes Executor

- Uses Celery & Kubernetes Combined
- Use Special Kubernetes Queue
- Specify Particular Tasks
- Advantage of Horizontal Auto Worker Scaling
- Delegate Long Running and Heavy Tasks to Executor

Astronomer Registry

Building Blocks for Your Apache Airflow Data Pipelines

- Providers
- Modules
- DAGs



Airflow on Kubernetes

Official Deployment for Airflow on Kubernetes use of Helm
Chart for Easy Development and Deployment with
Kubernetes Native-Features



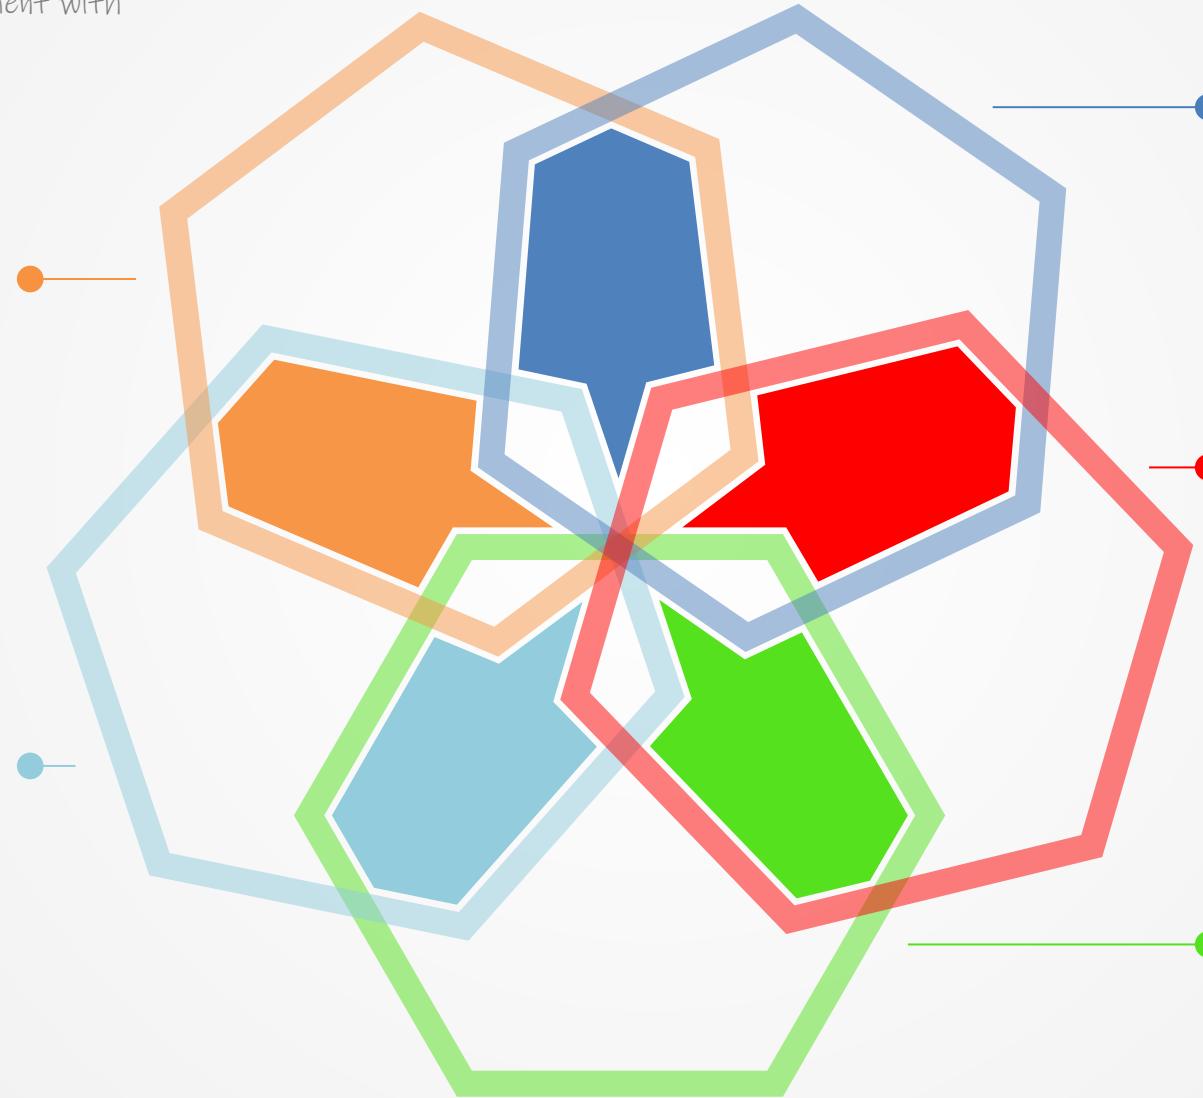
Auto Scaling

Using Celery Executor with Auto
Scaler of KEDA Embedded



Build Images

Build Images to Add Extra
Packages and Resources



Bullet Proof

Fully Integrated with Kubernetes Engine
Helm Chart Designed by Astronomer Team



Backend Databases

- Postgres (Recommended)
- MySQL
- SQLAlchemy

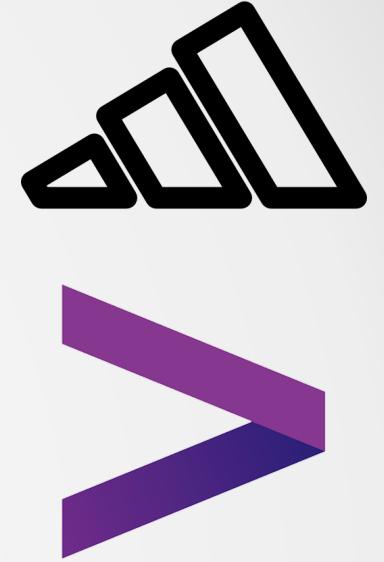
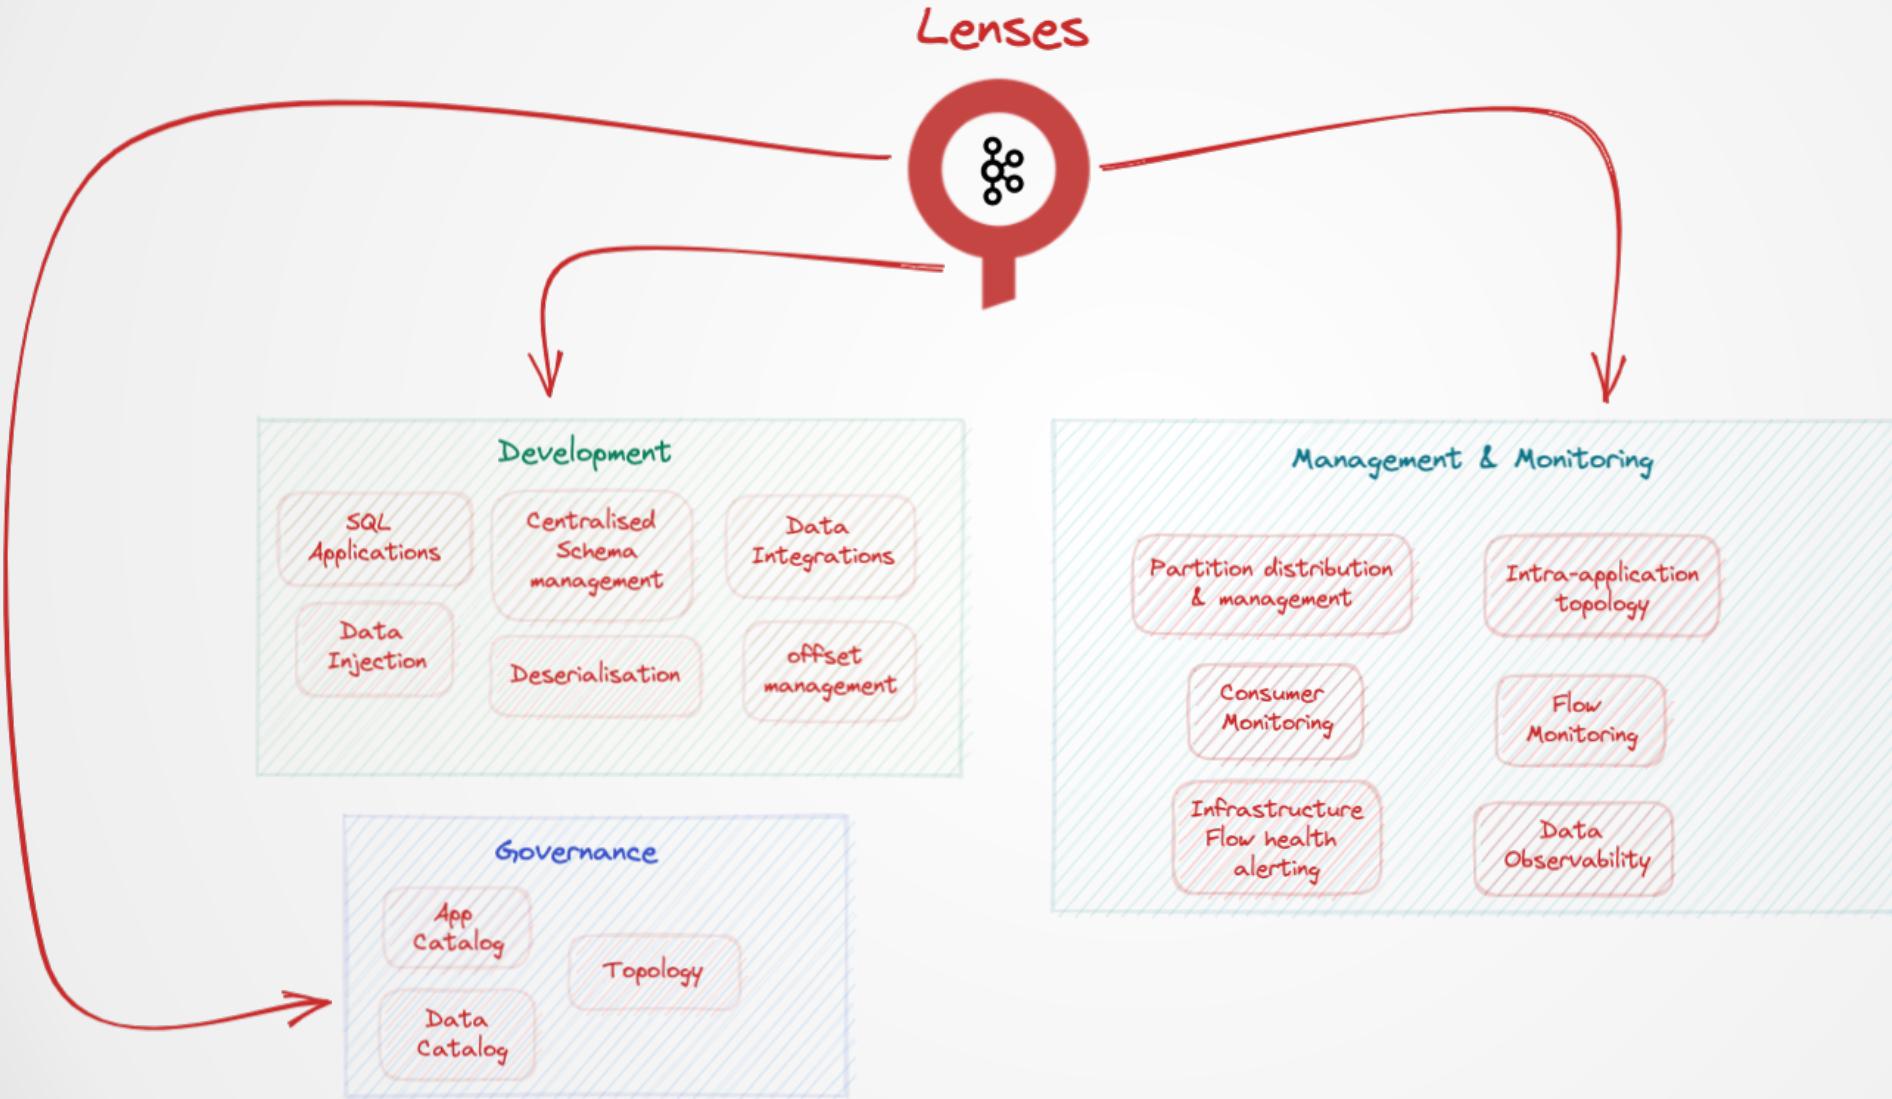


Supported Executors

- Local Executor
- Celery Executor
- Celery Kubernetes Executor
- Kubernetes Executor

Overview & Features

Real-Time Data without PFD in Apache Kafka



Lenses on Kubernetes

Enabling Lenses for Monitoring and Watchdog Apache Kafka Installations on Kubernetes

