



ONEWAY
SOLUTION



One Way Solution

The Sinks & Serving

Data Engineering – [Day 4]



JUAN MORENO

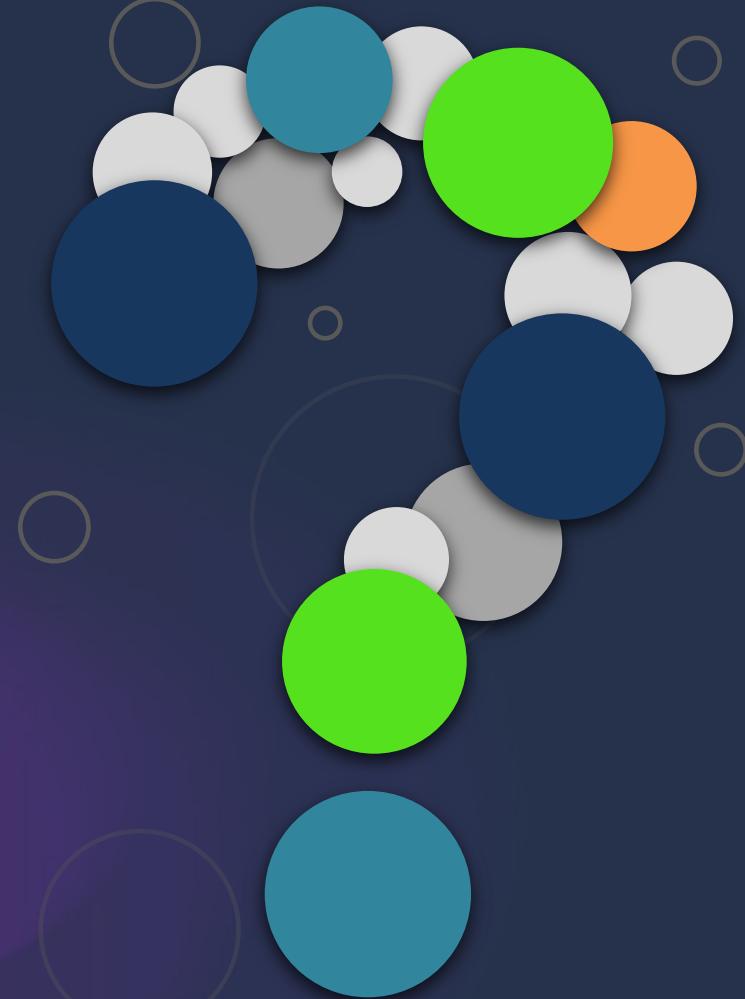
CEO & Data Architect
Data Engineer & MVP



MATEUS OLIVEIRA

Big Data Architect
Data In-Motion Specialist

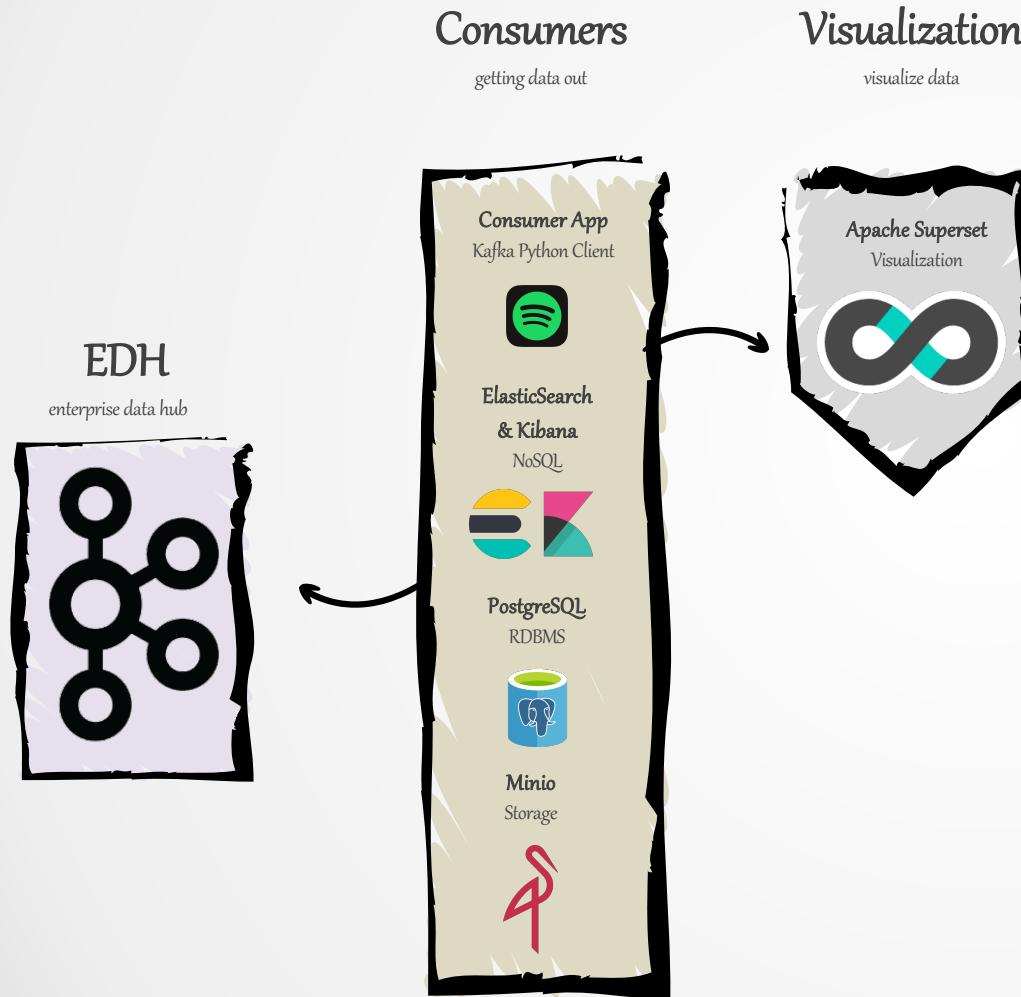




Shallow men believe in luck or
in circumstance. Strong men
believe in cause and effect.

Ralph Waldo Emerson

Use Case ~ [Data Sinks & Serving]



A dark, moody photograph of a modern skyscraper with a curved, ribbed facade, set against a hazy sky.

**With enough courage, you
can do without a reputation.**

Margaret Mitchell

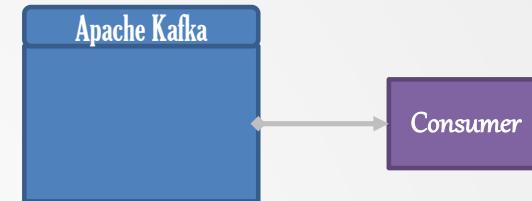
Kafka Consumer [Basics]



Consumer API [101]

subscribe to one or more topics and process the stream of records produced to them in an application

- **Consumers** = Subscribers or Readers in a Traditional Messaging Systems
- Reads Messages [Events] from a Topic(s)



Programming Languages [Kafka Clients]

Librdkafka is a C library of the Apache Kafka Protocol, providing Producer, Consumer & Admin Clients. Designed with Message Delivery Reliability and High Performance in mind.

Producers = 1 Million of Messages per Second
Consumers = 3 Million of Messages per Second

Confluent

- C, C++, Java, .Net, Python & Scala



Open Source Community

- Erlang, Groovy, Haskell, Kotlin, Lua, NodeJS, OCaml, PHP, Ruby, Rust, TCL, Swift

Kafka Consumer & Consumer Groups



Consumer Groups

suppose you have an application that needs to read messages from a kafka topic, run some validations against them, and write the results to another data store. in this case your application will create a consumer object, subscribe to the appropriate topic, and start receiving messages, validating them and writing the results. this may work well for a while, but what if the rate at which producers write messages to the topic exceeds the rate at which your application can validate them? if you are limited to a single consumer reading and processing the data, your application may fall farther and farther behind, unable to keep up with the rate of incoming messages. obviously, there is a need to scale consumption from topics. Just like multiple producers can write to the same topic, we need to allow multiple consumers to read from the same topic, splitting the data between them.



Info

Consumer = Reader
Consumer Group = Group of Readers

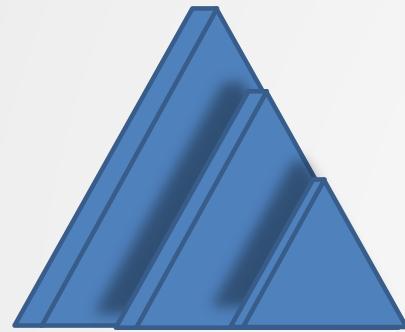
1 2 3 4 5 6



Consumer Group



Kafka Consumer Groups [Case 1]



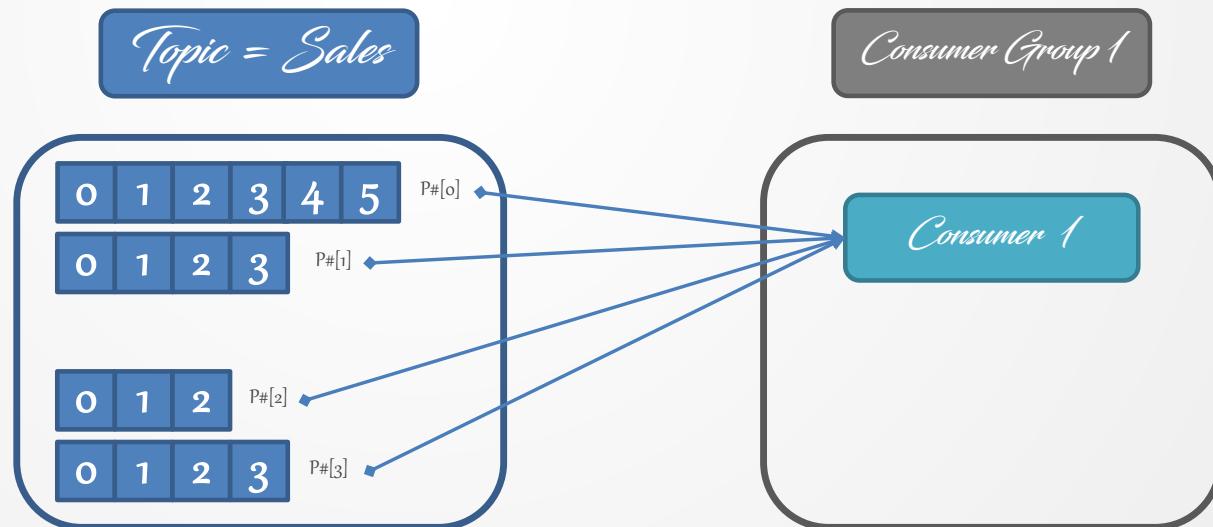
One Consumer Group with 4 Partitions

let's take topic T_1 with four partitions. now suppose we created a new consumer, C_1 , which is the only consumer in group G_1 , and use it to subscribe to topic T_1 . consumer C_1 will get all messages from all four T_1 partitions

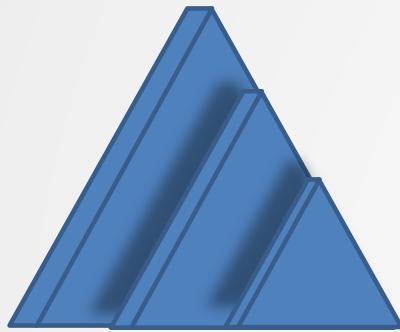


Info

unable to keep up with the rate of incoming messages,
always a need to scale consumption from topics



Kafka Consumer Groups [Case 2]

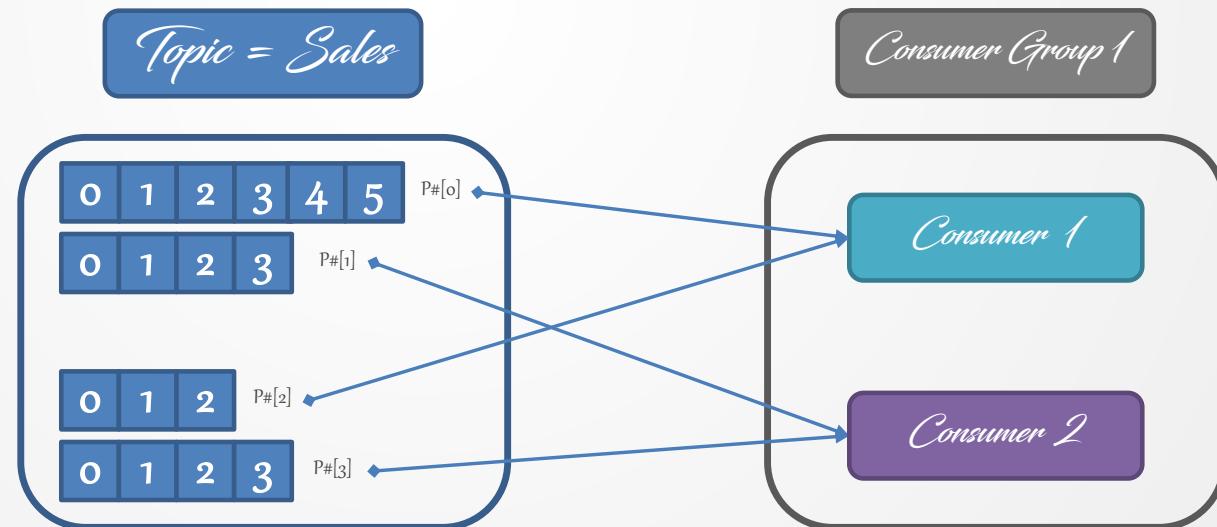


Info

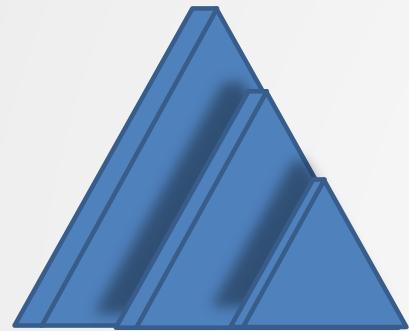
read messages faster since there is now two consumers requesting data from the same topic in different partitions

Four Partitions Split to Two Consumer Groups

if we add another consumer, C₂, to group G₁, each consumer will only get messages from two partitions.
perhaps messages from partition 0 and 2 go to C₁ and messages from partitions 1 and 3 go to consumer C₂



Kafka Consumer Groups [Case 3]

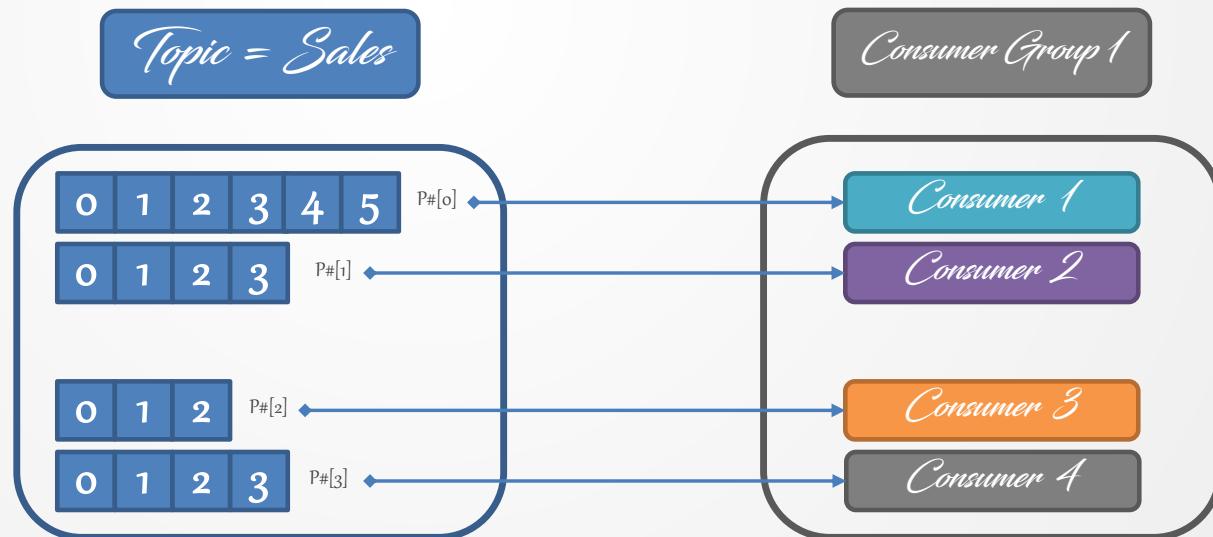


Four Consumer Groups to One Partition Each

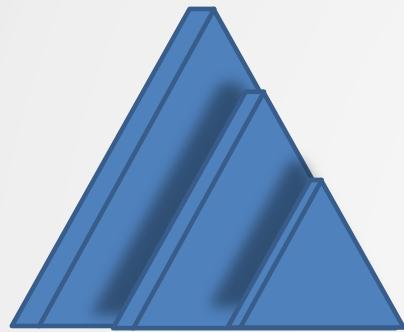
if G1 has four consumers, then each will read messages from a single partition



every consumer now can pull data from a partition making the read lot faster



Kafka Consumer Groups [Case 4]



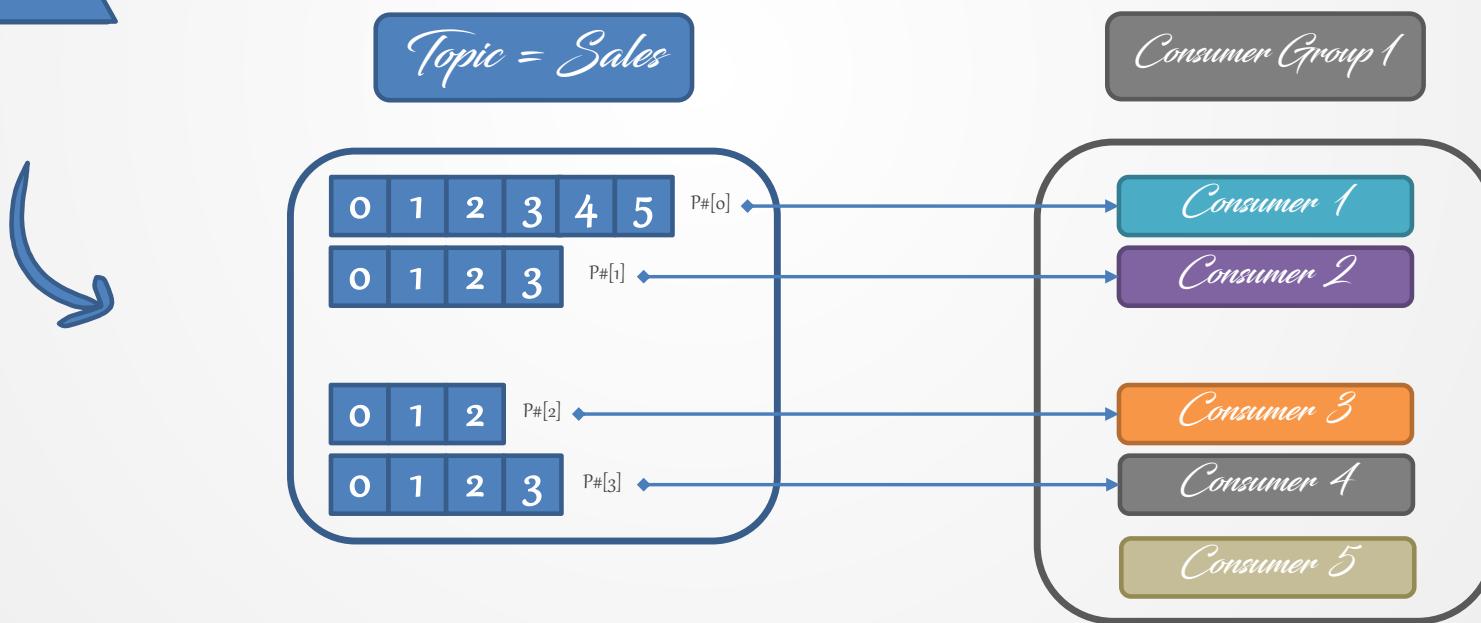
More Consumers Groups than Partitions

if we add more consumers to a single group with a single topic than we have partitions, some of the consumers will be idle and get no messages at all

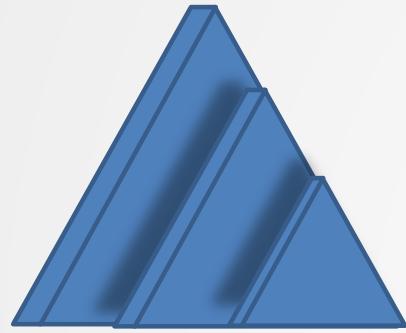


Info

there is no point in adding more consumers than you have partitions in a topic, some consumers will get idle



Kafka Consumer Groups [Case 5]



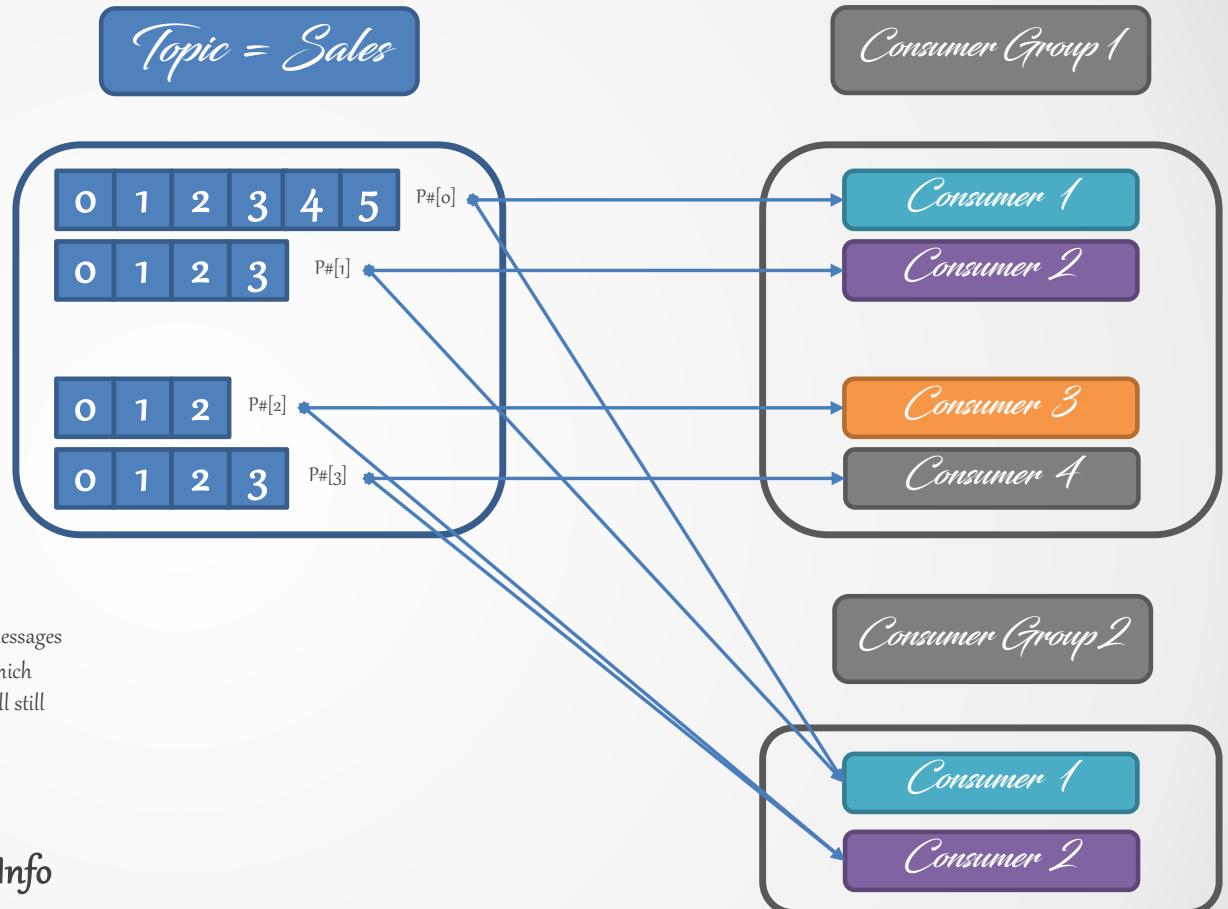
Adding New Consumer Group

If we add a new consumer group G_2 with a single consumer, this consumer will get all the messages in topic T_1 independent of what G_1 is doing. G_2 can have more than a single consumer, in which case they will each get a subset of partitions, just like we showed for G_1 , but G_2 as a whole will still get all the messages regardless of other consumer groups.



Info

New consumer groups won't impact others that are consuming from the same topic. Consumer groups are independent of each other.

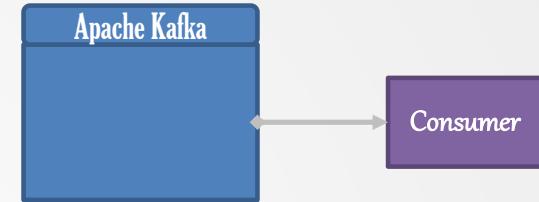


Kafka Consumer [Dissecting Read Process]



Kafka Consumer

the first step to start consuming records is to create a KafkaConsumer instance. creating a KafkaConsumer is very similar to creating a KafkaProducer - you create an instance with the properties you want to pass to the consumer. we will discuss all the properties. to start we just need to use the three mandatory properties: bootstrap.servers, key.deserializer, and value.deserializer.



Initial Configuration

start by adding the mandatory properties

- 1) Bootstrap.Servers
- 2) Key.Deserializer
- 3) Value.Deserializer

Recommendation = Group.Id



Topic Subscription

subscribe to one or more topics (use regex expression) for that

- 1) Topic Name



Poll Loop

the heart of the consumer API is a simple loop for polling data, it handles automatically

- 1) Coordination
- 2) Partition Rebalance
- 3) Heartbeats
- 4) Data Fetching

infinite loop = long-running application that polls from kafka.

poll() = needs to keep polling data otherwise will be considered dead



first time consumer calls a poll() with a new consumer, it is responsible to

- 1) Find Group Coordinator
- 2) Join Consumer Group
- 3) Receive Partition Assignment
- 4) One Consumer per Thread



Building a Kafka Consumer Application using Python



Consumer Settings

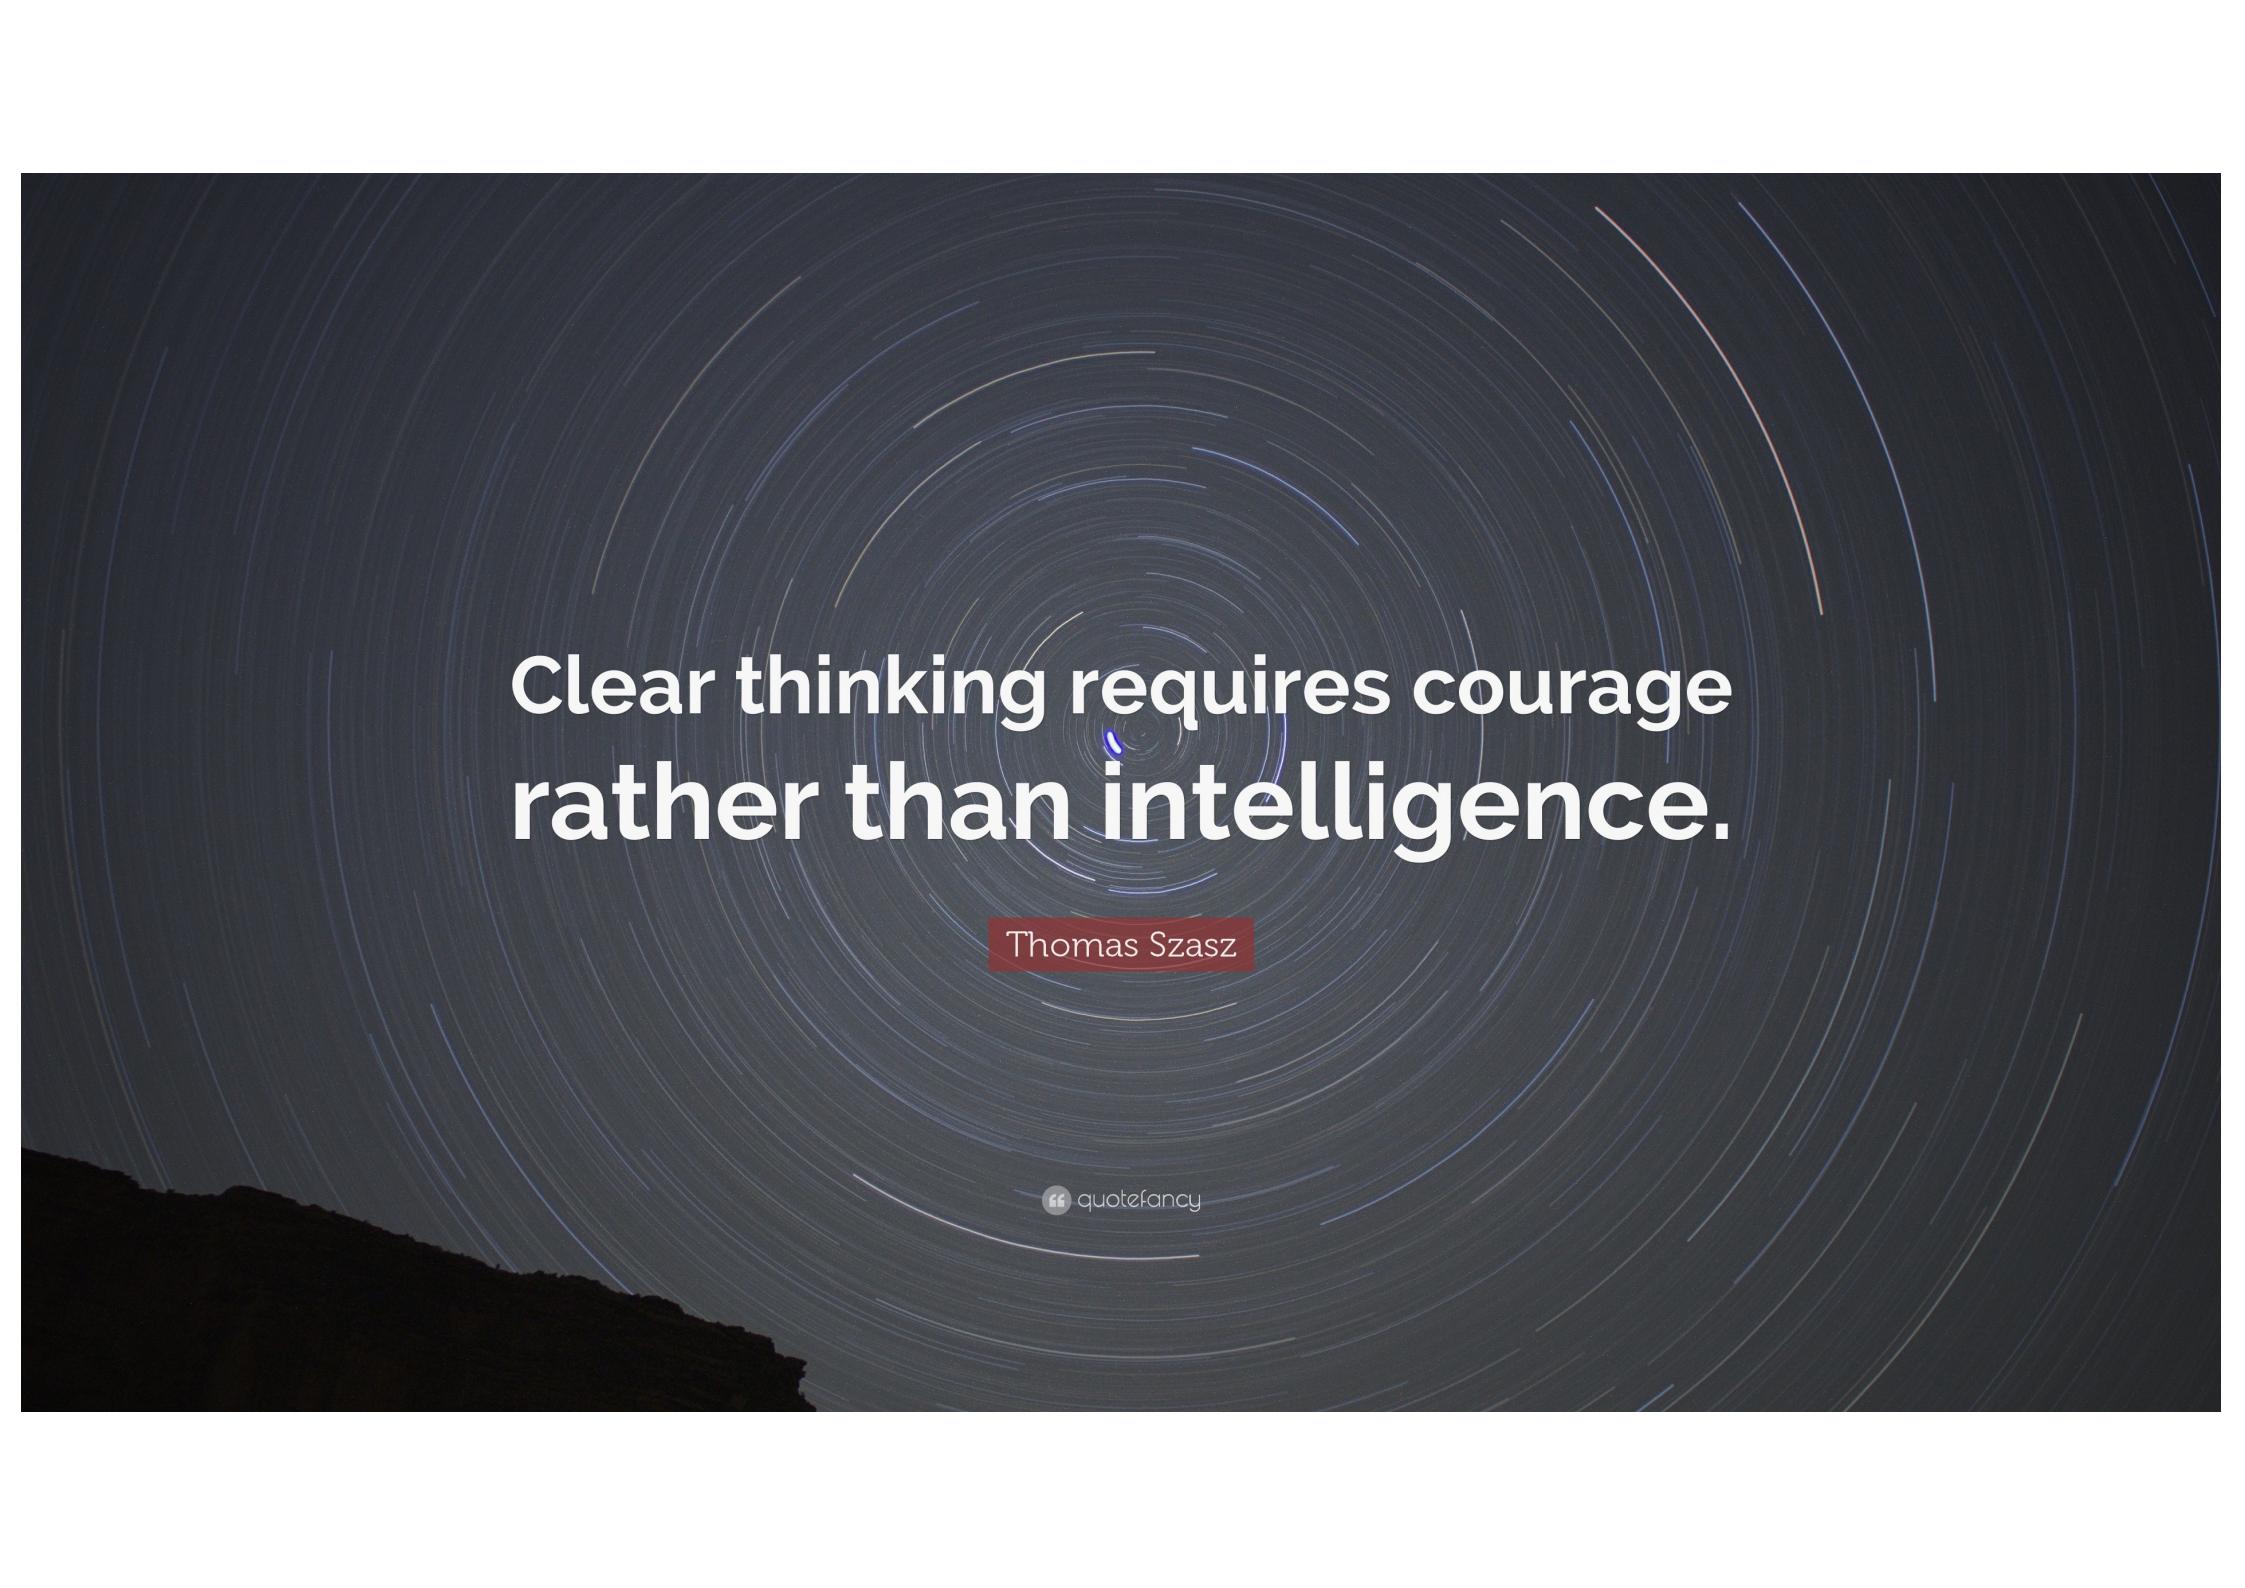


Kafka Consumer API



Consume Data





Clear thinking requires courage
rather than intelligence.

Thomas Szasz

Kafka Consumer Groups & Partition Rebalance



Rebalance

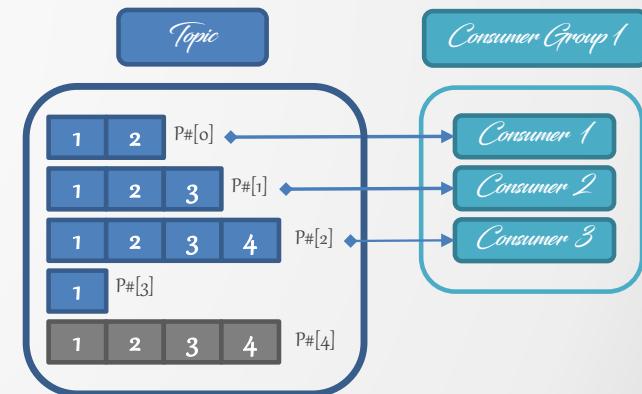
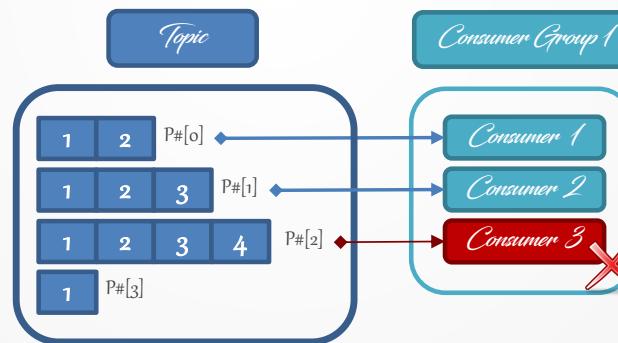
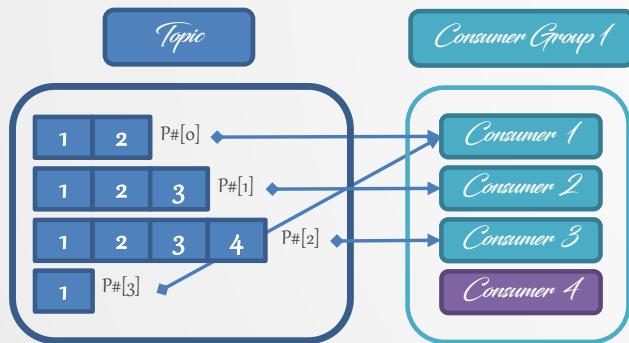
moving partition ownership from one consumer to another is called a rebalance. rebalances are important because they provide the consumer group with high availability and scalability (allowing us to easily and safely add and remove consumers), but in the normal course of events they are fairly undesirable. during a rebalance, consumers can't consume messages, so a rebalance is basically a short window of unavailability of the entire consumer group. in addition, when partitions are moved from one consumer to another, the consumer loses its current state; if it was caching any data, it will need to refresh its caches - slowing down the application until the consumer sets up its state again



Info

- Share Ownership of Partitions ~ Topics
- Rebalance is a Window of Unavailability

1. Add New Consumer
2. Consumer Shuts Down or Crashes
3. Topic is Modified = New Partitions Added



1

adding new consumer into the consumer group triggers a rebalance

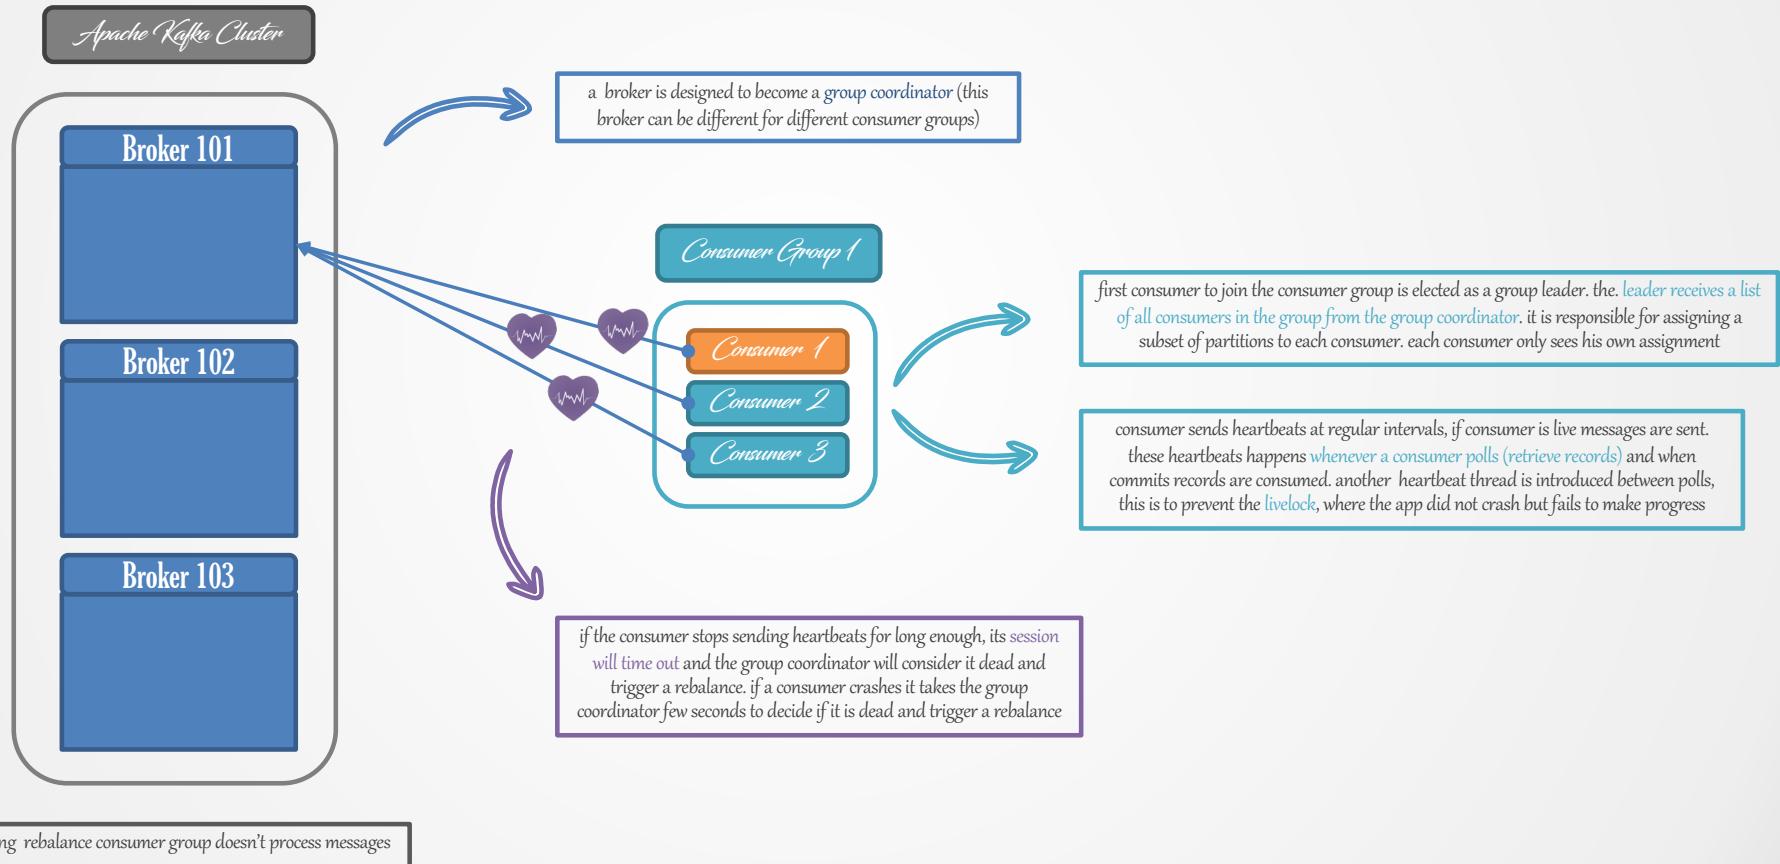
2

consumer shuts down or crashes triggering a rebalance process

3

adding a new partition into the topic triggers a rebalance

Kafka Consumer Groups & Membership



Kafka Consumer Groups [Commits & Offsets]



Commits & Offsets

whenever we call `poll()`, it returns records written to kafka that consumers in our group have not read yet. this means that we have a way of tracking which records were read by a consumer of the group. As discussed before, one of kafka's unique characteristics is that it does not track acknowledgments from consumers the way many jms queues do. instead, it allows consumers to use kafka to track their position (offset) in each partition.

1

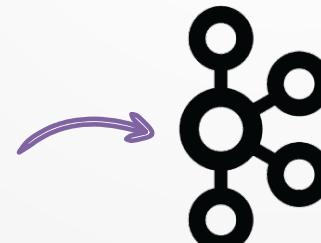
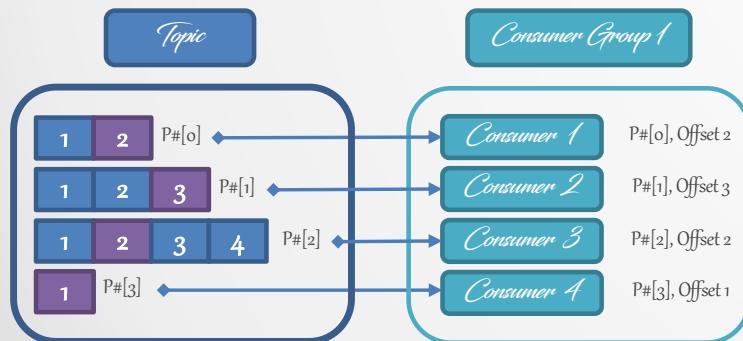
Commit

consumers use kafka to track their position, the action of updating the current position is what is call commit

2

Offset

it produces a message to kafka, to a special `_consumer_offset` topic, with the committed offset for each partition



3

Rebalance

consumer will read latest committed offset of each partition

- 1) App Crash or Unexpected Shutdown
- 2) New Consumer into Consumer Group
- 3) Topic Changes

Reprocessing

4

Reprocessing Events [Duplicates]

Processing Events Now

1 2 3 4 5 6 7 8

Last Committed Offset
(`_consumer_offset`)

Events ~ Last Poll()

2

3

1 Processing Events Now

Events ~ Last Poll()

1 2 3 4 5 6 7 8

Missed Events

Lost Events

2 Last Committed Offset
(`_consumer_offset`)

Kafka Consumer Groups [Commit Offset Options]



Modes

ways of committing offsets

Automatic Commit [Enable. Auto.Commit = True]

every five seconds the consumer will commit the largest offset your client received from poll(). the seconds can be controlled by setting Auto.Commit.Interval.ms

- does not give to developer enough control to avoid duplicate messages

Commit Current Offset [Enable. Auto.Commit = False]

used to reduce the possibility of missing messages and to reduce the number of duplicated messages during rebalancing. committed explicitly only and use commitSync(), throw an error if commit fails

- gives total control over the commit although it's a sync process (limit throughput)

Async Commit

another option is the async option where instead of waiting to respond to a commit, we just send the request and continue to process upcoming messages, problem is that commitAsync() won't retry, possibility to add a callback to log errors

- it's faster but does not allow retry, if the commit fails can leads to duplicated messages

Combine Synchronous & Asynchronous Commits

common pattern to combine the both methods before shutdown, if everything is fine the Async() is called since it's faster but if fails the Sync() is invoked and will retry until it succeeds or suffers unrecoverable failure

- speediness for normal writes and security whenever a failure or wait happens

commit current offset

```
while (true) {  
    ConsumerRecords<String, String> records = consumer.poll(100);  
    for (ConsumerRecord<String, String> record : records)  
    {  
        System.out.printf("topic = %s, partition = %s, offset = %d, customer = %s, country = %s\n",  
            record.topic(), record.partition(),  
            record.offset(), record.key(), record.value());  
    } try {  
        consumer.commitSync(); ←  
    } catch (CommitFailedException e) {  
        log.error("commit failed", e);  
    } }
```

async commit

```
while (true) {  
    ConsumerRecords<String, String> records = consumer.poll(100);  
    for (ConsumerRecord<String, String> record : records)  
    {  
        System.out.printf("topic = %s, partition = %s, offset = %d, customer = %s, country = %s\n",  
            record.topic(), record.partition(), record.offset(), record.key(), record.value());  
    } consumer.commitAsync(); ←  
}
```

sync & async commit

```
try {  
    while (true) {  
        ConsumerRecords<String, String> records = consumer.poll(100);  
        for (ConsumerRecord<String, String> record : records)  
        {  
            System.out.printf("topic = %s, partition = %s, offset = %d, customer = %s, country = %s\n",  
                record.topic(), record.partition(),  
                record.offset(), record.key(), record.value());  
        } consumer.commitAsync(); ←  
    } } catch (Exception e) { log.error("Unexpected error", e); } finally {  
    try {  
        consumer.commitSync(); ←  
    } finally {  
        consumer.close();  
    } }
```

Scaling-Out Kafka Consumer Application

1

Spin Up Python App Consumer

2

Verify Consumer Group ~ Scaling App

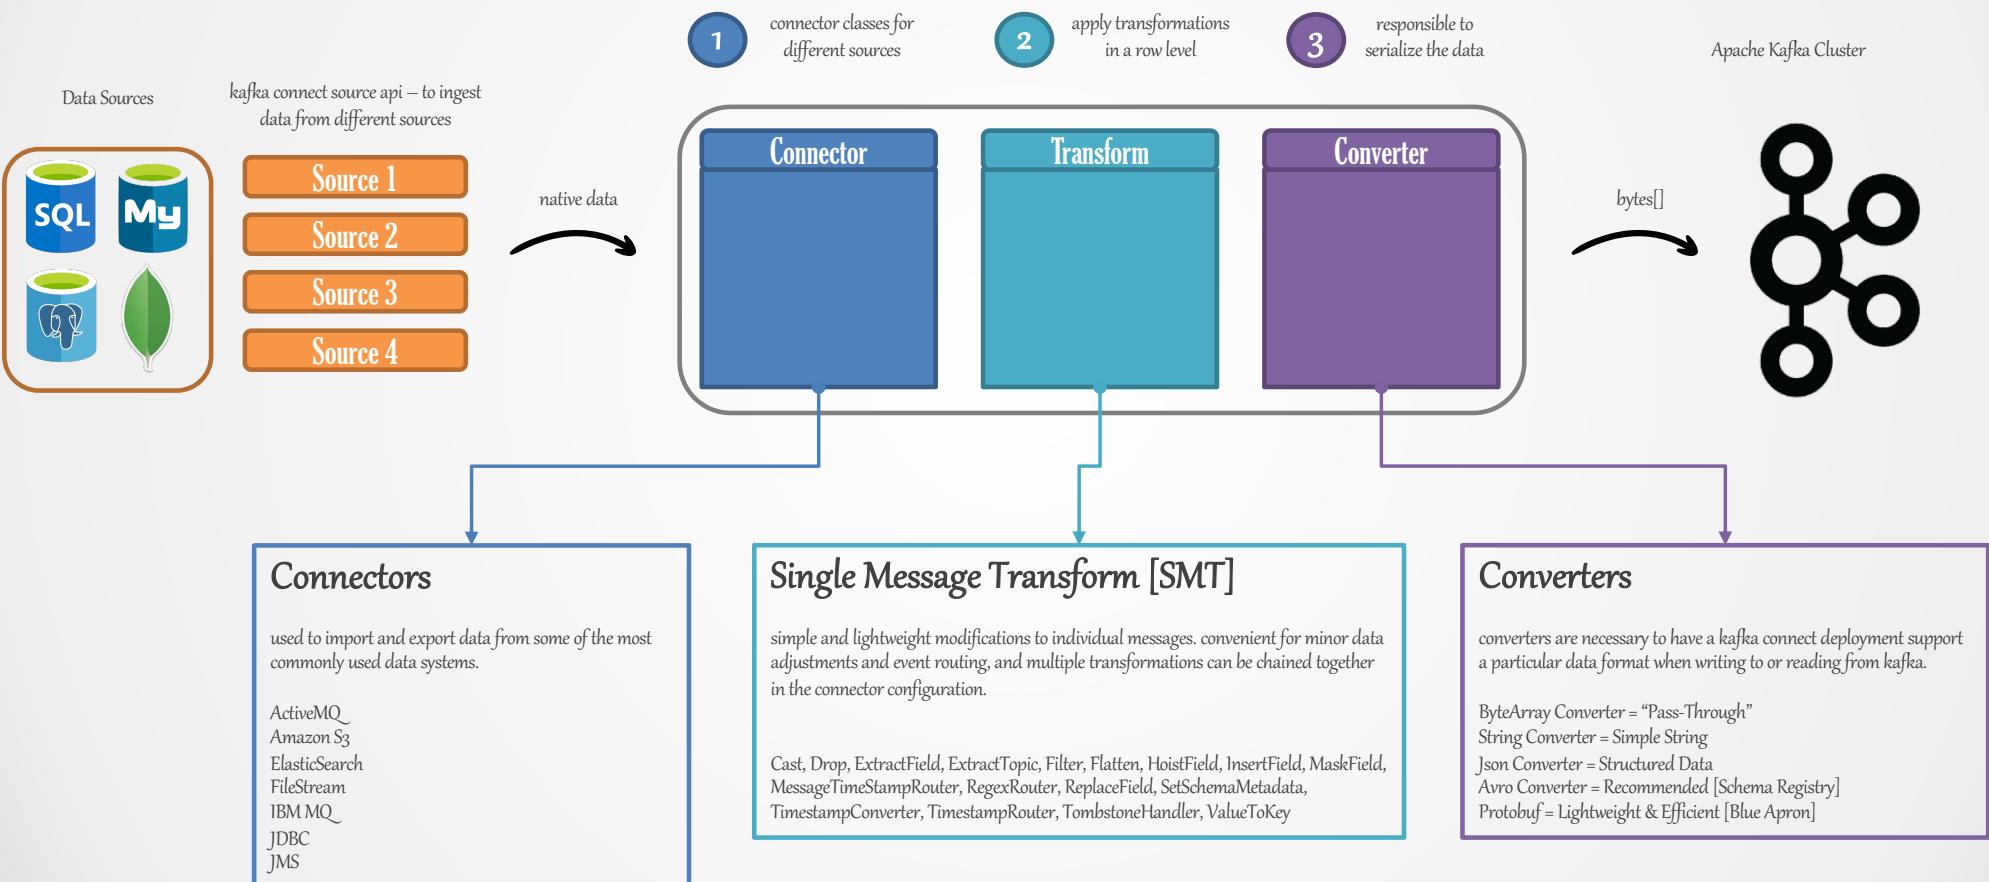


The background of the image is a wide-angle photograph of a mountain range at sunset. The sky is filled with horizontal clouds, with the upper portion in shades of orange and yellow, transitioning into a darker blue and purple towards the horizon. The mountain peaks are silhouetted against the bright sky.

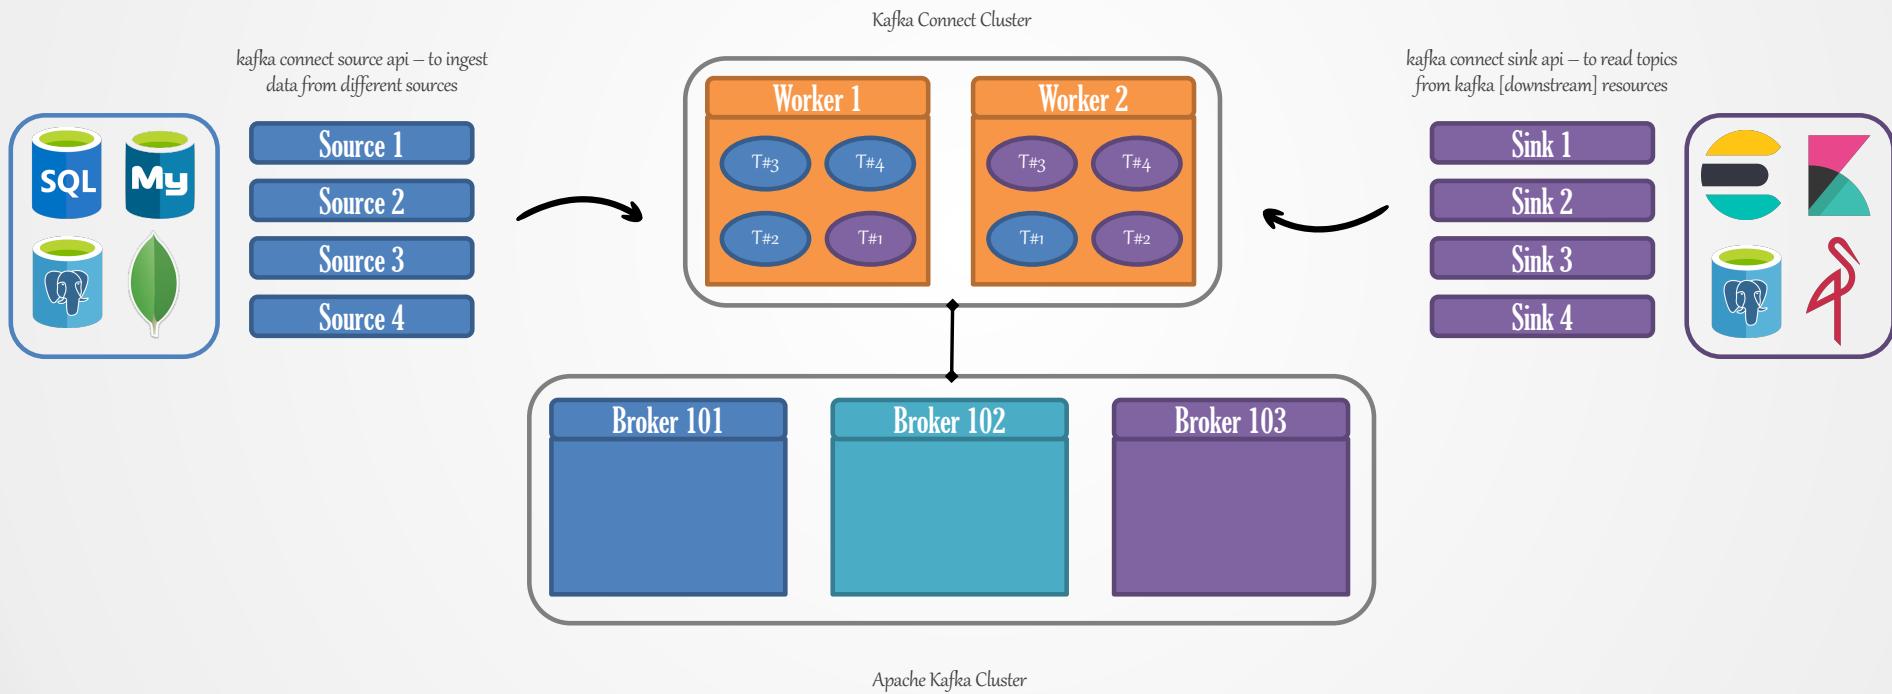
**Courage is never to let
your actions be
influenced by your fears.**

Arthur Koestler

Kafka Connect [Data Movement]

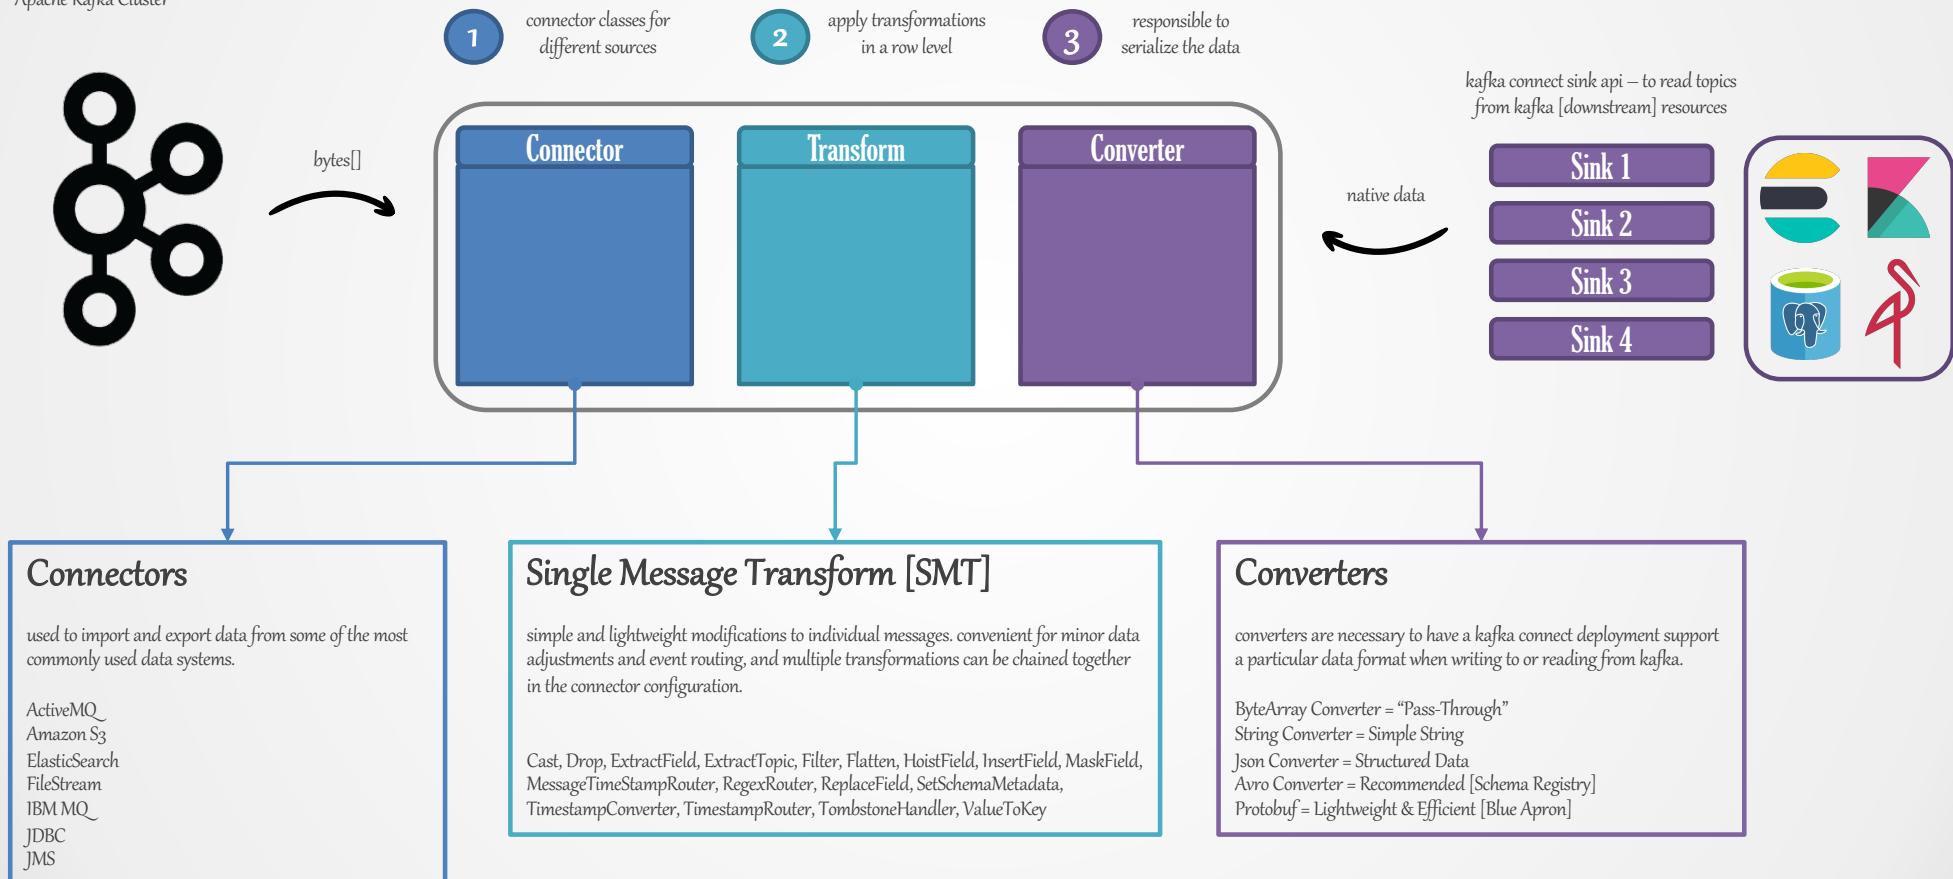


Kafka Connect [Sources & Sinks]



Kafka Connect Sink [Movement]

Apache Kafka Cluster





**It takes a lot of courage to show
your dreams to someone else.**

Erma Bombeck

Kafka Connect [ElasticSearch] ~ Sink



ElasticSearch Service Sink Connector for Confluent Platform

the service moves data from apache kafka to elasticsearch. it writes data from a topic to an index in elasticsearch. all data for a topic have the same type in elasticsearch. this allows an independent evolution of schemas for data from different topics. this simplifies the schema evolution because elasticsearch has one enforcement on mappings; that is, all fields with the same name in the same index must have the same mapping type.



ElasticSearch

- Distributed Open-Source Analytics Engine
- Built-On Apache Lucene ~ Released in 2010
- Simple REST APIs
- Powerful Search Engine for = App, Website, Logging, Business



Features

Mapping Inference

- Infer Mappings from Connect Schemas

Schema Evolution

- Backward, Forward & Fully Compatible Schema

Delivery Semantics

- Batching & Pipelined Writes ~ Boost Throughput
- Concurrent Processing of Multiple Batches

Automatic Retries

- Retries Requests Automatically



Kafka Connect [JDBC Sink Connector]



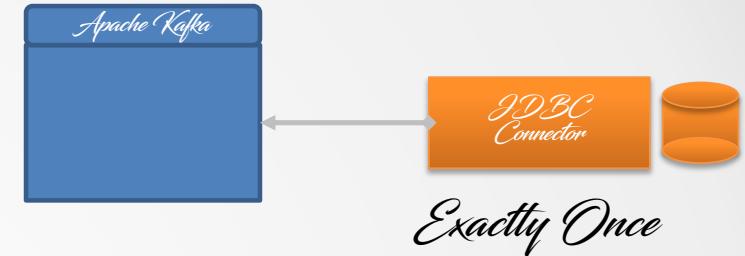
JDBC Connector [Source & Sink] for Confluent Platform

you can use the kafka connect jdbc source connector to import data from any relational database with a jdbc driver into apache kafka® topics. you can use the jdbc sink connector to export data from kafka topics to any relational database with a jdbc driver. the jdbc connector supports a wide variety of databases without requiring custom code for each one.



Relational Database Systems [RDBMS]

- Microsoft SQL Server
- PostgreSQL
- Oracle Database
- IBM DB2
- MySQL Server
- SAP HANA
- SQLite Embedded Database



Features

Data Mapping

- Apache Avro & Schema Registry
- JSON with Embedded Schema

Key Handling

- PK.Mode = None, Kafka [Topic, Partition, Offset], Record Key & Record Value

Delete Mode

- Consume Tombstone Record [Key with Null Value & PK.Mode = Record Key]

Idempotent Writes

- Insert Mode = Insert or Upsert

Auto-Creation & Auto-Evolution

- Auto.Create = Automatic Table Creation
- Auto.Evolve = Alter Table Statement

Kafka Connect [MongoDB] ~ Sink



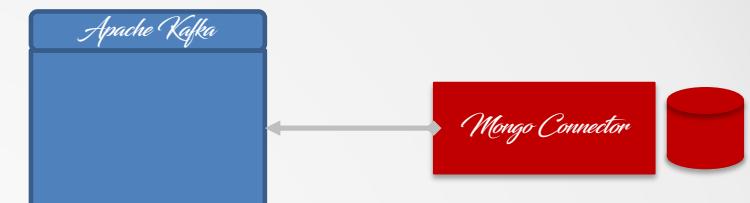
MongoDB Connector

The MongoDB Connector Officially Supported By MongoDB, Can Provide ~ Source and Sink Connectors Install One Connector and Ingestion and Output in MongoDB Database



NoSQL ~ MongoDB

- CAP Base Theorem
- Document Database [JSON]
- Fast Development Time
- High Performance
- Rich Query Language
- High Availability
- Horizontal Scalability
- Storage Engines – WiredTiger & In-Memory



Features

Unique Document ID Strategy

- document.id.strategy
- document.id.strategy.overwrite.existing
- delete.on.null.values

Schemas:

- Apache Avro & JSON Format

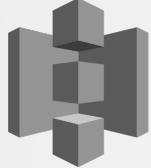
Dead Letter Queue Configuration

- Specifies Invalid Messages

Enable Mapper

- Setting namespace.mapper Using Class:
com.mongodb.kafka.connect.sink.namespaceMapper
ce.mapping.FieldPathNamespaceMapper

Kafka Connect [S3] ~ Sink



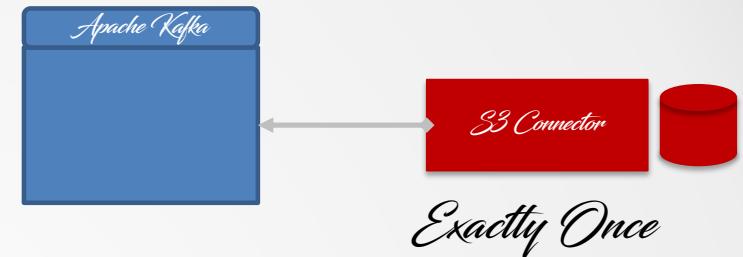
Amazon S3 Sink Connector for Confluent Platform

The S3 connector, currently available as a sink, allows you to export data from kafka topics to S3 objects in either avro, json and parquet formats. In addition, for certain data layouts, S3 connector exports data by guaranteeing exactly-once delivery semantics to consumers of the S3 objects it produces.



MinIO

- High Performance Kubernetes Native Object Storage
- Designed ~ S3 API
- 100% Open-Source
- Erasure Coding [Written in Assembly]
- Encryption
- Identity Management
- Continuous Replication
- Federation
- Multi-Cloud Gateway [S3, Blob Storage & GCS]



Features

Exactly Once Delivery

- Using a Deterministic Partitioner
- Eventual Consistency

Pluggable Data Format with or without Schema

- Apache Avro, Apache Parquet & JSON Format

Pluggable Partitioner

- Default, Custom & Time Based

Non-AWS Object Storage Support

- Minio

Partitioning Records

- Default
- Field
- Time
- Daily
- Hourly



Kafka Connect Sink

[Deploying Connectors]



Meet Data Sinks



Deploying Connectors



MDW & OLAP Systems

Cloud-Based Data Warehouses Systems

- Amazon Redshift
- Azure Synapse Analytics
- Google BigQuery
- Snowflake



Physical Hardware

- Massively Parallel Processing [MPP]
- “Loosely Coupled” & “Shared Nothing”



Columnar Data Storage Type

- Batch-Processing Mode
- Compression Benefits
- I/O Reduction Operations
- Index Bitmap



MDW

*OLAP
Real-time*

Real-Time Ingest

- Exactly Once Semantics
- Directly Source from Apache Kafka as Consumer
- Designed for Streaming Technologies



Users

- SQL using Apache Calcite
- API using RestAPI Calls

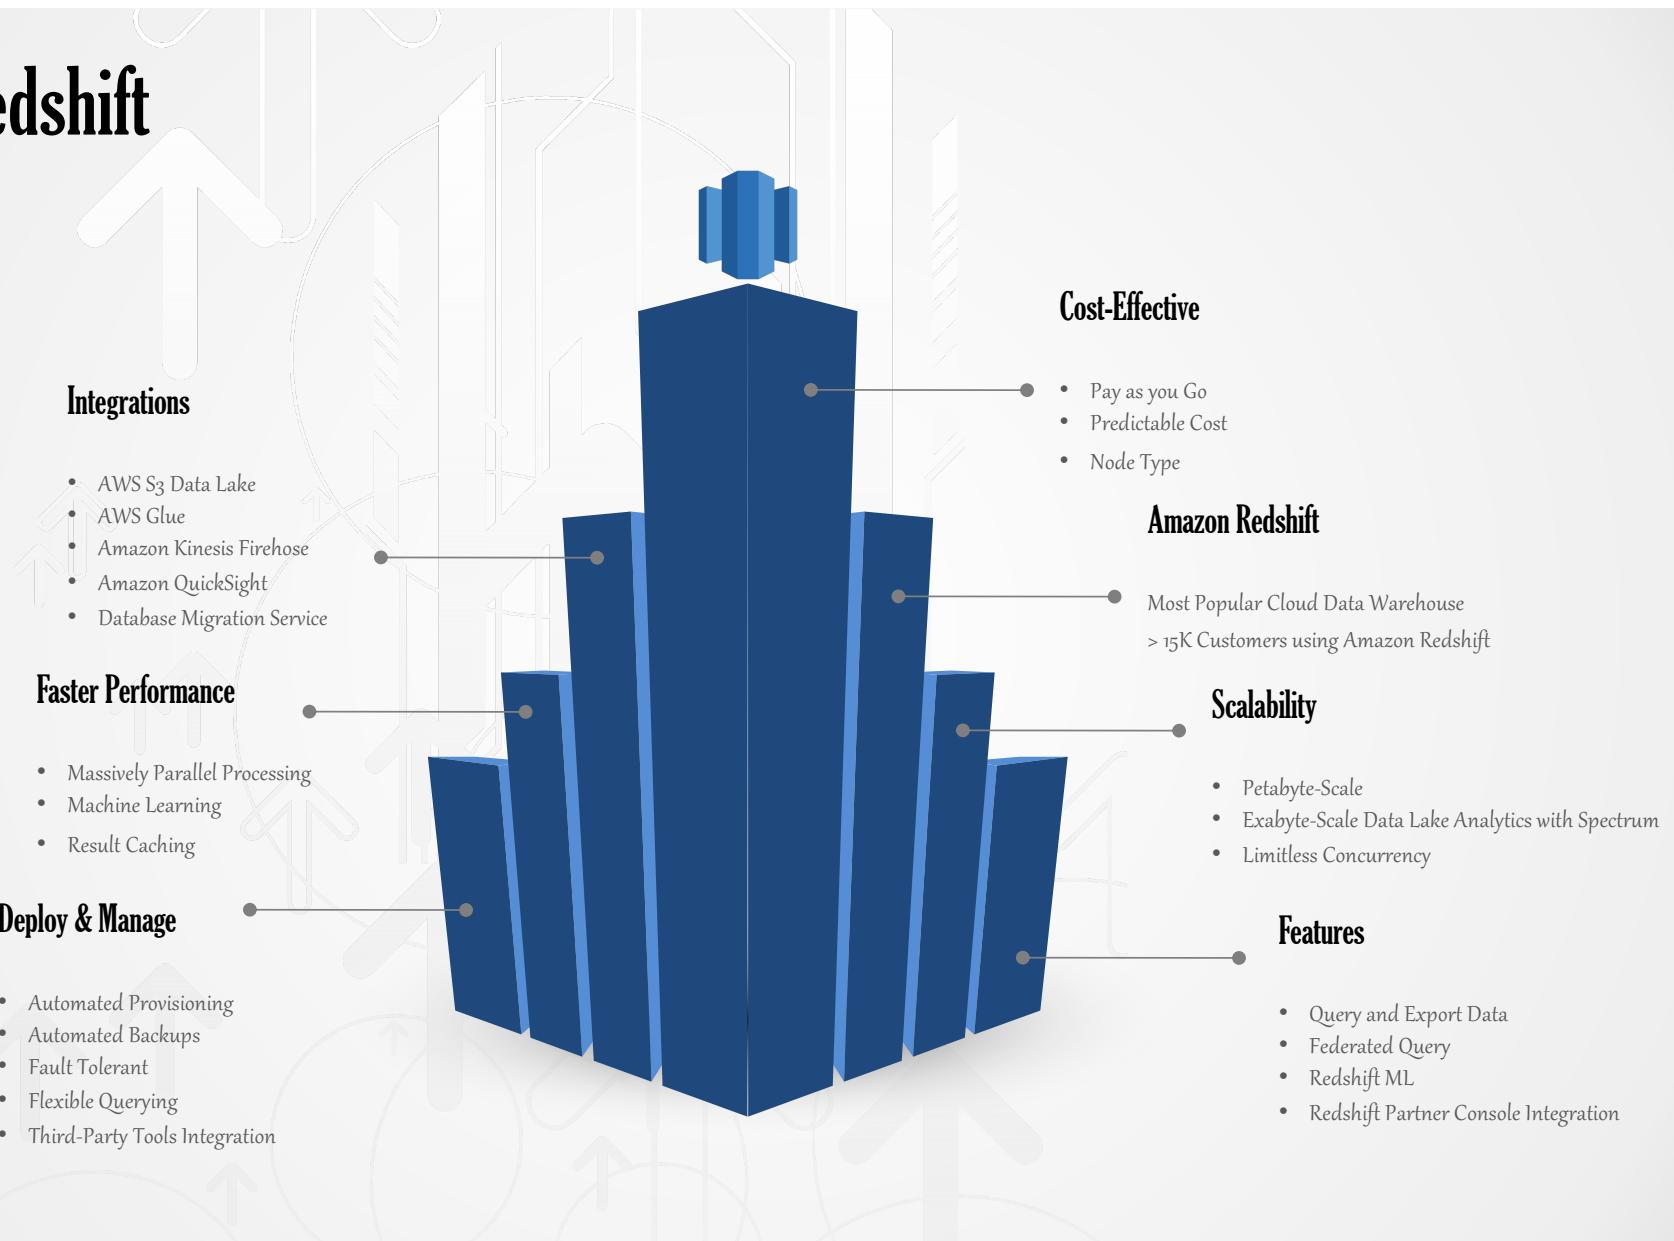


Key Capabilities

- Optimized for Multiple Connections
- Columnar-Oriented Storage
- Horizontally Scale & Fault-Tolerant
- Joins using Presto & Trino



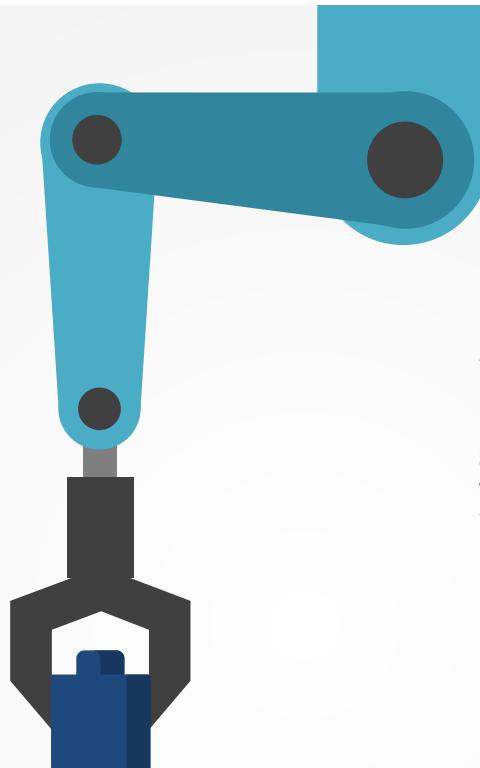
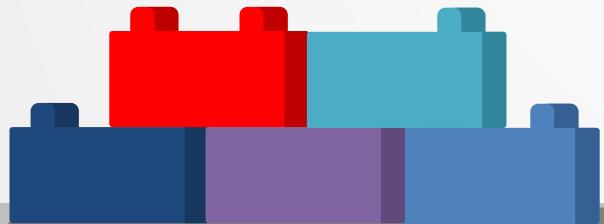
Amazon Redshift



Azure Synapse Analytics

Features

- EDW - Azure SQL Dw Engine
- Data Lake Exploration
- Languages – T-SQL, Python, Scala, Spark SQL & .NET
- Orchestration – Azure Data Factory
- Streaming Ingestion & Analytics
- Integrated AI & BI



Components

- Built-In SQL Pool
- DW using - SQL Dw
- Apache Spark Pools
- Data Integration & Studio

Unified Experience

Significantly Reduce Project Development Time with a Unified Experience for Developing End-to-End Analytics Solutions



Google BigQuery



Google's Serverless Offering

- Automatic Resource Provisioning
- SaaS Offering for Dw

Storage & Computation

- Separated Storage & Compute
- Choose Storage Tier
- Control Costs

Programmatic Interaction

- REST API
- Java
- Python
- Node.js
- C#
- Ruby
- PHP

Big Data Ecosystem Integration

- Cloud DataProc
- Cloud DataFlow
- Apache Big Data Ecosystem
- Apache Hadoop & Apache Beam

Foundation for BI & AI

- EDW Google's Offering
- Integration, Transformation & Analyzes
- TensorFlow & BigQuery ML

Real-Time Analytics with SQL

- High-Speed Streaming Insertion API
- Standard ANSI:2011 SQL Support
- ODBC & JDBC Drivers

Snowflake

Fast and Easy SQL Analytics

- In Seconds, Provision Compute Clusters Ranging in Size from Extra-Small ~ 6XL
- Fully ANSI SQL Compatible, with Native Support for Semi-Structured Data
- Optimized Direct Connectors for Popular BI and Analytics Tools

Use All of Your Data

- Near-Unlimited, Low-Cost Cloud Storage with 2-3x Compression
- A Single Copy of Your Data Available Immediately Everywhere
- Native Support for Geospatial Data and Analysis



Accelerate BI & Analytics for All Users

Support a Virtually Unlimited Number of Concurrent Users and Queries with Near-Unlimited, Dedicated Compute Resources. Query Semi-Structured Data Directly with SQL, and with Your Favorite Business Intelligence and Machine Learning Tools

Single Source for All Data

Creates a Single, Governed, and Immediately Queryable Source for Effectively All Your Data, Including JSON and XML, with Near-Unlimited and Cost-Effective Storage

Managed Service

- Automatic Query Caching, Planning, Parsing, and Optimization
- Automatic Updates with No Scheduled Downtime
- Cross-Cloud Data Replication for Seamless, Global Data Access

Security, Governance & Privacy

- Automatic Data Encryption At-Rest and In-Transit
- Dynamic Data Masking and Tokenization
- Leverage Functionality ~ GDPR and CCPA Compliance
- Certifications for SOC2 Type 2, ISO 27001, PCI, HIPAA, FedRAMP

Apache Pinot

Characteristics

- Blazing Fast OLAP Engine
- LinkedIn, Uber, Target, Slack, Stripe
- Designed for Real-Time Analytics

History

- Pinot Noir ~ Name of Grape for Wine
- Developed Internally at LinkedIn in 2014
- Open-Sourced in 2015
- Entered in Apache Incubator in 2018

Key Capabilities

- Columnar-Oriented Storage
- Pluggable Indexing Technologies
- Horizontally Scale & Fault-Tolerant
- Performs Anomaly Detection using ThirdEye
- Joins using Presto & Trino

Apache Pinot

Realtime Distributed OLAP DataStore, Designed for Answering OLAP Queries with Low-Latency

Performance

- 100K + Queries per Second at ms Latency
- Ingestion of Million of Events per Second
- 50+ User-Facing Analytics

Users

- PQL ~ SQL using Apache Calcite
- API ~ RestAPIs for Broker & Controller
- External Clients ~ JDBC, Java, Python & Go

Core Components

- Controller (Helix & Zookeeper) ~ Manage State & Health
- Broker ~ Route Queries
- Server ~ Host Segments (Data) ~ Realtime & Offline Tables
- Minion ~ Purge Data



Data Virtualization

Access Data in Any Source, Shape or Structure without Copying from Source



Processing Query Engine

Connects, Discover, Plan and Analyze Different Data Stores for Query at Scale



Modern Data Warehouses

Redshift, BigQuery



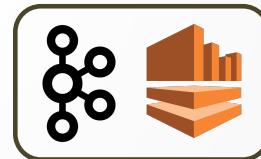
Real-Time OLAP

Apache Druid & Apache Pinot



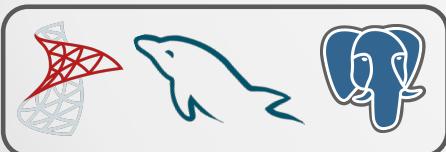
Real-Time Ingestion

Apache Kafka, Kinesis



Relational Databases

SQL Server, MySQL, Postgres



NoSQL

Cassandra, MongoDB, Redis, ElasticSearch



Data Lake

HDFS, S3, Blob Storage, GCS & MinIO [S3*]



Trino

Fastest Bunny on Forest, Eating Any Source of Data



Trino

- Speed & Scale
- Simplicity
- Versatile
- In-Place Analysis
- Query Federation
- Runs Everywhere
- Trusted
- Trino Software Foundation [OSS]



Designed For

- Access Data using SQL
- Data Warehousing
- Analytics
- OLAP

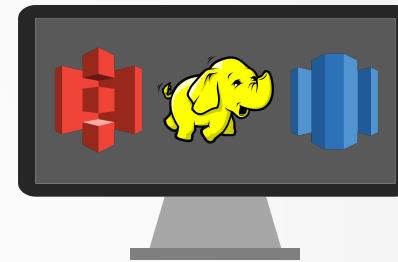


Trino is a Tool Designed to Efficiently Query Vast Amounts of Data using Distributed Queries in a Terabytes & Petabytes Range



Query Execution Model

Turn SQL Statements into Queries,
Executes Across a Distributed
Cluster of Coordinator and Workers



Features

- Memory Management & Spilling using Compression and Encryption
- Resource Groups
- Distributed Sort
- Dynamic Filtering ~ Query Performance, Network Traffic Reduction & Remote Load Reduction
- Cost Based Optimization - CBO
- Query Pushdown



Apache ORC ~ Optimized Row Columnar

- Smallest, Fastest Columnar Storage for Apache Hadoop Workloads
- ACID Support
- Built-In Indexes
- Complex Types
- Integration with Apache Spark & Trino

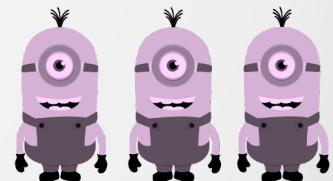
Coordinator a.k.a The Brain

Parsing Statements, Planning Queries, &
Managing Trino Worker Nodes



Workers a.k.a The Slave

Responsible for Executing Tasks & Processing Data. Fetch Data from Connectors and Exchange Intermediate Data in Parallel





Success is not final, failure
is not fatal: it is the courage
to continue that counts.

Winston Churchill

 quotefancy



ONEWAY
SOLUTION