



**ONEWAY**  
SOLUTION



One Way Solution

# Best Practices & Guidelines

Data Engineering – [Day 5]



JUAN MORENO

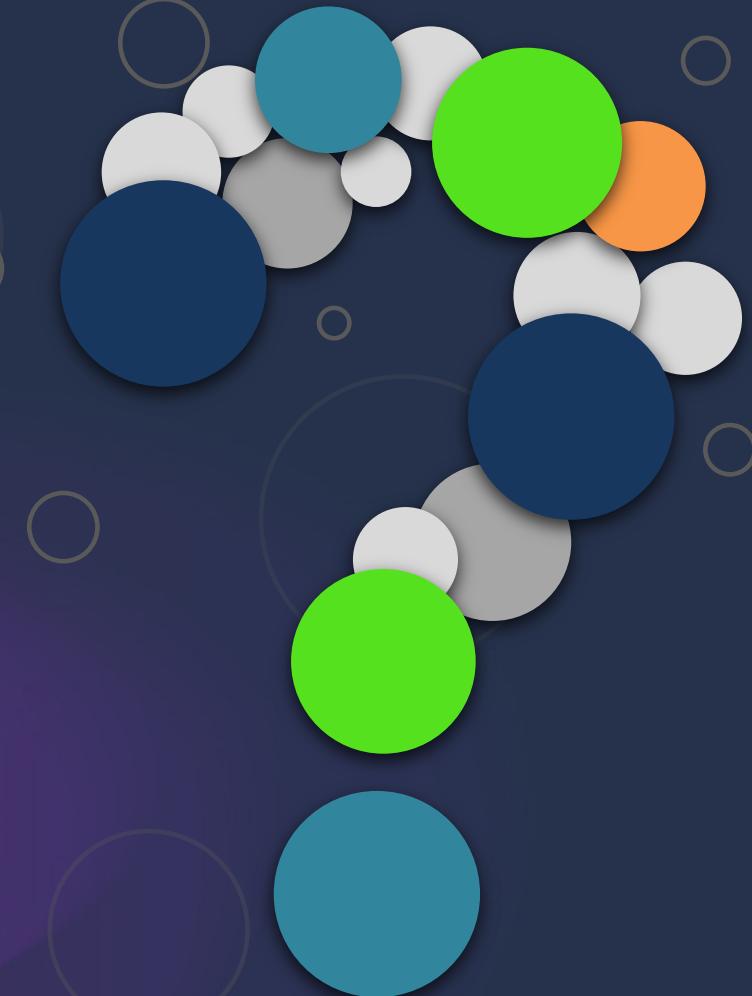
CEO & Data Architect  
Data Engineer & MVP



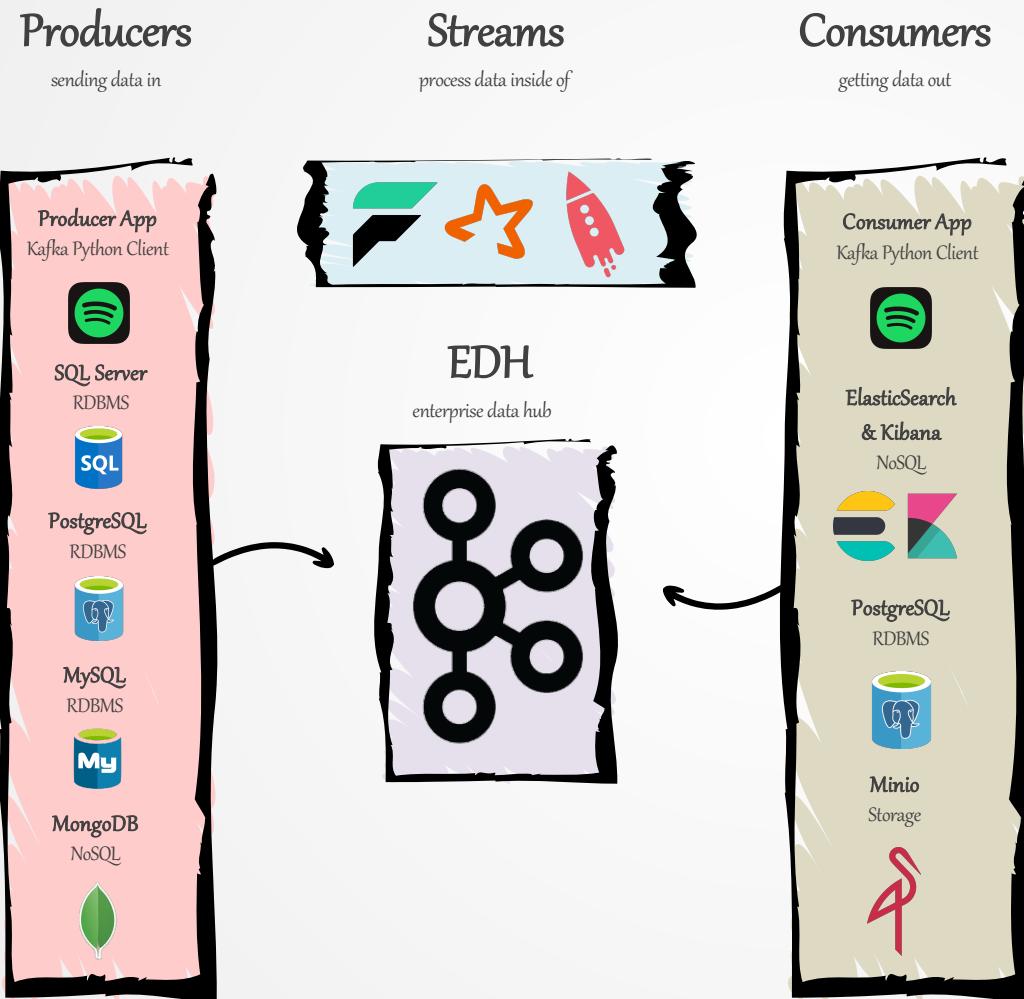
MATEUS OLIVEIRA

Big Data Architect  
Data In-Motion Specialist





# Use Case ~ Music Streaming using [Apache Kafka]





Courage is what it takes to stand  
up and speak; courage is also what  
it takes to sit down and listen.

Winston Churchill

Being deeply loved by someone gives you strength, while loving someone deeply gives you courage.

Lao Tzu



# *Advanced Strategies for Apache Kafka Development*

# Message Ordering in Apache Kafka

*Tip:* if ordering is necessary and we can work in producer configs then we must use only 1 partition



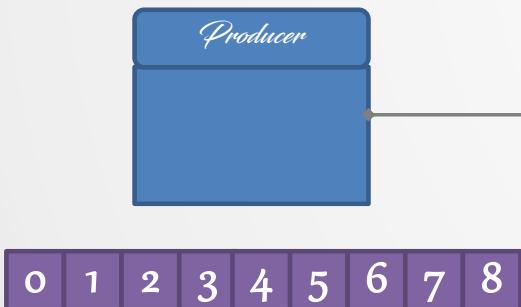
*Apache Kafka:* guarantee message ordering in two levels, the producer level with eos configuration or in partition level, for example if we have only one partition, we always will receive the events in the right order



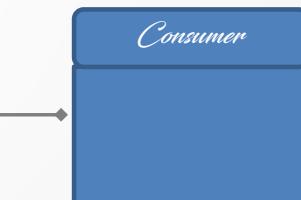
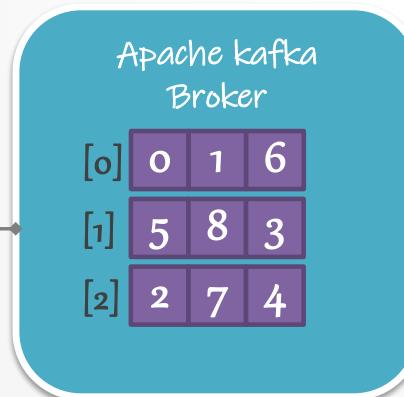
*Downside:* with only 1 partition its not possible to scale write and reads applications because is based on partition strategy

## Producer Configs

- Exactly-Once Enabled
- max.inflight.connectors <= 5
- acks=all



events that is going to be written into kafka sequence of number 0 - 8



events will be read in sequence

## Consumer Configs

order is placed on the producer side, without the need to specify on the consumption

consumer reads based on the offset and partition in this case we have the ordering because the producer writes in the right order into the topic, segregate by partition, when consume kicks in it reads the last offset which has the last number of the sequence

## Apache Broker

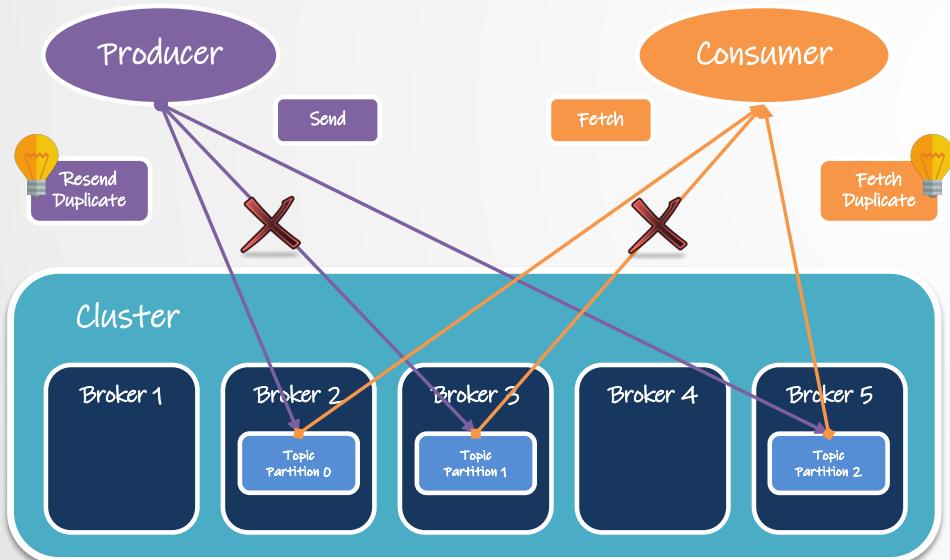
if the producer doesn't have eos enabled with the producer configurations, then kafka can't guarantee ordering with more than one partition

# Performing Transactions with Producers & Consumers



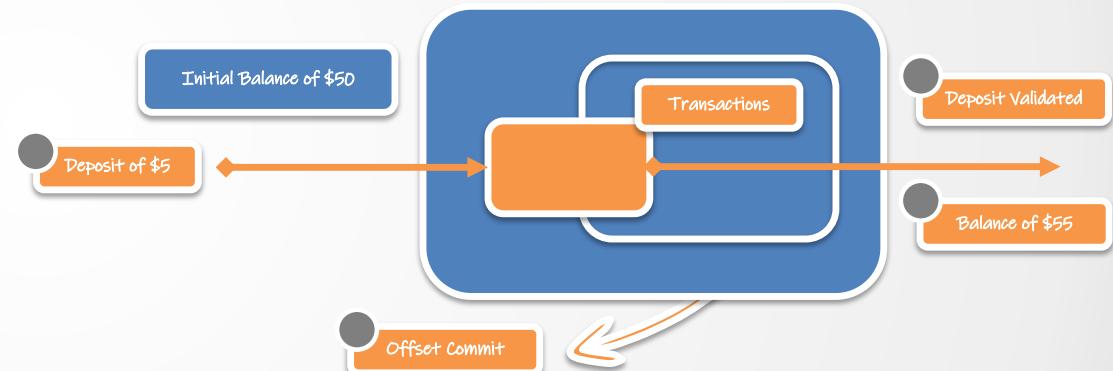
## Two Opportunities for Duplicates

- **Producer** ~ Duplication Upon Retrying Write Operation
- **Consumer** ~ Failure During Increment Offset Write



## Transactions ~ Multi-Operation-Commits

- **State Store** ~ Disk-Residence Hash Table Backed By Apache Kafka
- Input and Output Topics in an Atomic Way
- Stream Processors ~ Kafka Streams and KSQLDB



## Transactions ~ In a Nutshell

- Stream Applications with **Read-Process-Write** Pattern
- Exactly-Once-Semantics with Atomicity & Facilitating Zombie Fencing
- Transaction Coordinator (Kafka Broker) & Log (Kafka Topic)
- 1KB Records, 100 MS ~ 3% of Degradation in Throughput

# Stream-Table Duality

## Stream

event stream records the history of what has happened in the world as a sequence of events

- 1) Immutable Data
- 2) Only Inserting (Appending) New Events
- 3) Persistent, Durable & Fault Tolerant
- 4) Keyed Stream (Key)

Stream



Immutable with  
Historic Information

## Table

represents the state of the world (snapshot)

- 1) Mutable Data
- 2) Insert, Update, Delete
- 3) Keyed Event Identifies Row
- 4) Persistent, Durable & Fault Tolerant

Table



Mutable without  
Historic Information

## Stream-Table Duality

### Stream-Table

- Stream into a Table by Aggregating Events using COUNT() or SUM()
- Streams = Apache Kafka Topics with Enforced Schema

### Table-Stream

- Capture Changes (Insert, Update & Delete) into a Change Stream, (CDC)
- Tables = Topics using Log Compression (Key)

Streams = Record History

Table = Represent State

Agg = SUM | COUNT



Duality



Table Changes

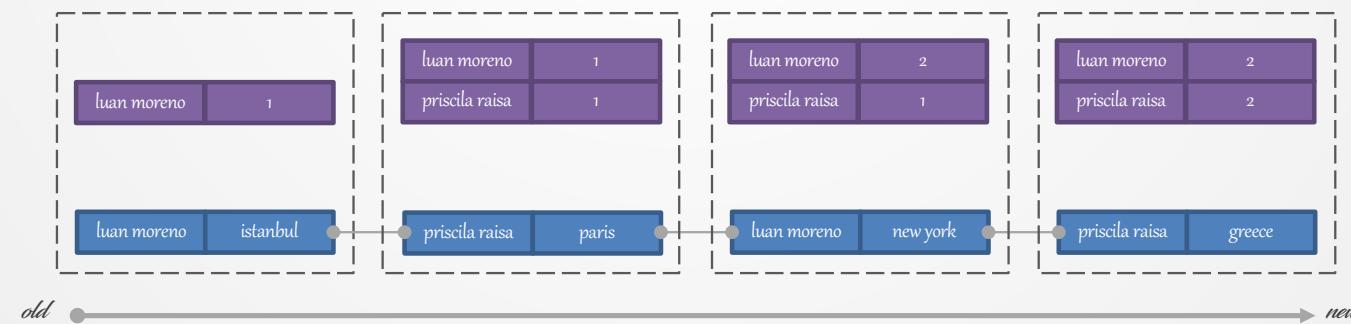


Table

Stream

Last State

Sequence of Events



# Advanced Windowing [Stream-Stream Join]



## Window Types

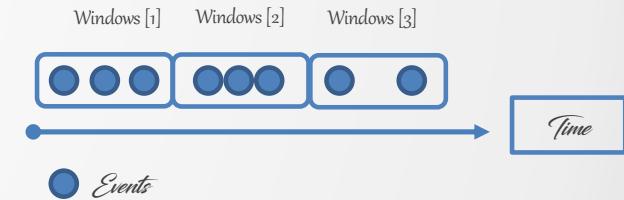
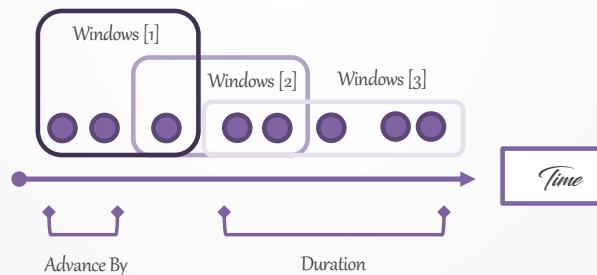
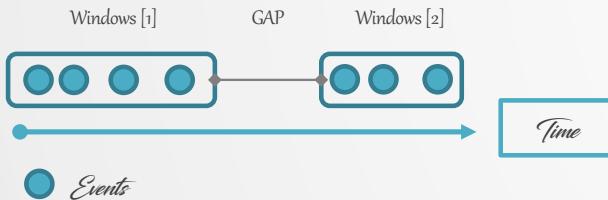
most operations on streams are windowed operations

**Session** = aggregates records into a session, which represents a period of activity separated by a specific gap of inactivity (idleness). If a record's timestamp occurs outside of the session gap, new session is created

**Hopping** = based on time intervals. it is defined by two properties, duration and it's advance [HOP] interval, this HOP specifies how far a window moves forward in time relative to the previous window

**Tumbling** = special case of hopping, it is based in time interval too. it is defined by a single property: windows duration. this type of window never overlap a record belong once and only one window

Window Type	Mode	Description
Session	Session-Based	Dynamically-Sized, Non-Overlapping, Data-Driven Windows
Hopping	Time-Based	Fixed-Duration, Overlapping Windows
Tumbling	Time-Based	Fixed-Duration, Non-Overlapping, Gap-Less Windows



```
--session
SELECT regionid, COUNT(*)
FROM pageviews
WINDOW SESSION (60 SECONDS)
GROUP BY regionid
EMIT CHANGES;
```

```
--hopping
SELECT regionid, COUNT(*)
FROM pageviews
WINDOW HOPPING (SIZE 30 SECONDS, ADVANCE BY 10 SECONDS)
GROUP BY regionid
EMIT CHANGES;
```

```
--tumbling
SELECT orderzip_code, TOPK(order_total, 5)
FROM orders
WINDOW TUMBLING (SIZE 1 HOUR)
GROUP BY order_zipcode
EMIT CHANGES;
```

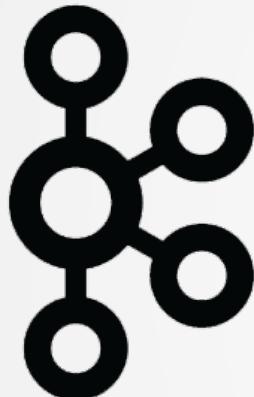


# *Advanced Strategies for Apache Kafka Administration*

# Log Skimming on Apache Kafka using ElasticSearch

## Apache Kafka

produce logs of all his activities



**Info**  
information about activities, some troublesome behaviors



**Warn**  
information about out of normal behavior



**Error**  
information about critical error, most import ones

## Filebeat

access log folders and copy the logs info inside elasticsearch



## Elasticsearch

database where logs are stored, fully integrated with filebeat and kibana



## Kibana

visualize and explore logs of data captured by elasticsearch

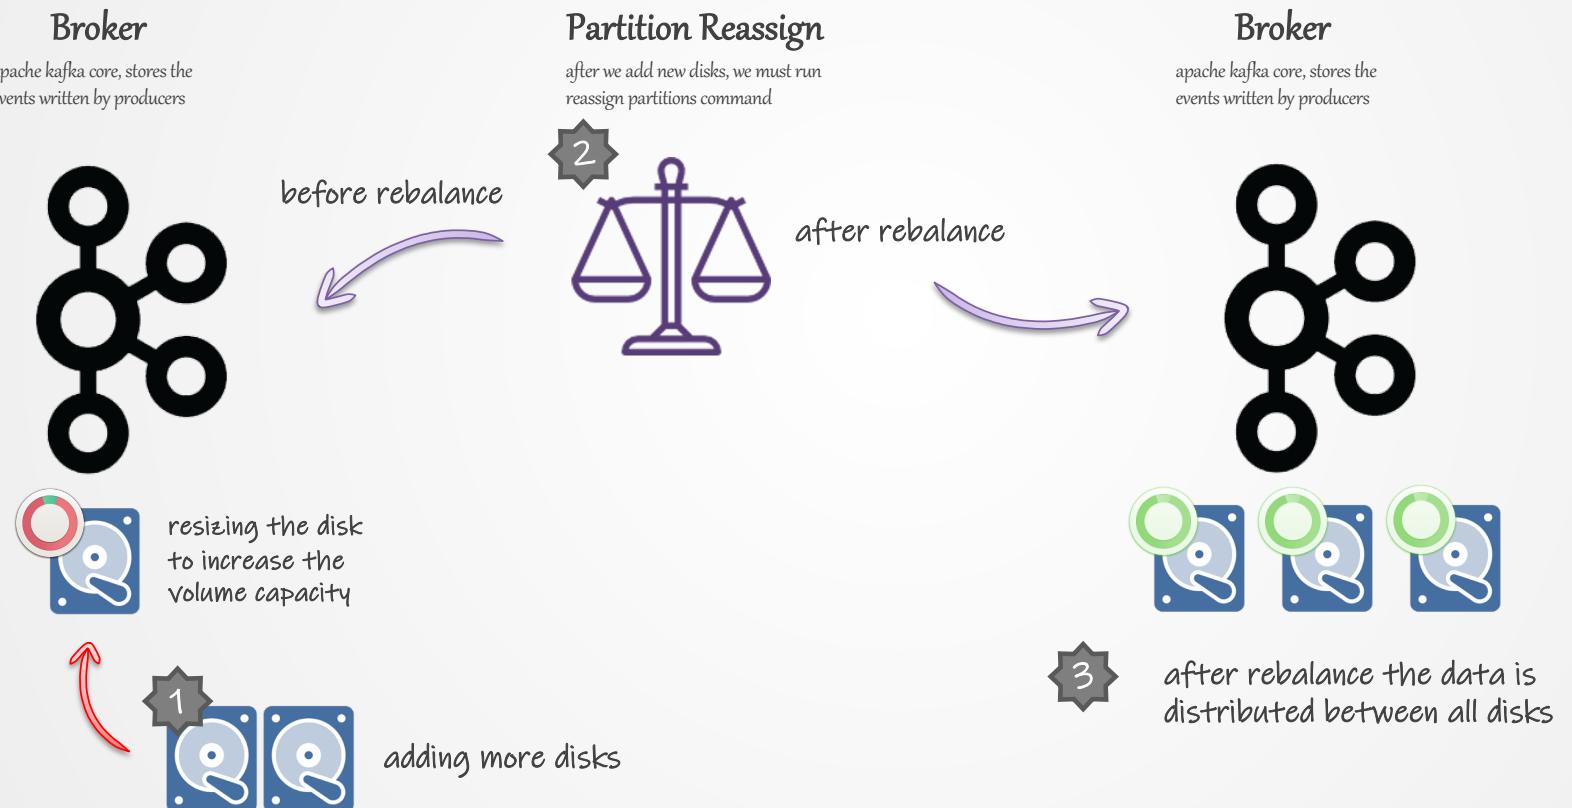


# Storage Resizing for Apache Kafka



**JBOD=Just a Bunch of Disks**

in KIP-112 we have jbod handle fix, and we recommend the use in kafka deployments, specially if we are talking about kubernetes deployments, we don't need to resize only add more disk



# Apache Kafka Partition Reassign - Rebalance



## Partition Reassign

Apache Kafka is a storage in real-time for streaming events which depend on disks to store segments, and we must be careful to not reach 100% disk occupation, but if we need to add more disks, we need to execute a rebalance equals to partition reassign command line after add the new disks



## Warnings

- only execute out of business hours, the topics becomes offline during the rebalance, nobody writes or reads
- simple math 10 topics with 8 partitions each and data from 1mi to 1.000 can take something close to 30 seconds to rebalance, but this depends a lot of how unbalanced are things and if we are using SSD

## How partition reassign works



### List Topics & Create JSON File

create a json file base in the topics in this pattern:  
`{ "topics": [ { "topic": "_consumer_offsets" }, { "topic": "_schemas" } ] }`



### Create a Reassign JSON File with Output from Command Line

```
"kubectl exec edh-kafka-0 -c kafka -i -t -- bin/kafka-reassign-partitions.sh --bootstrap-server localhost:9092 --broker-list "0,1,2" --topics-to-move-json-file /tmp/topics-to-move.json --generate"
```



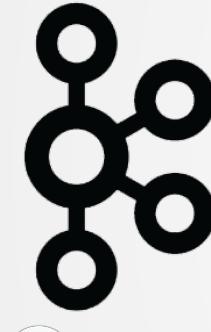
### Execute Command with JSON Reassign File

```
"kubectl exec edh-kafka-0 -c kafka -i -t -- bin/kafka-reassign-partitions.sh --bootstrap-server localhost:9092 --reassignment-json-file /tmp/cluster-reassign.json --execute"
```

# Automating Administrative Tasks using Cruise Control

## Broker

apache kafka core, stores the events written by producers



resizing the disk  
to increase the volume capacity



adding more disks

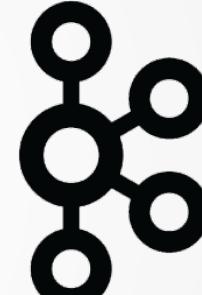
## Cruise Control

based on policies cruise control oversee the cluster and propose plans to rebalance when needed, and we can approve this plans and apply the changes



## Broker

apache kafka core, stores the events written by producers



3 the data is distributed between all disks

# Monitoring [Apache Kafka] with Prometheus & Grafana



## Prometheus

- Open-Source System for Monitoring & Alerting ~ 2012
- Time-Series DB ~ Streams of Timestamped Values
- Built-In ~ Scraping, Storing, Querying & Graphing System
- PromQL ~ Flexible Query Language
- Pull over HTTP
- Written in Go



*jmx*



## Grafana

- Open-Source Visualization & Analytics Software
- Query, Visualize, Alert & Metrics
- Time Series Database ~ Graphs & Visualizations
- Data Sources ~ Prometheus, ElasticSearch, MySQL, PostgreSQL



*jmx*



*metrics*



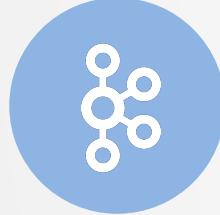
# Monitoring Apache Kafka with ElasticSearch and Prometheus



# DataOps with Lenses

## Fancy UI

Explore Topics, Events, Schemas, Partitions, Health of Connectors, Apache Brokers & Apache Zookeepers



## Observability Apache Kafka

Lenses Gives a 360° View of Apache Kafka Components:

- Apache Broker
- Apache Zookeeper
- Kafka Connect
- Schema Registry

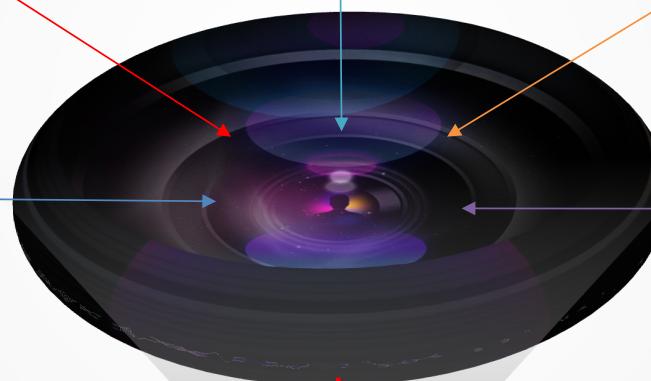
## Deploy Everywhere

- Virtual Machines (VMs)
- Docker & Kubernetes
- Public Cloud



## Explore Events using SQL

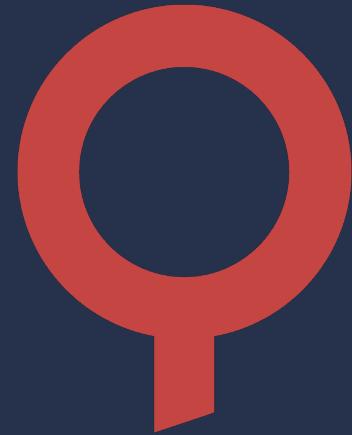
Query Data using SQL with Abstraction of Streams and Tables with Similar KSQLDB Experience



## Features

- App Catalog
- Centralized Schema Management
- Consumer Lag Monitoring
- Custom Alert Routing
- Data Catalog
- Data Discovery
- Data Injection
- Data Masking Rule

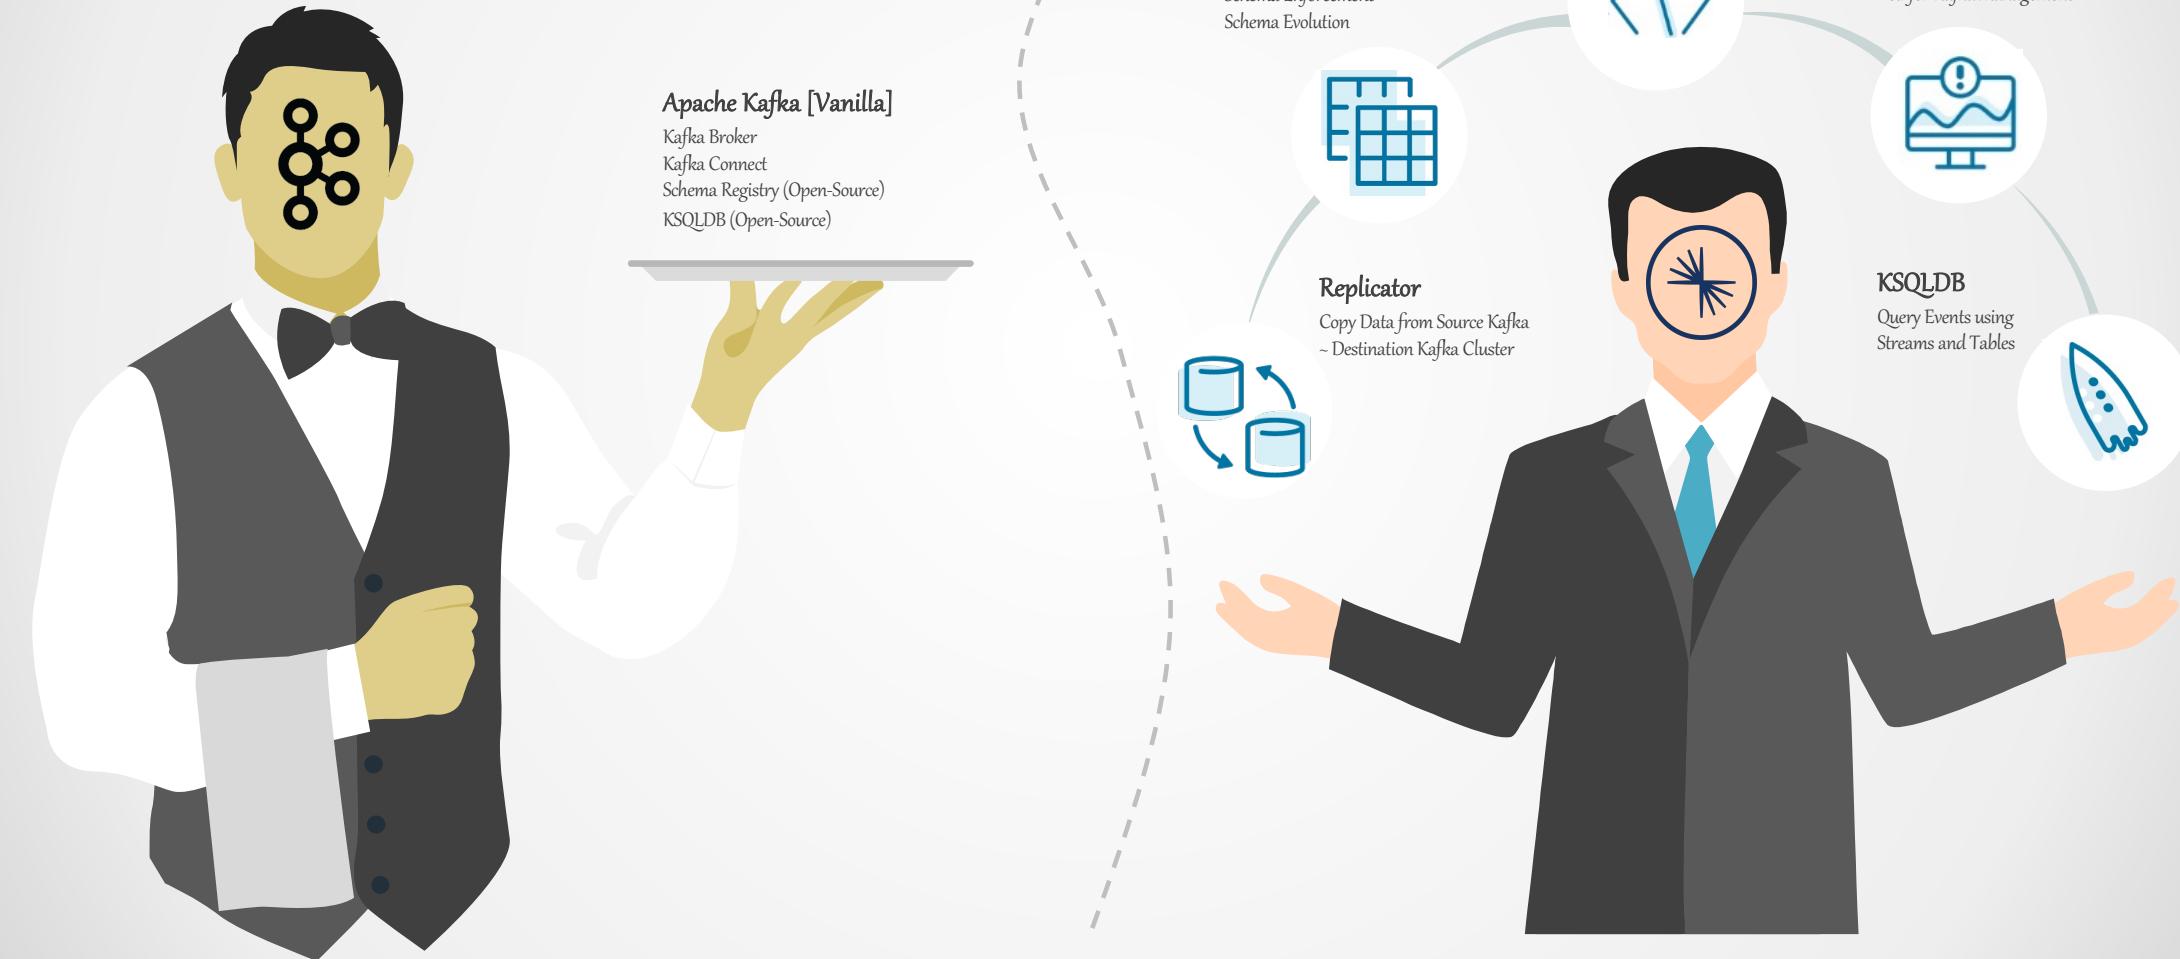
# DataOps using Lenses



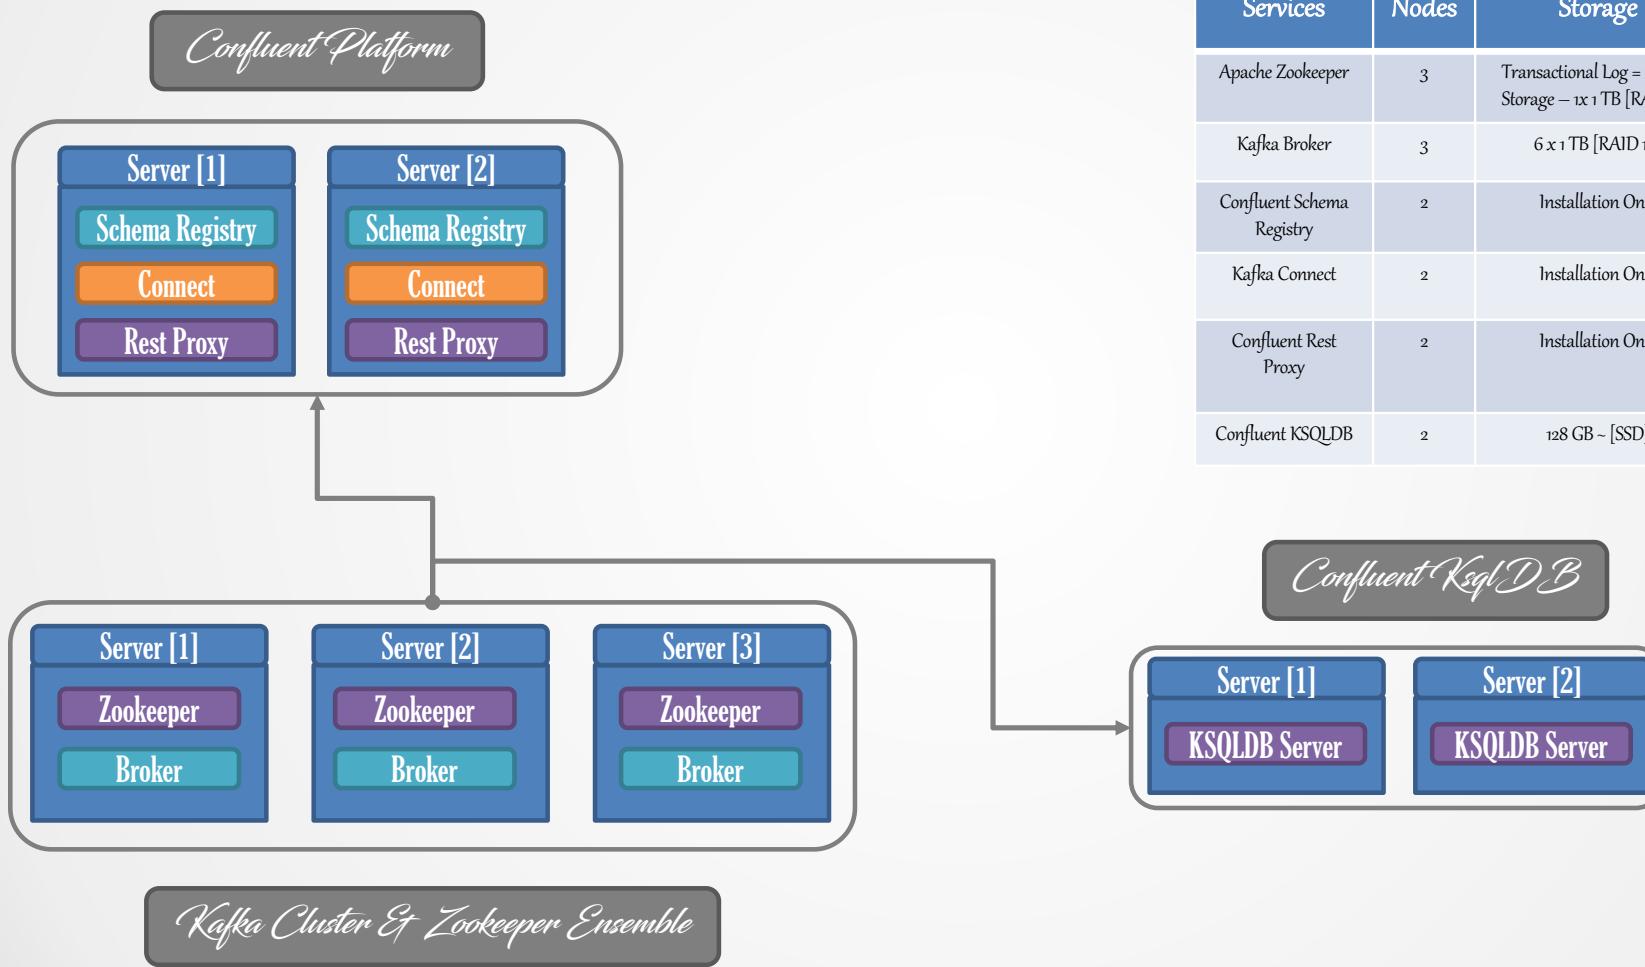


# *Advanced Deployment Techniques for Apache Kafka*

# Apache Kafka & Confluent [Services]

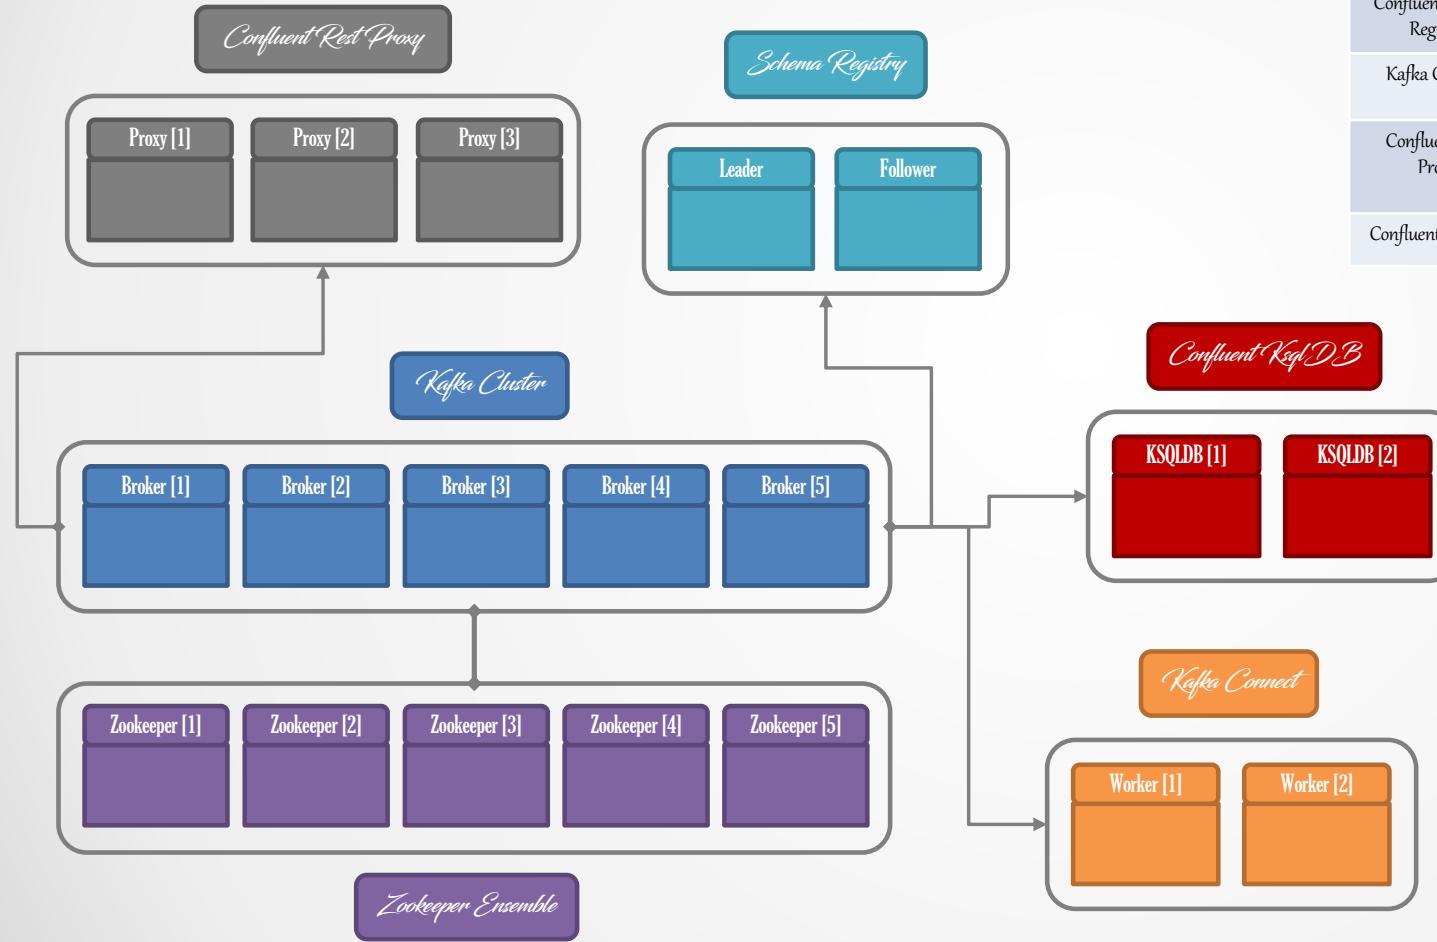


# Apache Kafka & Confluent [Architecture] ~ Small



Services	Nodes	Storage	Memory	CPU
Apache Zookeeper	3	Transactional Log = 256 GB Storage – 1x 1 TB [RAID 10]	16 GB	2~4
Kafka Broker	3	6 x 1 TB [RAID 10]	32 GB +	6
Confluent Schema Registry	2	Installation Only	1 GB Heap Size	2~4
Kafka Connect	2	Installation Only	2 GB Heap Size	2~4
Confluent Rest Proxy	2	Installation Only	Producer = 512 MB Consumer = 16 MB	16
Confluent KSQLDB	2	128 GB ~ [SSD]	10 GB	2

# Apache Kafka & Confluent [Architecture] ~ Large



Services	Nodes	Storage	Memory	CPU
Apache Zookeeper	5	Transactional Log = 512 GB Storage – 2x 1 TB [RAID 10]	32 GB	2~4
Kafka Broker	5	12 x 1 TB [RAID 10]	64 GB +	12
Confluent Schema Registry	2	Installation Only	2 GB Heap Size	2~4
Kafka Connect	2	Installation Only	4 GB Heap Size	2~4
Confluent Rest Proxy	3	Installation Only	Producer = 1 GB Consumer = 32 MB	16
Confluent KSQLDB	2	256 GB ~ [SSD]	20 GB	4

# Kafka Deployment on Kubernetes using Strimzi



kafka-yaml files



cruise\_control-yaml files



connect-yaml files



bridge-yaml files



connector-yaml files



5 minutes deployment

## Strimzi Deployment



Broker



Connect



Zookeeper



Cruise Control



Bridge



Connector

## Custom Resources Definition



operator control the kafka kind objects



modify statefulsets without stop services



logs of all crds can be seen in the operator

## Kubernetes Benefits



scale as business need, as much as you like



easy management & deployment process



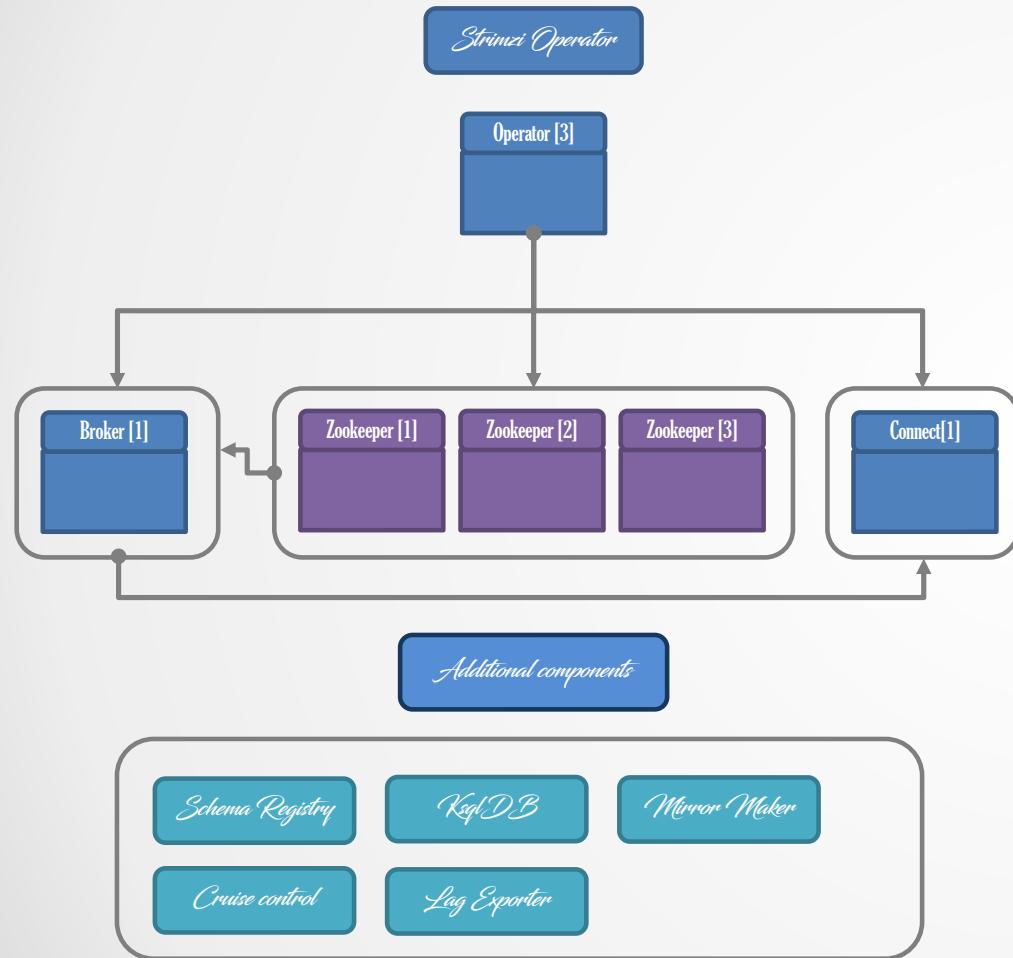
huge amount of big data products going to kubernetes



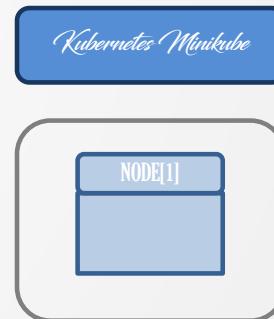
centralize log with filebeat and elasticsearch



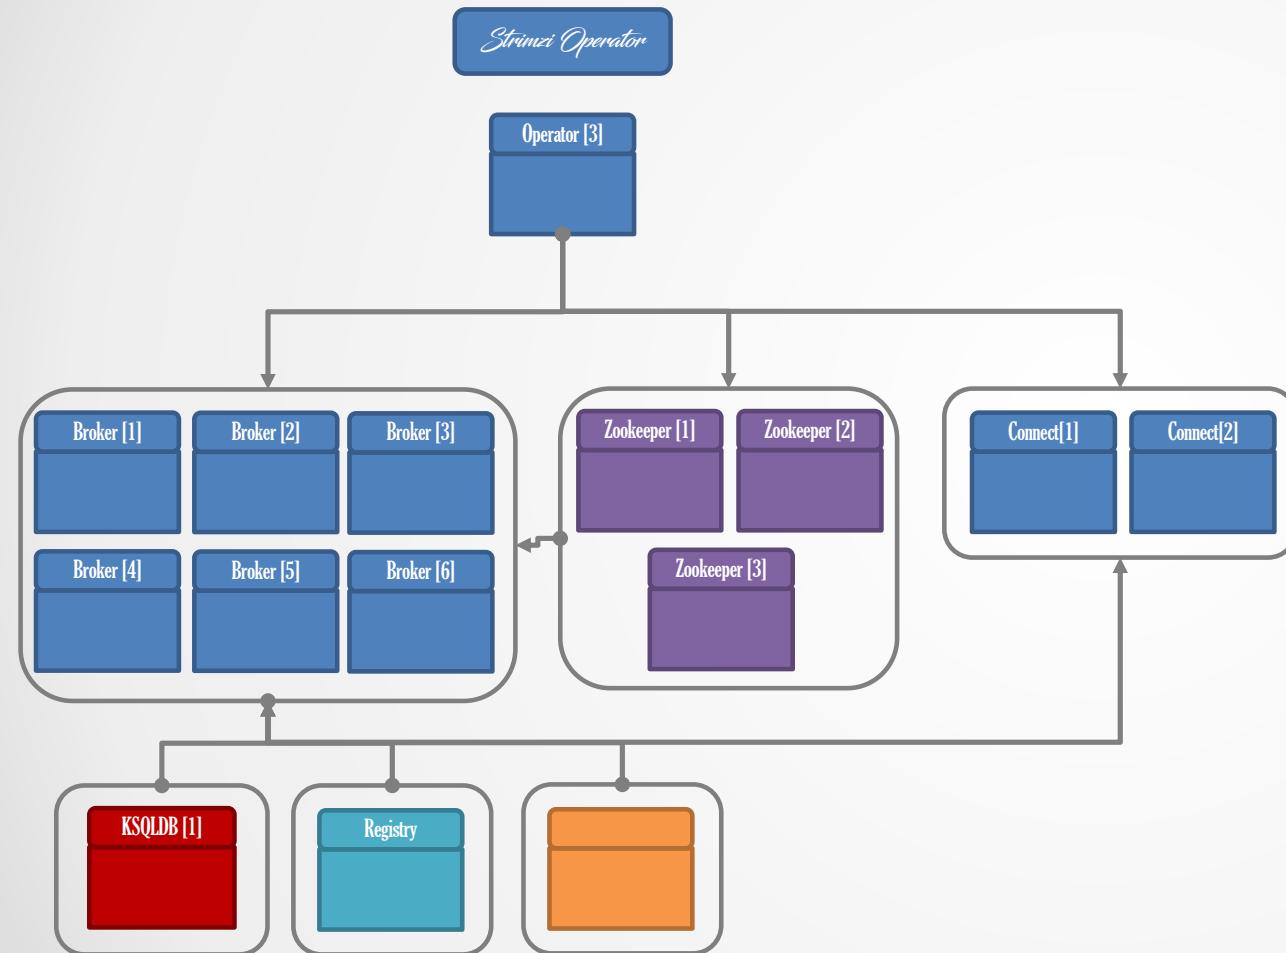
# Kafka Deployment on Kubernetes [K\*S] ~ Small



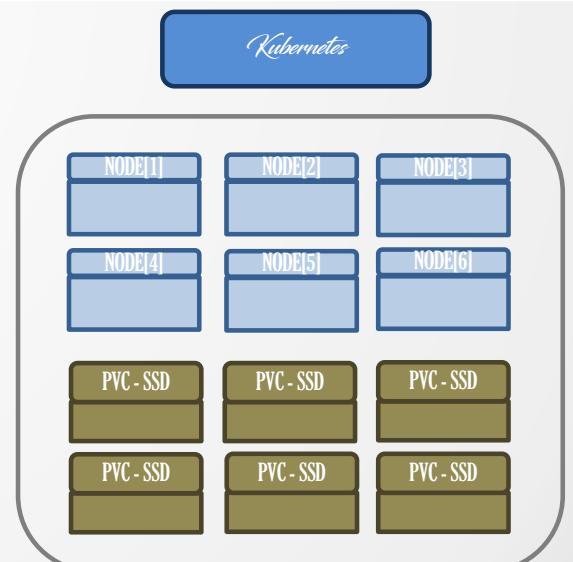
Services	Pods	Storage	Memory	CPU
Strimzi Operator	1	Ephemeral	Requests = 348Mi Limits = 348Mi	Requests = 200m Limits = 1000m
Apache Zookeeper	3	5Gi	Requests = 256Mi Limits = 512Mi	Requests = 250m Limits = 500m
Kafka Broker	1	40Gi or Ephemeral	Requests = 500Mi Limits = 1Gi	Requests = 250m Limits = 500m
Kafka Connect	1	Ephemeral	Requests = 500Mi Limits = 1Gi	Requests = 250m Limits = 500m
Strimzi bridge	1	Ephemeral	Requests = 100Mi Limits = 200Mi	Requests = 100m Limits = 200m
Confluent Schema Registry	1	Ephemeral	Requests = 100Mi Limits = 100Mi	Requests = 100m Limits = 200m
Confluent KSQL	1	Ephemeral	Requests = 512Mi Limits = 2Gi	Requests = 250m Limits = 500m



# Kafka Deployment on Kubernetes [K\*S] ~ Large



Services	Pods	Storage	Memory	CPU
Strimzi Operator	1	Ephemeral	Requests = 348Mi Limits = 348Mi	Requests = 200m Limits = 1000m
Apache Zookeeper	3	5Gi	Requests = 256Mi Limits = 512Mi	Requests = 250m Limits = 500m
Kafka Broker	6	100GiJBOD	Requests = 1Gi Limits = 2Gi	Requests = 250m Limits = 500m
Kafka Connect	2	Ephemeral	Requests = 1Gi Limits = 2Gi	Requests = 250m Limits = 500m
Strimzi bridge	1	Ephemeral	Requests = 100Mi Limits = 200Mi	Requests = 100m Limits = 200m
Confluent Schema Registry	1	Ephemeral	Requests = 100Mi Limits = 100Mi	Requests = 100m Limits = 200m
Confluent KSQL	1	Ephemeral	Requests = 1Gi Limits = 3Gi	Requests = 500m Limits = 1000m





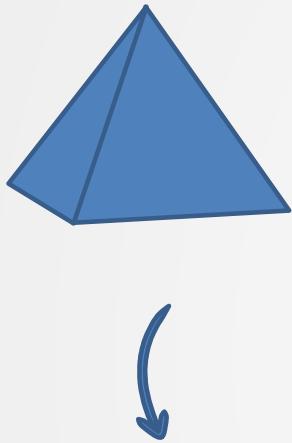
**It takes a lot of courage to show  
your dreams to someone else.**

Erma Bombeck



# *Best Practices for Apache Kafka*

# Optimizing Kafka for [Throughput]



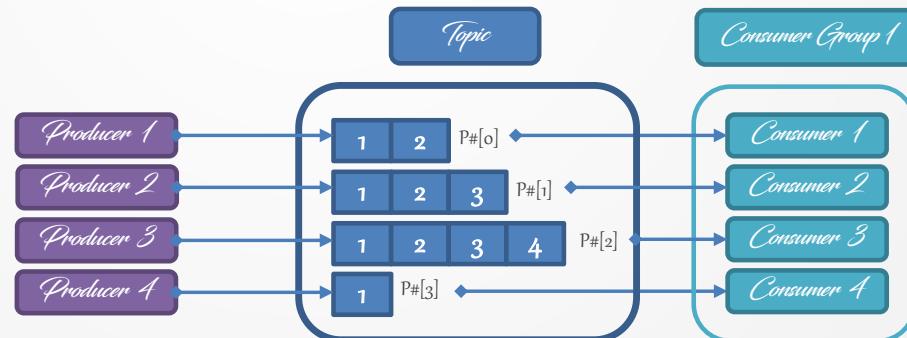
## Optimize for Throughput

to optimize for throughput, the producers, brokers, and consumers need to move as much data as they can within a given amount of time. for high throughput, you are trying to maximize the rate at which this data moves. this data rate should be as fast as possible. a topic **partition** is the unit of parallelism in kafka. messages to different partitions can be sent in parallel by producers, written in parallel by different brokers, and read in parallel by different consumers. in general, a higher number of topic partitions results in higher throughput, and to maximize throughput, you want enough partitions to utilize all brokers in the cluster.

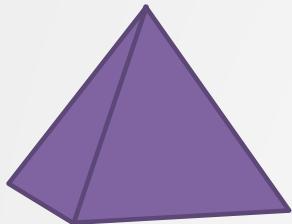
Producer	Value	Default
Batch.Size	100.000 ~ 200.000	16.384
Linger.Ms	10 ~ 100	0
Compression.Type	LZ4	Without Compression
ACKS	1	1
Buffer Memory	Increase [Partitions]	33.554.432

Consumer	Value	Default
Fetch.Min.Bytes	100.000	0

*Throughput*



# Optimizing Kafka for [Latency]



## Optimize for Latency

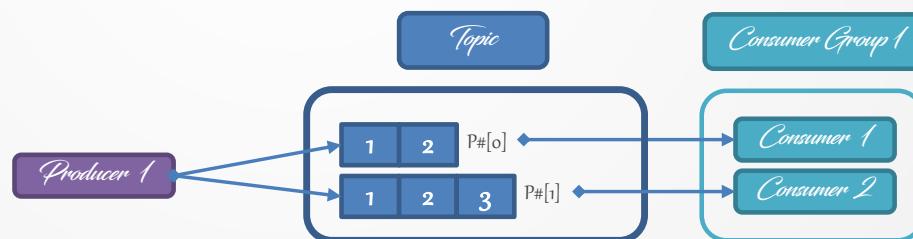
because a [partition](#) is a unit of parallelism in kafka, an increased number of partitions may increase throughput. however, there is a tradeoff in that an increased number of partitions may also increase latency. a broker by default uses a [single thread](#) to replicate data from another broker, so it may take longer to replicate a lot of partitions shared between each pair of brokers and consequently take longer for messages to be considered committed. no message can be consumed until it is committed, so this can ultimately increase end-to-end latency.

Broker	Value	Default
Num.Relicas.Fetchers	Increase [Followers]	1

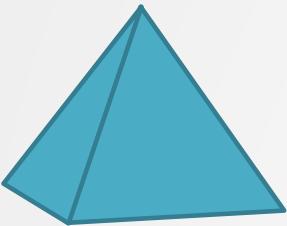
Producer	Value	Default
Linger.Ms	0	0
Compression.Type	NONE	Without Compression
ACKS	1	1

Consumer	Value	Default
Fetch.Min.Bytes	1	1

*Latency*



# Optimizing Kafka for [Durability]



## Optimize for Durability

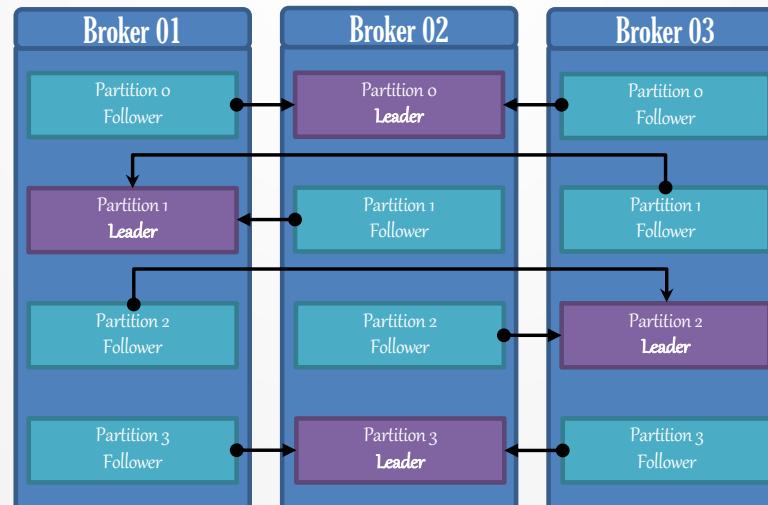
durability is all about *reducing the chance for a message to get lost*. the most important feature that enables durability is *replication*, which ensures that messages are copied to multiple brokers. if a broker has a failure, the data is available from at least one other broker. topics with high durability requirements should have the configuration parameter which will ensure that the cluster can handle a loss of two brokers without losing the data.



*Durability*

**Leader**

**Followers**



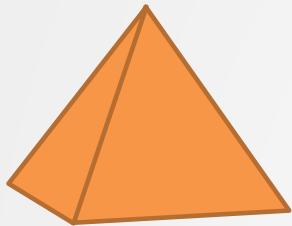
Producer	Value	Default
ACKS	ALL	1
Enable.Idempotence	True	False
Max.InFlight.Requests .per.Connection	1	5

Streams	Value	Default
Replication Factor Configuration	3	1
Exactly Once Semantics	True	At Least Once

Consumer	Value	Default
Enable.Auto.Commit	False	True
Isolation.Level	Read Committed [EOS]	NONE

Broker	Value	Default
Default.Replication.Factor	3	1
Auto.Create.Topics.Enable	False	True
Min.InSync.Replicas	2	1
Unclean.Leader.Election.Enable	False	False
Log.Flush.Interval Message & MS	Low	OS

# Optimizing Kafka for [Availability]



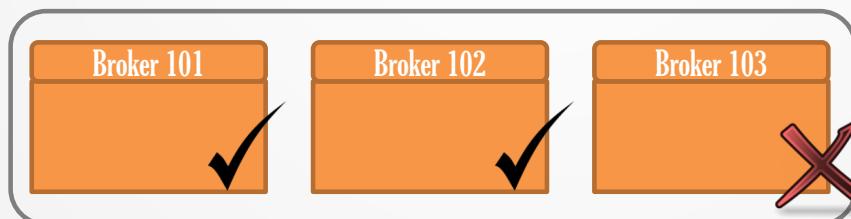
## Optimize for Availability

optimize for high availability, you should tune kafka to recover as quickly as possible from failure scenarios. higher partition counts may increase parallelism but having [more partitions can also increase recovery time](#) in the event of a broker failure. all produce and consume requests will pause until the leader election completes, and these leader elections happen per partition. so, take this recovery time into consideration when choosing partition counts.



*Availability*

*Kafka Cluster*

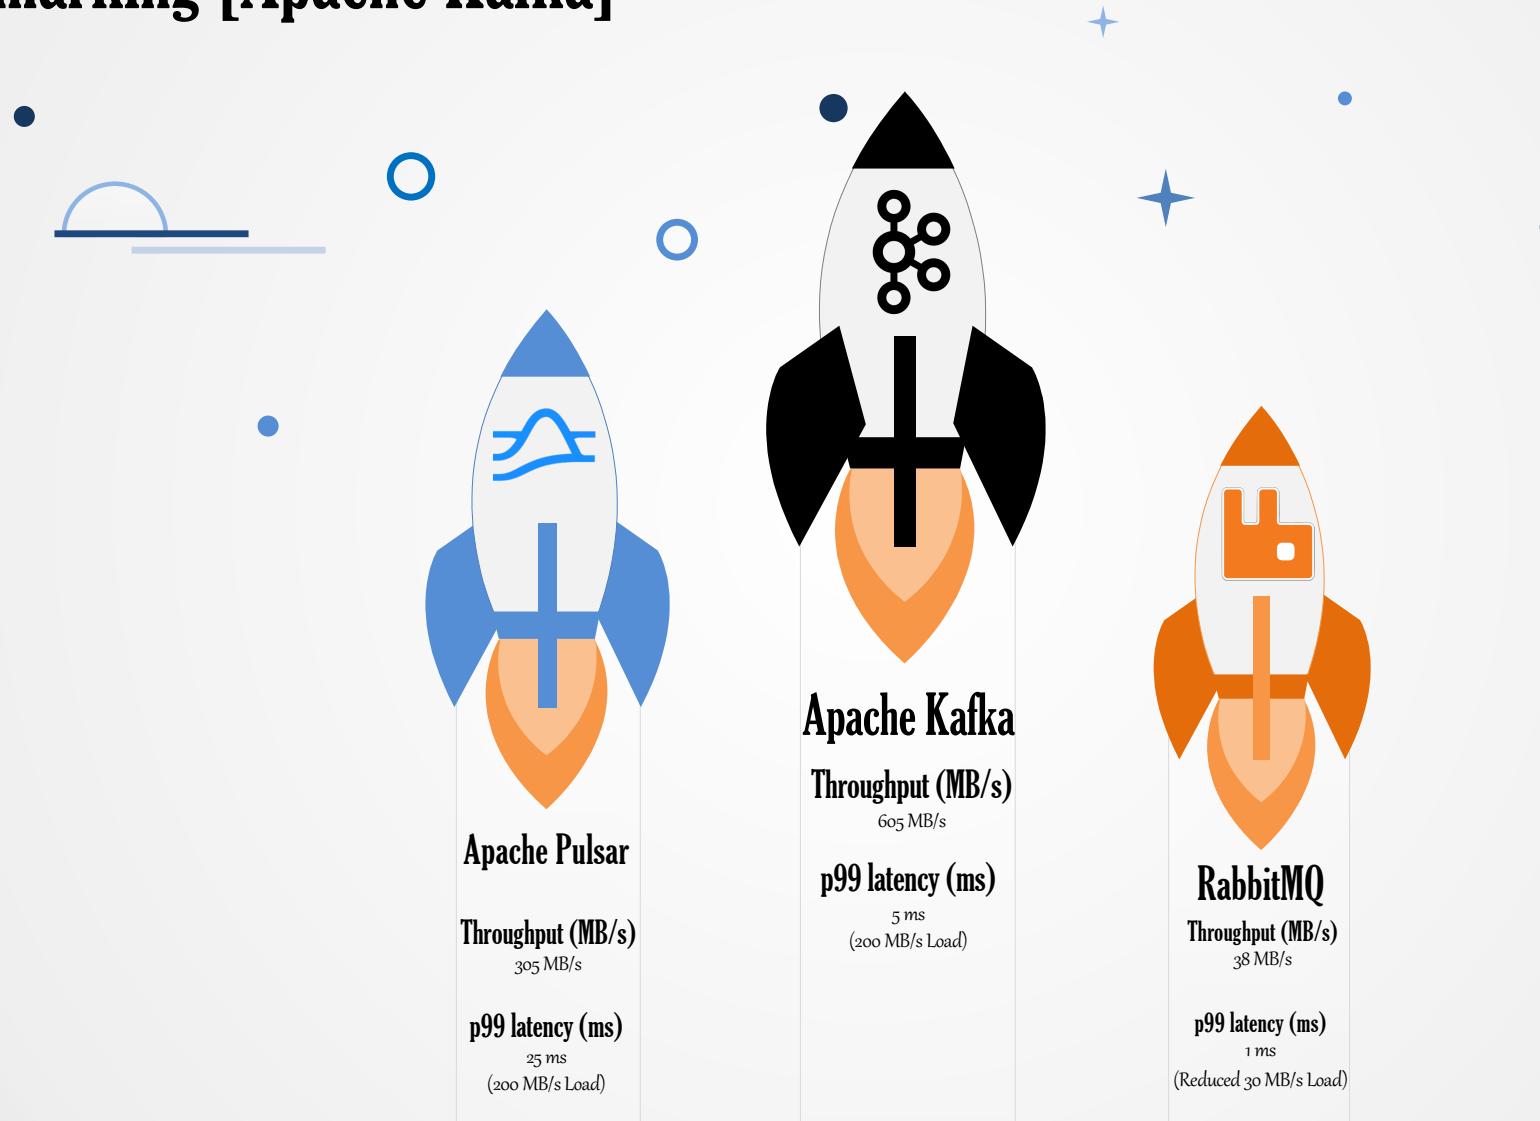


Streams	Value	Default
Stand By Replicas Configuration	1	0

Consumer	Value	Default
Session.Timeout.Ms	3.000	10.000

Broker	Value	Default
Unclean.Leader.Election.Enable	True	False
Min.InSync.Replicas	1	1
Num.Recovery.Threads.per.Data.Dir	5	1

# Benchmarking [Apache Kafka]



# Apache Kafka Security [Features]

Apache Kafka is the Wild-West without Security

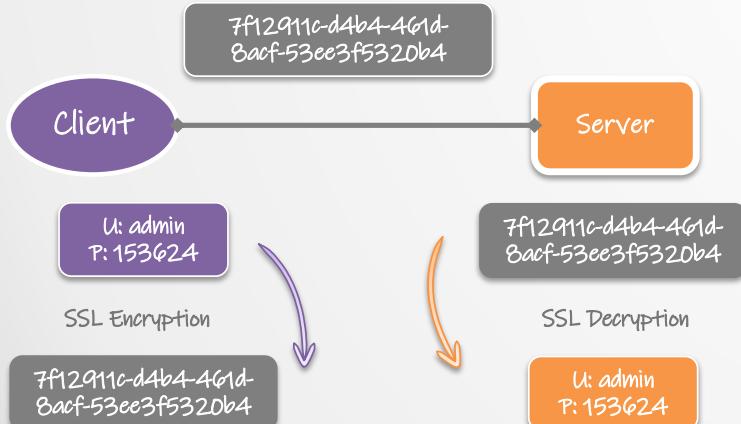


## Problems Security Solves

- **Encryption of Data In-Flight using SSL & TLS** ~ Encrypted Between Producers, Apache Kafka & Consumers
- **Authentication using SSL or SASL** ~ Producers & Consumers Authenticate on Apache Kafka Cluster. Identify Verification
- **Authorization using ACLs** ~ Access Control List for Authorized ~ Read & Write Topics

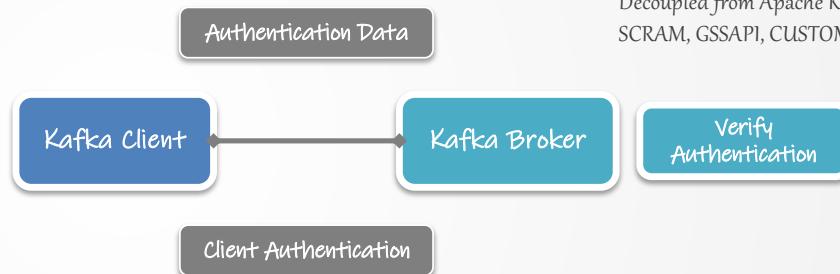
## Encryption SSL

for in-flight encryption only, disk is unencrypted



## Authentication (SSL & SASL)

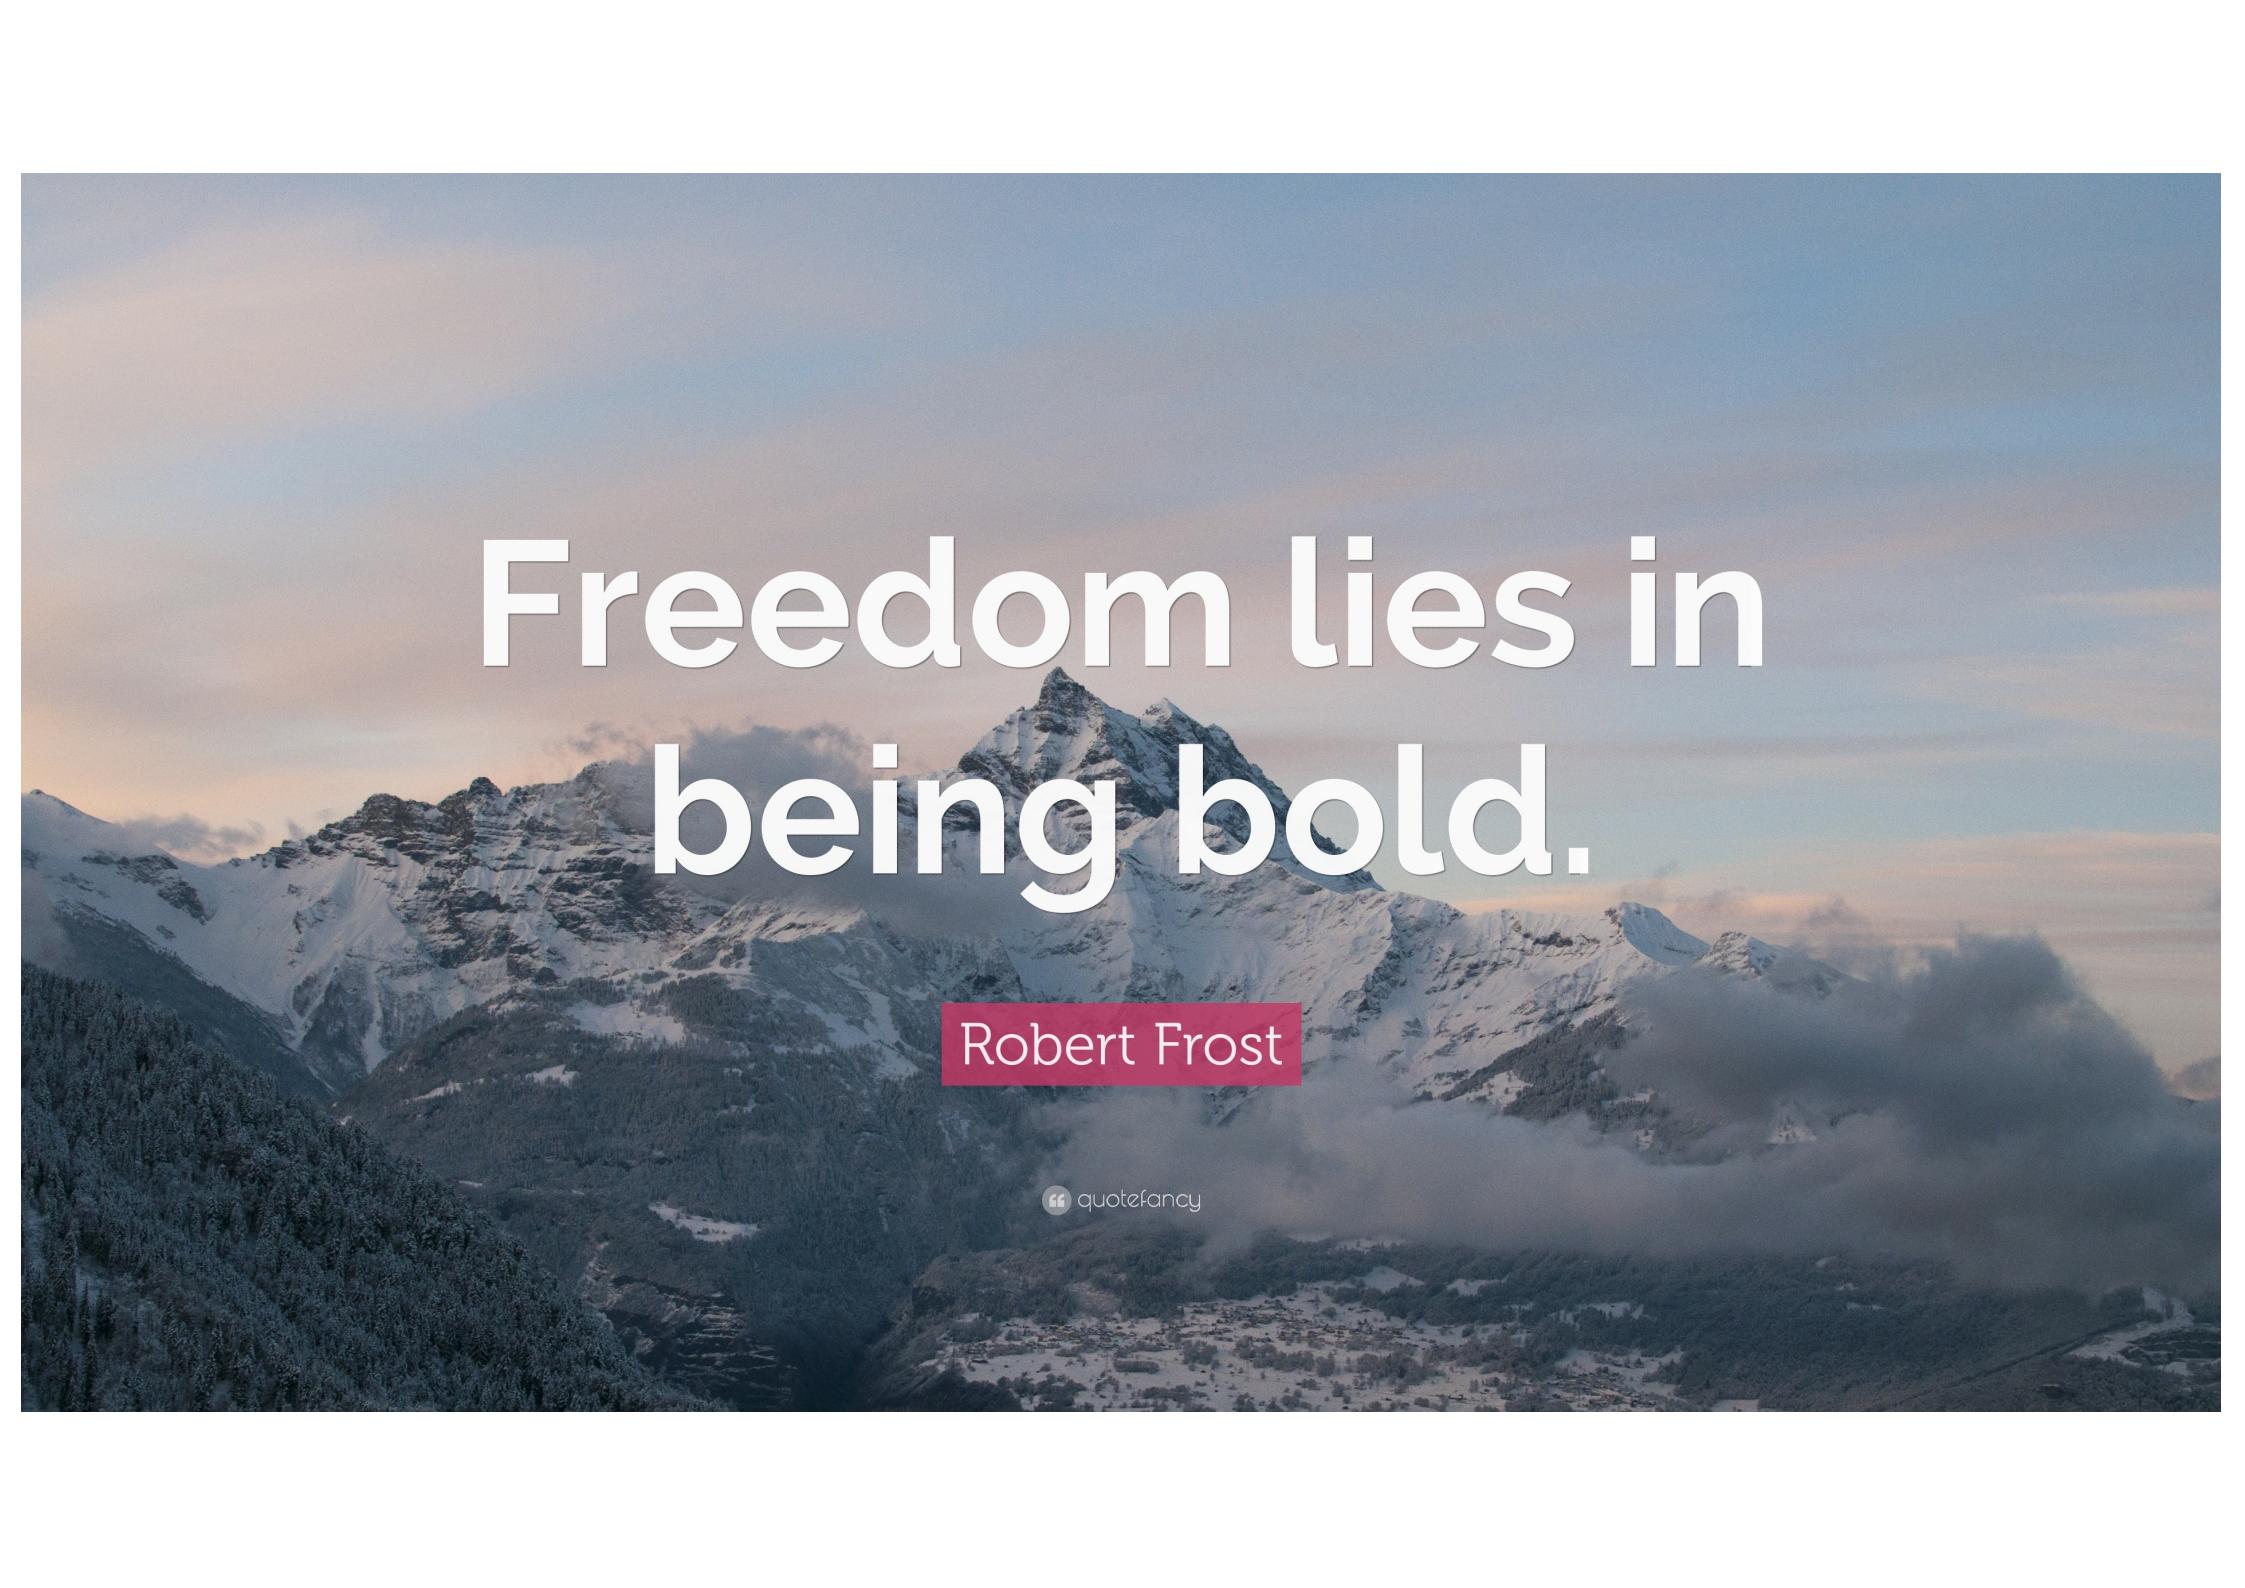
two ways to authenticate the kafka clients



## Authorization (ACL)

permissions to access resources

- user luan.moreno can view topic transactions
- user mateus.oliveira cannot view topic users

A wide-angle photograph of a mountain range at dusk or dawn. The peaks are covered in snow, and the sky is filled with soft, pastel-colored clouds transitioning from blue to orange and yellow. The foreground shows a valley with some snow and dark evergreen trees.

Freedom lies in  
being bold.

Robert Frost

# Business Use-Cases - Architecture Scenarios for Apache Kafka

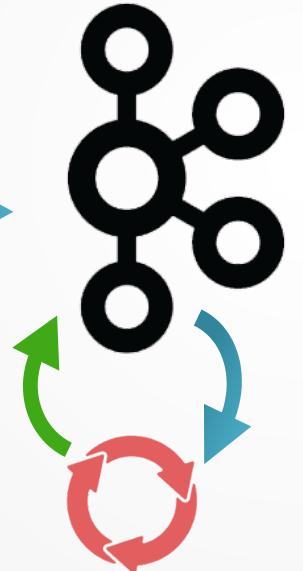
## Data Comes from Everywhere

data comes from databases, applications, need to work with financial data, geo-special data, commerce data, multiple formats and different schemas, sometimes even without schema, and we need in real-time but demand for batching



## Central Nervous System

with data comes from all places we need a platform that can ingest from many formats and be accessed from multiple consumers, all data in organization flows through this platform call apache kafka



## Streaming Processing

some process to transform the data into meaningful data for the organization

## APIs

applications consuming from kafka to access the output events, microservices era

## Data Warehouses

output data into modern data warehouses to be consumed by visualization apps

## Data Lake

output data or raw data sink into data lake to be consumed by data scientists and other data users.

# Apache Kafka as Integration Hub using KafkaConnect

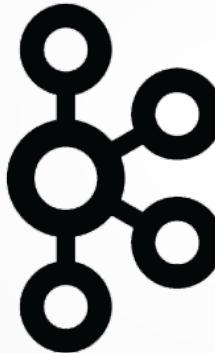
## Data Sources

databases such mongoDB, sql server, postgres and mysql



## Apache Kafka

apache kafka receive database data using kafka connect



### Kafka Connect Source

connectors build using kafka connect api to connect into databases and extract data



## Data Stores

databases such oracle, snowflake, bigquery



### Kafka Connect Sink

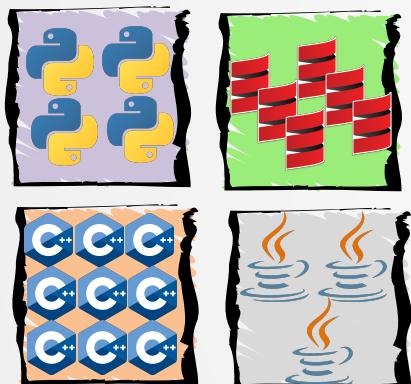
connectors build using kafka connect api to connect into databases and load data



# Apache Kafka as a Shock Absorber System

## APIs - Microservices

apis from different languages and services send data to data layer, somethings need to scale such as black Friday scenarios



## Apache Kafka

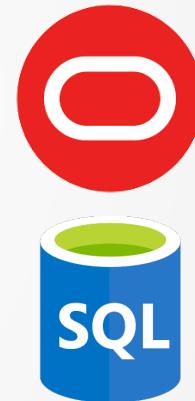
apache kafka works as shock absorber receiving the data from apis and scale if is needed



## Brokers Scale-Out

## Database

relational databases, not distributed systems





Without fear there  
cannot be courage.

Christopher Paolini

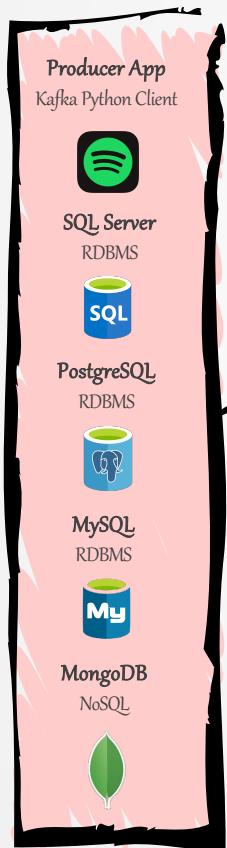


## Use-Case - Sneak Peak

Kappa Architecture

### Producers

sending data in



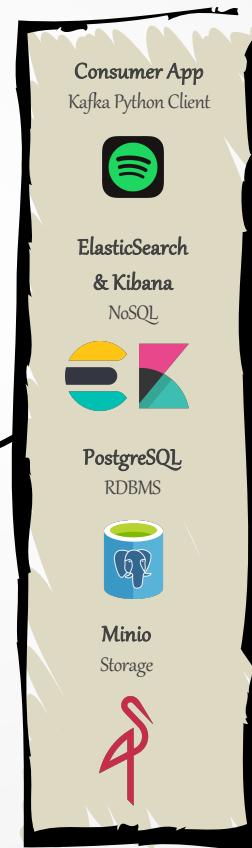
### Streams

process data inside of



### Consumers

getting data out



# [Orion] - Big Data as a Service

data-driven architecture for microservices & data pipelines. this process allows tighten and seamless integration between applications and data stores for analytics at scale.

**Application Ingestion Layer**  
Kafka Producer API



**Enterprise Data Hub [EDH]**  
Apache Kafka [Strimzi]



**Processing Layer**  
KSQLDB, Apache Spark & Faust



**OSS Data Serving**  
Apache Druid & YugaByteDB



**Data Visualization**  
Metabase & Superset



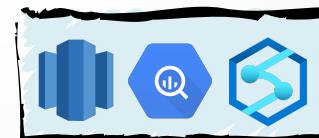
**Data Store Ingestion Layer**  
Kafka Connect Source API



**Deep Storage**  
MinIO



**Cloud Data Serving**  
Redshift, BigQuery, Synapse Analytics



**Continuous Delivery**  
ArgoCD - GitOps



**Orchestration Layer**  
Airflow



**Kubernetes**

Open-Source Container-Orchestration System



**Integration & Exploration Layer**  
Trino & Zeppelin



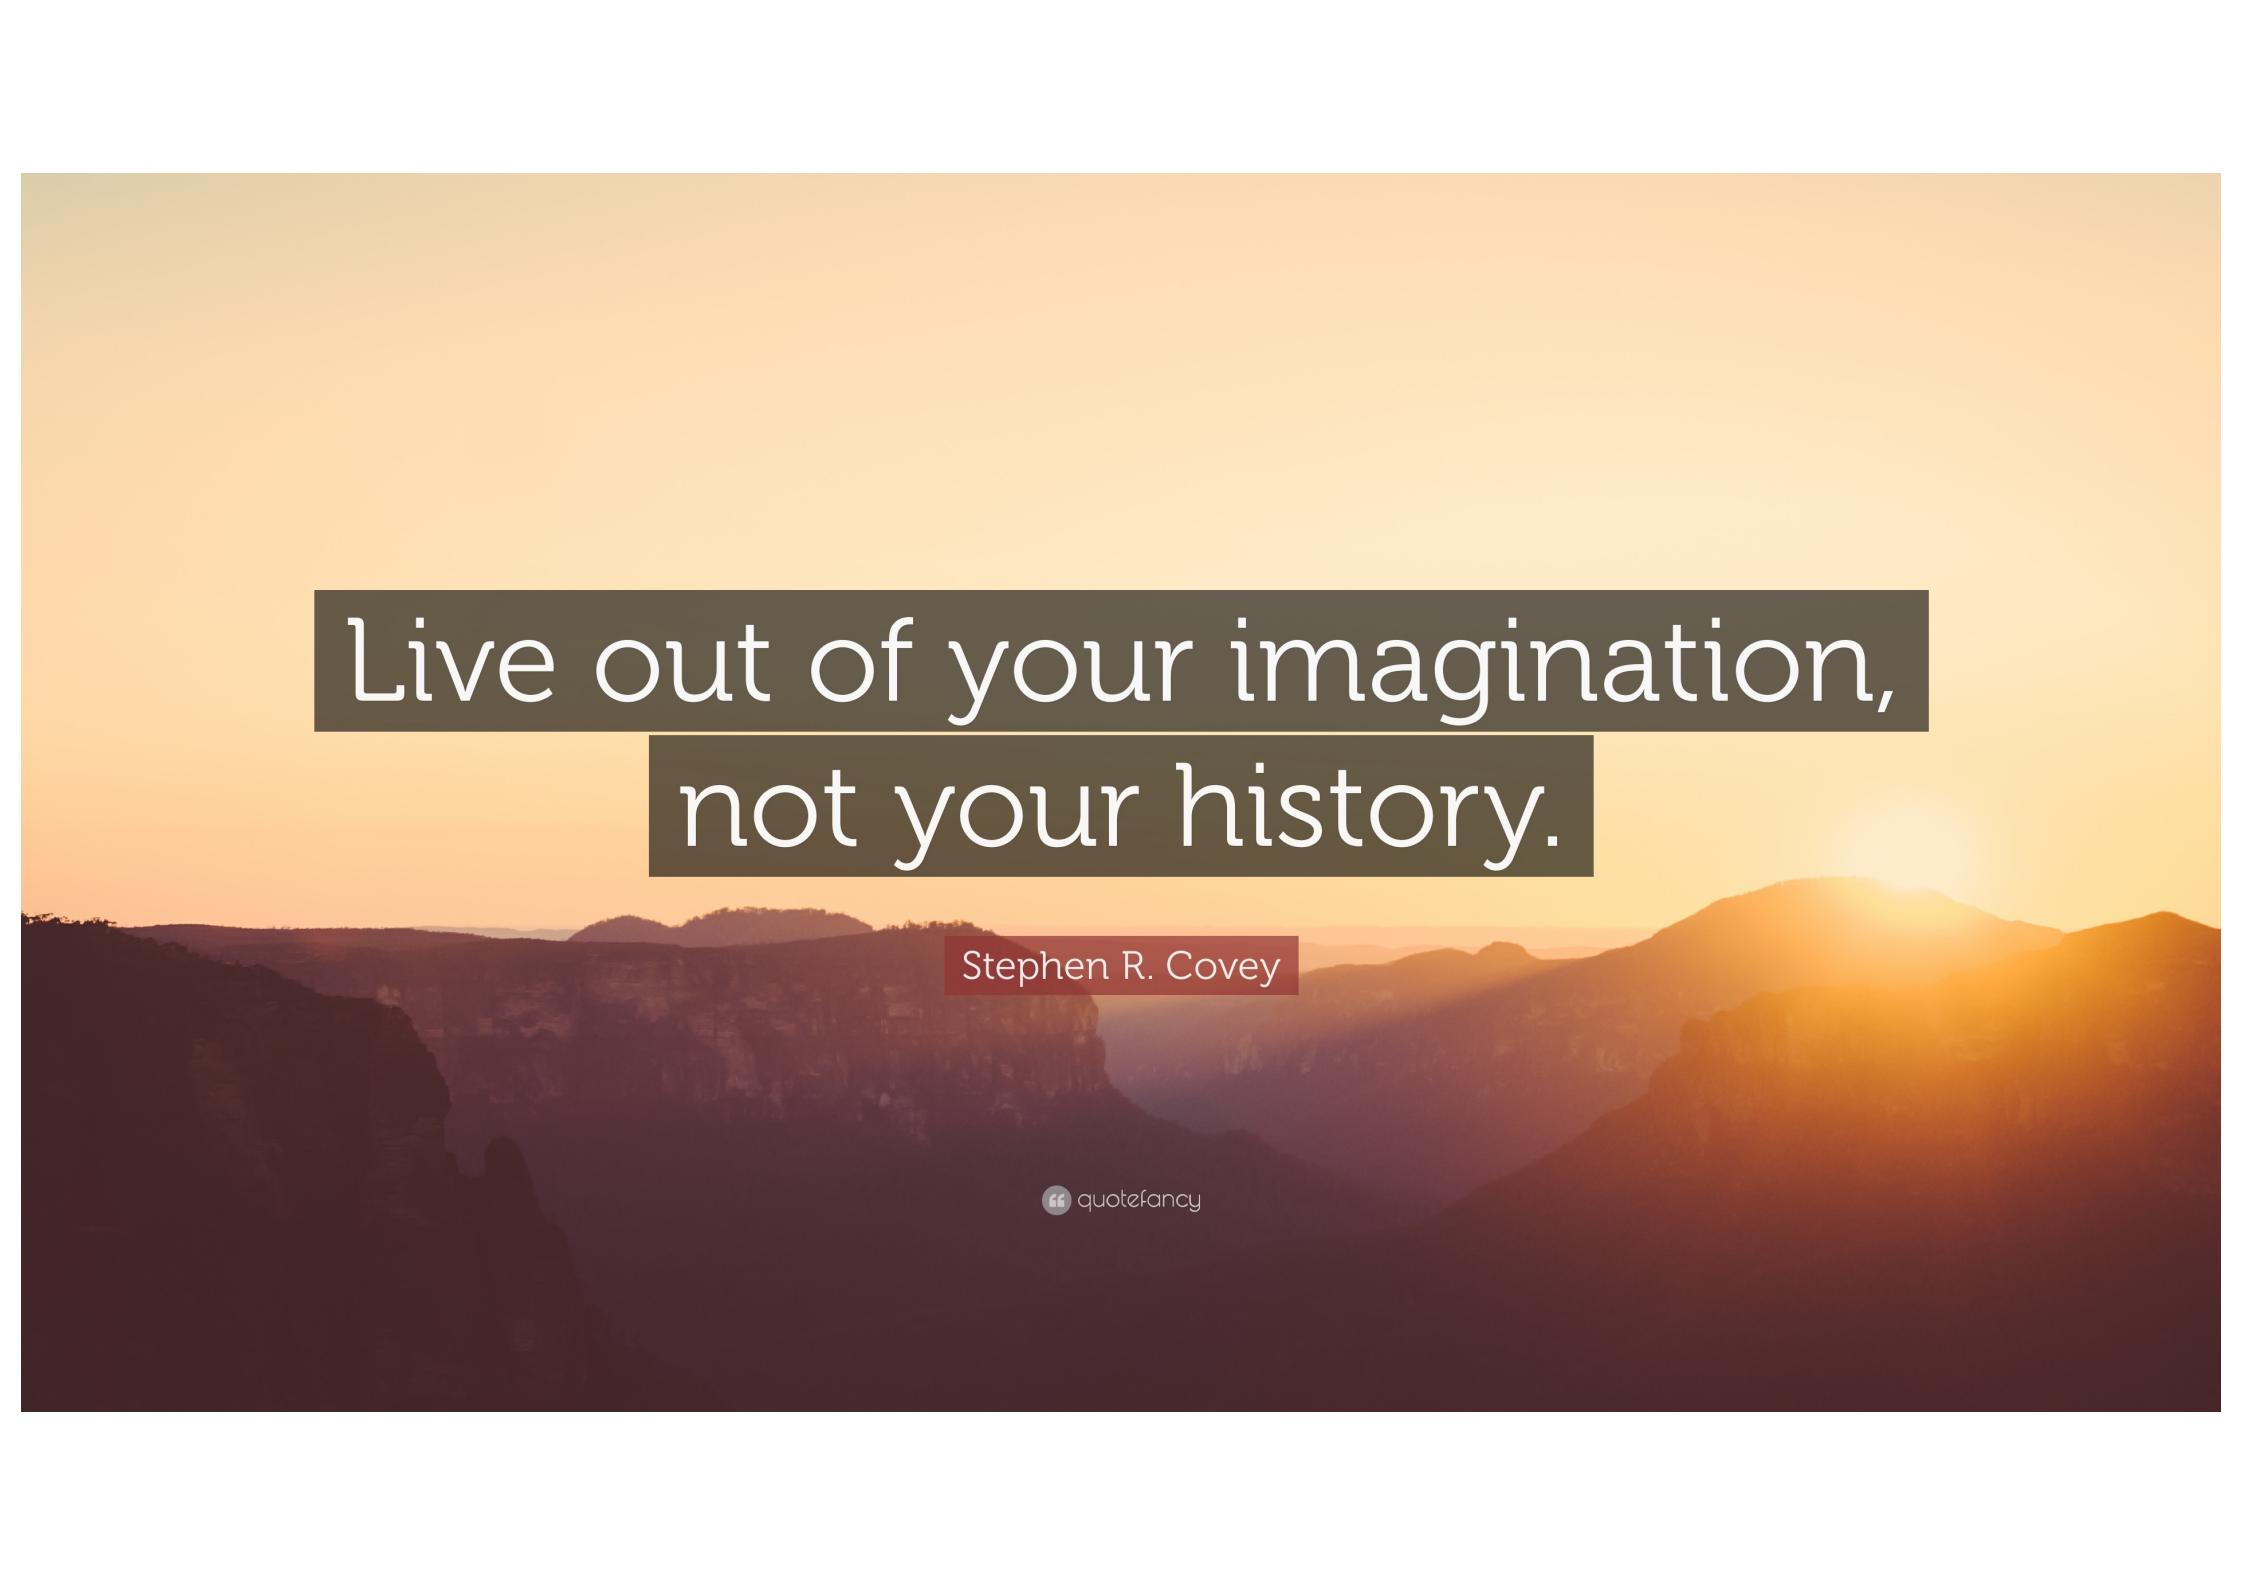
**Logging**  
FileBeat, ElasticSearch & Kibana



**Monitoring**  
Prometheus, Alert Manager & Grafana

## *Kappa Architecture*



The background of the image is a photograph of a mountainous landscape during a golden hour. The sky is filled with warm, orange and yellow hues, transitioning into darker blues and purples at the top. In the foreground, dark silhouettes of mountain peaks and valleys are visible against the bright sky. The overall atmosphere is serene and inspiring.

Live out of your imagination,  
not your history.

Stephen R. Covey



Focus on the solution,  
not on the problem.

Jim Rohn

# Data Engineer [Certifications]

Data Engineer Career ~ Part 1



## Microsoft Certified: Azure Data Engineer Associate

design and implement the management, monitoring, security, and privacy of data using the full stack of azure data services to satisfy business needs.



## Google Professional Data Engineer

enables data-driven decision making by collecting, transforming, and publishing data.



## AWS Certified Big Data - Specialty

certification is intended for individuals who perform complex big data analyses with at least two years of experience using aws technology.



## Confluent Certified Developer for Apache Kafka (CCDAK)

- Application Design ~ 40%
- Development ~ 30%
- Deployment, Testing & Monitoring ~ 30%



## Confluent Certified Operator for Apache Kafka (CCOAK)

- Kafka Fundamentals ~ 15%
- Managing, Configuring & Optimizing ~ 30%
- Kafka Security ~ 15%
- Designing, Troubleshooting & Integrating Systems ~ 40%

# Data Engineer [Study]

Data Engineer Career ~ Part 2



# *Luan Moreno M. Maciel*



*YouTube*  
luanmorenommaciel



*LinkedIn*  
Luan Moreno Medeiros Maciel



*Facebook*  
Luan Moreno Medeiros Maciel



*Instagram*  
engenhariadedados



*Meetup*  
BSB-Al-Big-Data-Analytics





# Thank You



*One Way Solution*



**ONEWAY**  
SOLUTION